

# Finishing Up CALL & Intro To Digital Circuits

# Example C Program: Hello.c

```
#include <stdio.h>
int main()
{
    printf("Hello, %s\n", "world");
    return 0;
}
```

# Compiled Hello.c: Hello.s

```
.text
.align 2
.globl main
main:
    addi sp,sp,-16
    sw    ra,12(sp)
    lui   a0,%hi(string1)
    addi a0,a0,%lo(string1)
    lui   a1,%hi(string2)
    addi a1,a1,%lo(string2)
    call printf
    lw    ra,12(sp)
    addi sp,sp,16
    li    a0,0
    ret
.section .rodata
.balign 4
string1:
    .string "Hello, %s!\n"
string2:
    .string "world"
```

# Directive: enter text section  
# Directive: align code to  $2^2$  bytes  
# Directive: declare global symbol main  
# label for start of main  
# allocate stack frame  
# save return address  
# compute address of  
# string1  
# compute address of  
# string2  
# call function printf  
# restore return address  
# deallocate stack frame  
# load return value 0  
# return  
# Directive: enter read-only data section  
# Directive: align data section to 4 bytes  
# label for first string  
# Directive: null-terminated string  
# label for second string  
# Directive: null-terminated string

# Assembled Hello.s: Linkable Hello.o

```
00000000 <main>:  
0: ff010113 addi sp,sp,-16  
4: 00112623 sw ra,12(sp)  
8: 00000537 lui a0,0x0      # addr placeholder  
c: 00050513 addi a0,a0,0    # addr placeholder  
10: 000005b7 lui a1,0x0      # addr placeholder  
14: 00058593 addi a1,a1,0    # addr placeholder  
18: 00000097 auipc ra,0x0    # addr placeholder  
1c: 000080e7 jalr ra        # addr placeholder  
20: 00c12083 lw ra,12(sp)  
24: 01010113 addi sp,sp,16  
28: 00000513 addi a0,a0,0  
2c: 00008067 jalr ra
```

# Linked Hello.o: a.out

```
000101b0 <main>:  
101b0: ff010113 addi sp,sp,-16  
101b4: 00112623 sw ra,12(sp)  
101b8: 00021537 lui a0,0x21  
101bc: a1050513 addi a0,a0,-1520 # 20a10 <string1>  
101c0: 000215b7 lui a1,0x21  
101c4: a1c58593 addi a1,a1,-1508 # 20a1c <string2>  
101c8: 288000ef jal ra,10450 # <printf>  
101cc: 00c12083 lw ra,12(sp)  
101d0: 01010113 addi sp,sp,16  
101d4: 00000513 addi a0,0,0  
101d8: 00008067 jalr ra
```

# LUI/ADDI Address Calculation in RISC-V

Target address of <string1> is **0x00020 A10**

Instruction sequence **LUI 0x00020, ADDI 0xA10** does not quite work because immediates in RISC-V are sign extended (and 0xA10 has a 1 in the high order bit)!

$$0x00020\ 000 + 0xFFFFF\ A10 = 0x0001F\ A10 \text{ (Off by } 0x00001\ 000)$$

So we get the right address if we calculate it as follows:

$$(0x00020\ 000 + 0x00001\ 000) + 0xFFFFF\ A10 = 0x00020\ A10$$

What is 0xFFFFF A10?

$$\text{Twos complement of } 0xFFFFF\ A10 = 0x00000\ 5EF + 1 = 0x00000\ 5F0 = 1520_{ten}$$

$$\text{So } 0xFFFFF\ A10 = -1520_{ten}$$

Instruction sequence **LUI 0x00021, ADDI -1520** calculates **0x00020 A10**

# Static vs. Dynamically Linked Libraries

- What we've described is the traditional way: **statically-linked approach**
  - Library is now part of the executable, so if the library updates, we don't get the fix (have to recompile if we have source)
  - Includes the entire library even if not all of it will be used
  - Executable is self-contained
- Alternative is **dynamically linked libraries (DLL)**, common on Windows & UNIX platforms

# Dynamically Linked Libraries

[en.wikipedia.org/wiki/Dynamic\\_linking](https://en.wikipedia.org/wiki/Dynamic_linking)

- Space/time issues
  - + Storing a program requires less disk space
  - + Sending a program requires less time
  - + Executing two programs requires less memory (if they share a library)
  - – At runtime, there's time overhead to do link
- Upgrades
  - + Replacing one file (libXYZ.so) upgrades every program that uses library “XYZ”
  - – Having the executable isn't enough anymore

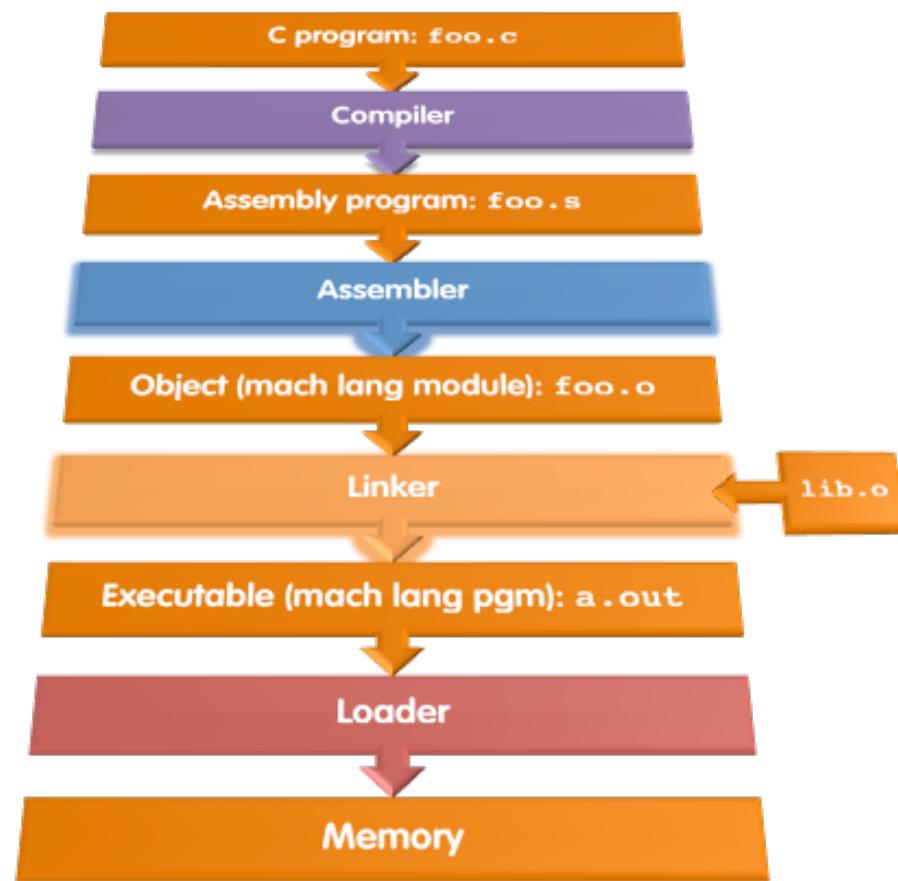
*Overall, dynamic linking adds quite a bit of complexity to the compiler, linker, and operating system. However, it provides many benefits that often outweigh these*

# Dynamically Linked Libraries

- Prevailing approach to dynamic linking uses machine code as the “lowest common denominator”
  - Linker does not use information about how the program or library was compiled (i.e., what compiler or language)
  - Can be described as “linking at the machine code level”
  - This isn’t the only way to do it ...
- Also these days will ***randomize layout*** (Address Space Layout Randomization)
  - Acts as a defense to make exploiting C memory errors substantially harder, as modern exploitation requires jumping to pieces of existing code (“Return oriented programming”) to counter another defense (marking heap & stack unexecutable, so attacker can’t write code into just anywhere in memory).

# In Conclusion...

- Compiler converts a single HLL file into a single assembly language file.
- Assembler removes pseudo-instructions, converts what it can to machine language, and creates a checklist for the linker (relocation table). A .s file becomes a .o file.
  - Does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses.
  - Enables separate compilation, libraries that need not be compiled, and resolves remaining addresses
- Loader loads executable into memory and begins execution.



# Announcements...

- We are nearly done grading the Midterm
  - Should be out by Friday Morning
- Project 2.1 and 2.2 are now out...
  - Remember, we have slip days:  
If something happens you don't need to ask permission to use one
- No lab/OH on Monday
  - Lab3 and Lab4 due dates pushed back 1 week
- And don't cheat on the project...
  - Nick went nuclear on a little less than a dozen violators on Project 1...
  - And if I do catch you, ***don't lie to me***

**Purported DPRK design for a ~150kT thermonuclear weapon**



# Clicker Question:

## Midterm 1 (In Modern American Written English)

- How was Midterm 1?

A. 😊

B. 😐

C. 😢

D. 😧

E. 💩

# Hardware Design

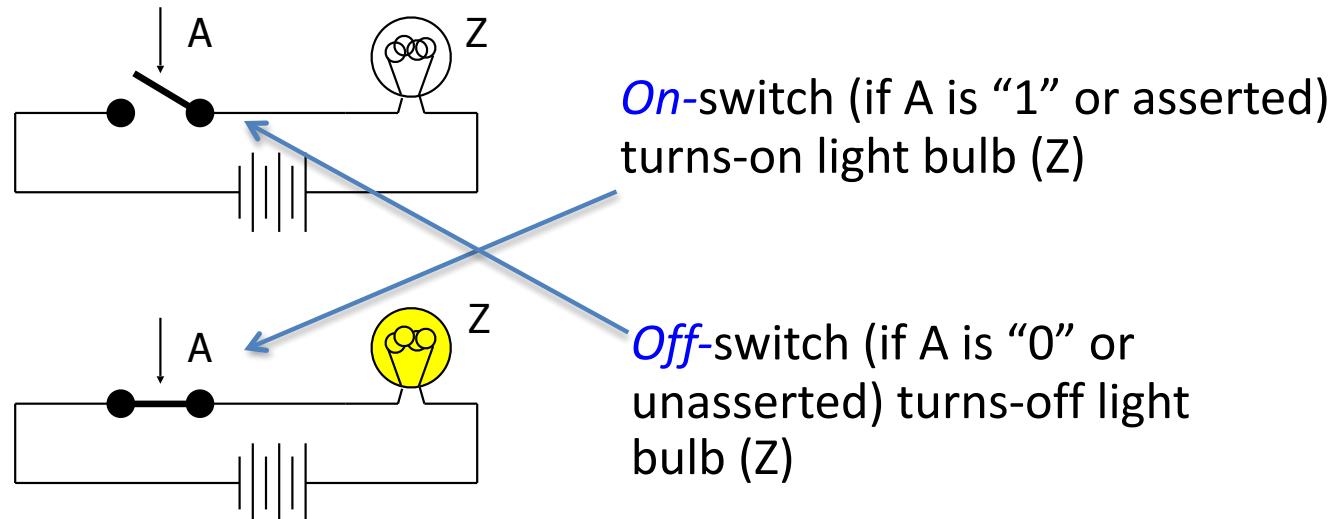
- Next several weeks: how a modern processor is built, starting with basic elements as building blocks
- Why study hardware design?
  - Understand capabilities and limitations of HW in general and processors in particular
  - What processors can do fast and what they can't do fast  
(avoid slow things if you want your code to run fast!)
  - Background for more in-depth HW courses (EECS 151, CS 152)
  - Hard to know what you'll need for next 30 years
  - There is only so much you can do with standard processors: you may need to design own custom HW for extra performance
    - Even some commercial processors today have customizable hardware!

# Synchronous Digital Systems: Almost Every Processor etc...

- ***Synchronous:***
  - All operations coordinated by a central clock
    - “Heartbeat” of the system!
    - **Anynchronous** systems much much much harder to design & debug
- ***Digital:***
  - Represent all values by discrete values
  - Two binary digits: **1** and **0**, or **True** and **False**
    - Electrical signals are treated as 1's and 0's
      - 1 and 0 are complements of each other
  - High/low voltage for true/false, 1/0
  - These days, even a lot of **analog** circuitry is best done as:
    - Put it through an analog to digital converter, do it all digitally in a synchronous circuit, and output to a digital-to-analog converter

# Switches: Basic Element of Physical Implementations

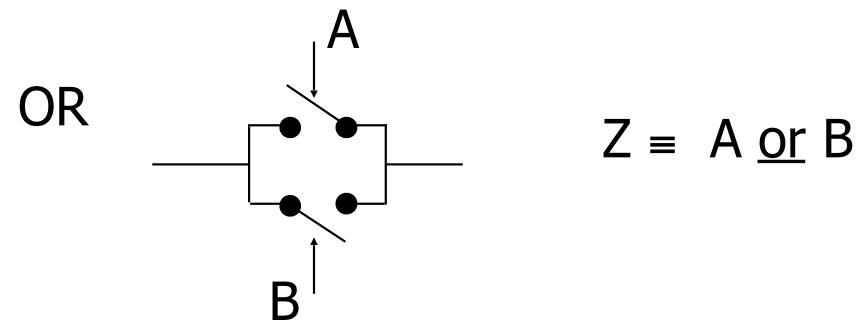
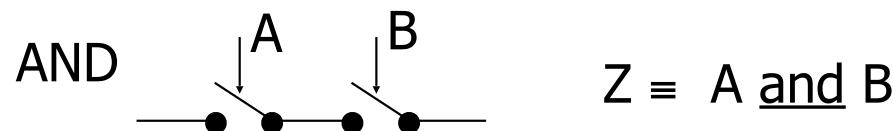
- Implementing a simple circuit (arrow shows action if wire changes to “1” or is *asserted*):



$$Z \equiv A$$

# Switches (cont'd)

- Compose switches into more complex ones (Boolean functions):



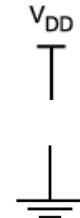
# Historical Note

- Early computer designers built ad hoc circuits from switches
  - Began to notice common patterns in their work: ANDs, ORs, ...
  - Master's thesis (by Claude Shannon, 1940) made link between work and 19th Century Mathematician George Boole
    - Called it "Boolean" in his honor
    - Could apply math to give theory to hardware design, minimization, ...



# Transistors

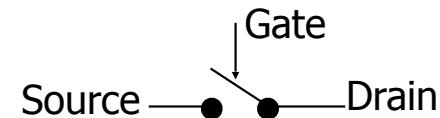
- High voltage ( $V_{dd}$ ) represents 1, or true
  - The Raspberry Pi:  $V_{dd} \sim 1.2$  Volt
- Low voltage (0 Volt or Ground) represents 0, or false
- Pick a midpoint voltage to decide if a 0 or a 1
  - Voltage greater than midpoint = 1
  - Voltage less than midpoint = 0
  - This removes noise as signals propagate –  
a big advantage of digital systems over analog systems
- If one switch can control another switch, we can build a computer!
- Our switches: CMOS transistors



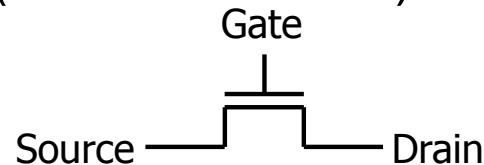
# CMOS Transistor Networks

- Modern digital systems designed in CMOS
  - MOS: Metal-Oxide on Semiconductor
    - Describes how the transistors are constructed
    - These are "Field Effect Transistors" (MOSFETs)
  - C for complementary: use **pairs** of normally-on and normally-off switches
- CMOS transistors act as voltage-controlled switches
  - Similar, though easier to work with, than electro-mechanical relay switches from earlier era
  - Use energy primarily when switching
    - But not completely: they do "leak" a bit

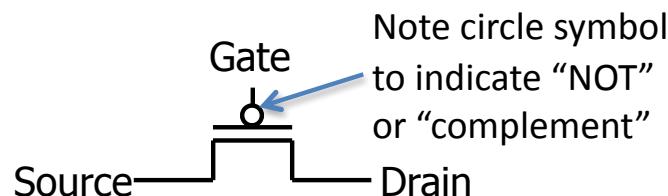
# CMOS Transistors



- Three terminals: source, gate, and drain
- Switch action:  
if voltage on gate terminal is (some amount) higher/lower than source terminal then conducting path established between drain and source terminals (switch is closed)

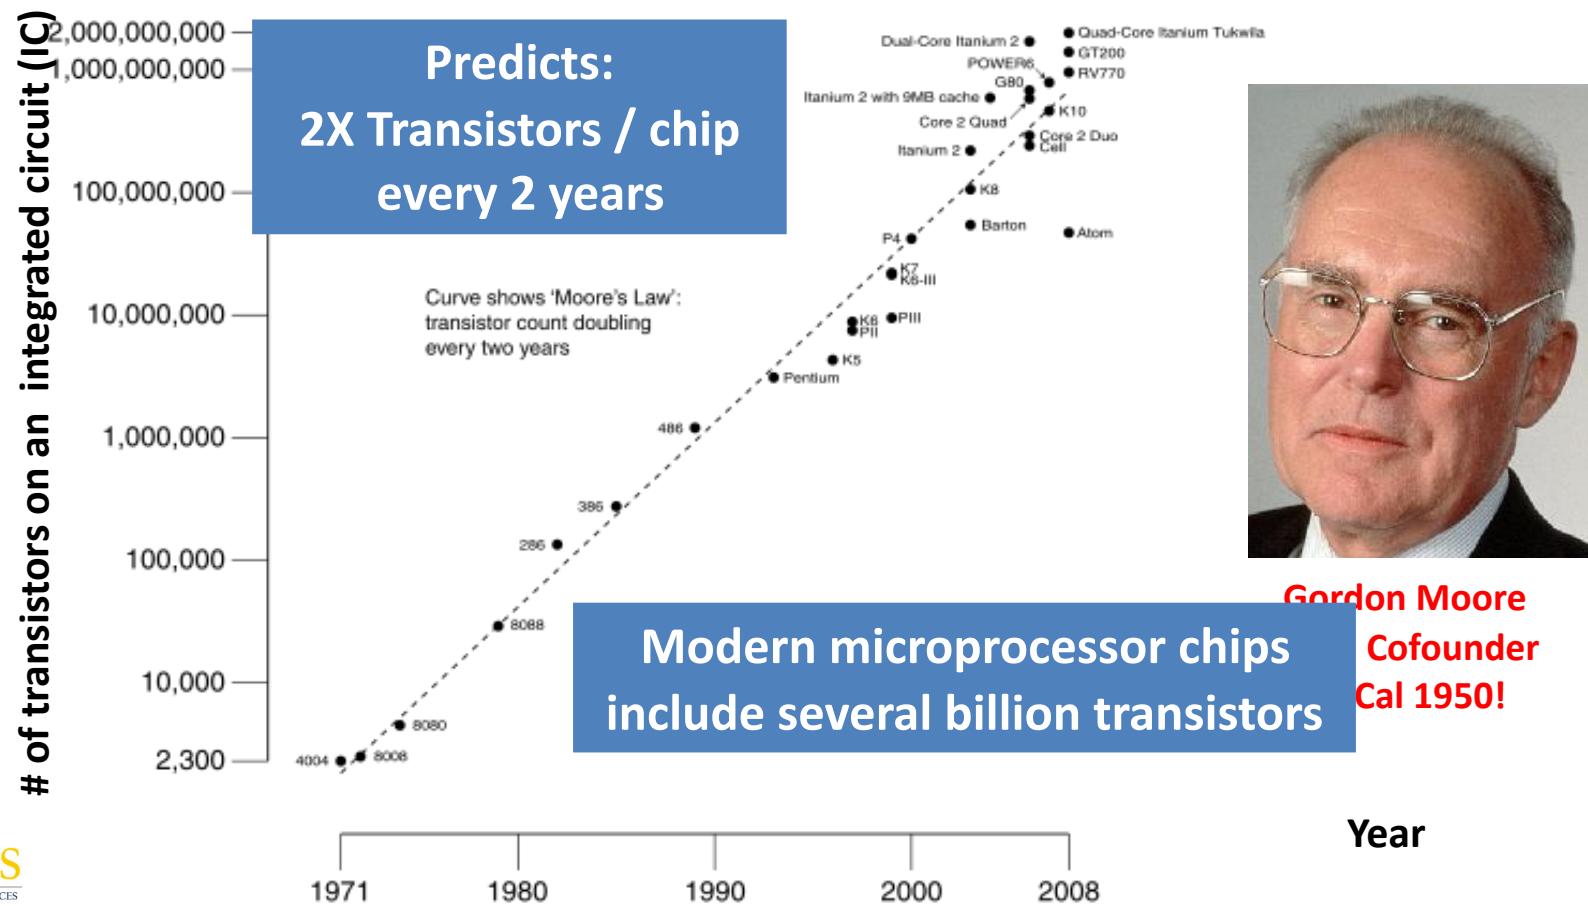


n-channel transistor  
off when voltage at Gate is low  
on when:  
 $voltage(Gate) > voltage(\text{Threshold})$

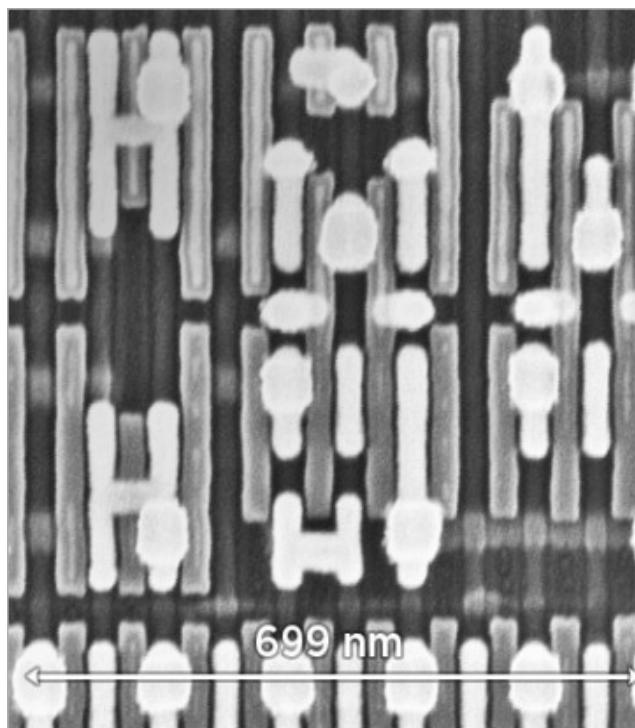


p-channel transistor  
on when voltage at Gate is low  
off when:  
 $voltage(Gate) > voltage(\text{Threshold})$

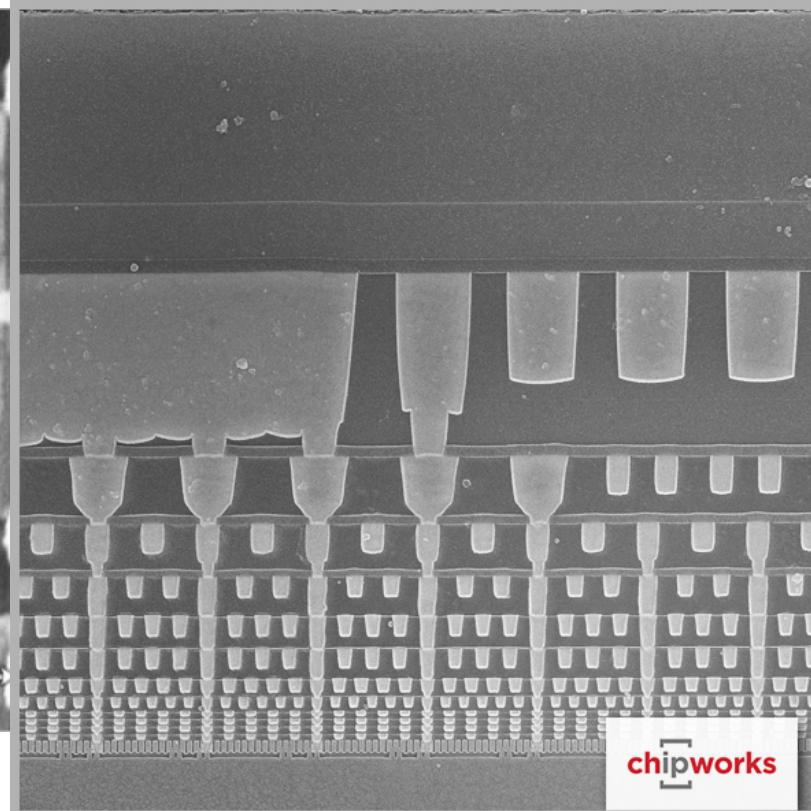
# #2: Moore's Law



# Intel 14nm Technology

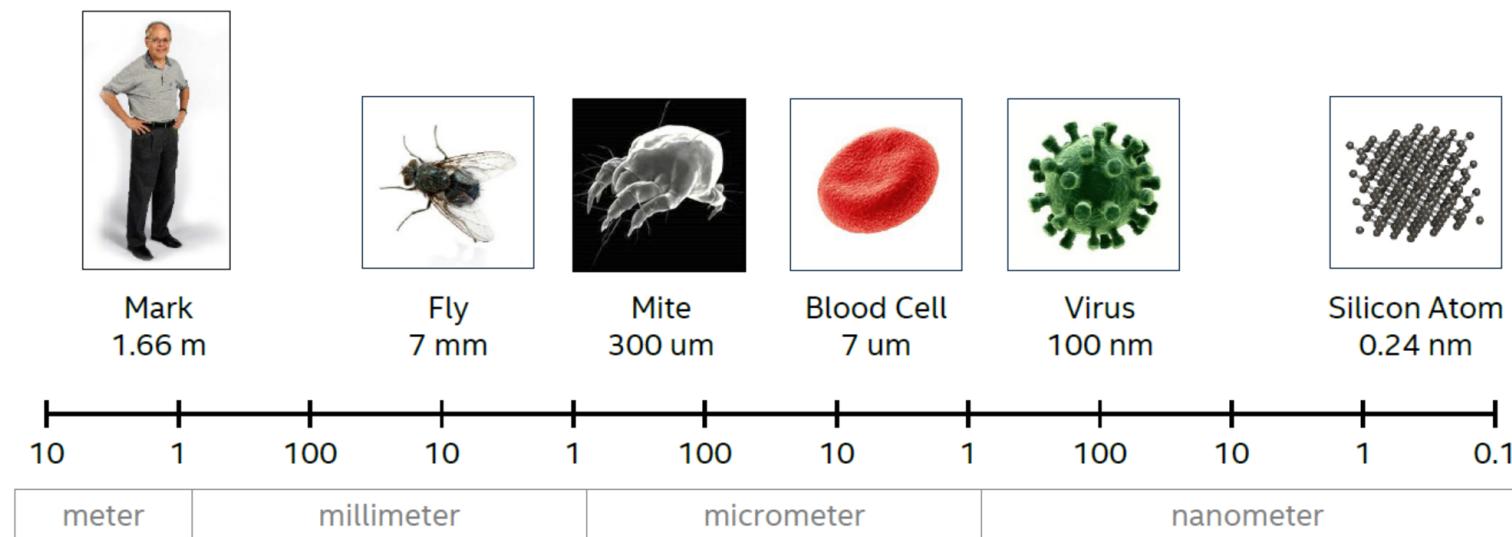


Plan view of transistors



Side view of wiring layers

# Sense of Scale



Source: Mark Bohr, IDF14

# CMOS Circuit Rules

- ***Don't pass*** weak values => Use Complementary Pairs
  - N-type transistors pass weak 1's ( $V_{dd} - V_{th}$ )
  - N-type transistors pass strong 0's (ground)
  - Use N-type transistors only to pass 0's (N for negative)
  - Converse for P-type transistors: Pass weak 0s, strong 1s
    - Pass weak 0's ( $V_{th}$ ), strong 1's ( $V_{dd}$ )
    - Use P-type transistors only to pass 1's (P for positive)
  - Use pairs of N-type and P-type to get strong values
- ***Never*** leave a wire undriven (in this class)
  - Make sure there's always a path to Vdd or GND
- ***Never*** create a path from Vdd to GND (ground)
  - This would short-circuit the power supply!

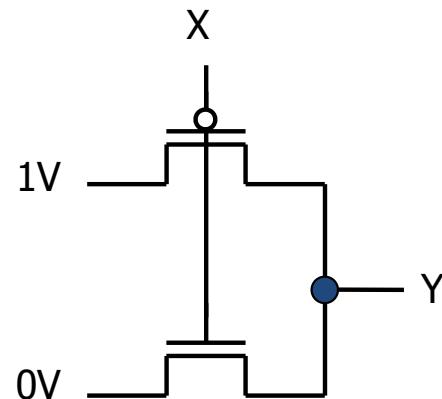
# CMOS Networks

p-channel transistor

on when voltage at Gate is low

off when:

voltage(Gate) > voltage (Threshold)



what is the  
relationship  
between x and y?

x	y
0 Volt (GND)	1 Volt (Vdd)
1 Volt (Vdd)	0 Volt (GND)

n-channel transistor

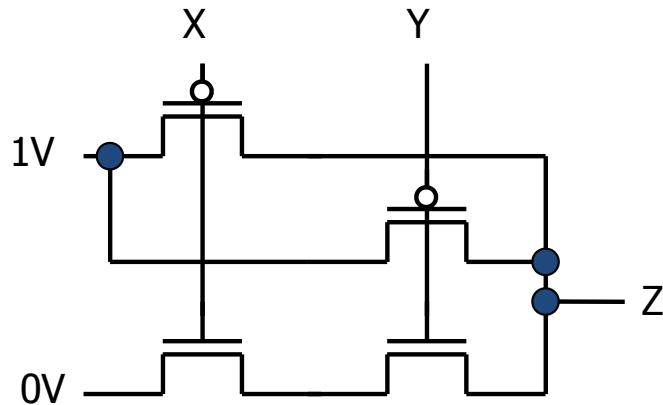
off when voltage at Gate is low

on when:

voltage(Gate) > voltage (Threshold)

Called an *inverter* or *not gate*

# Two-Input Networks

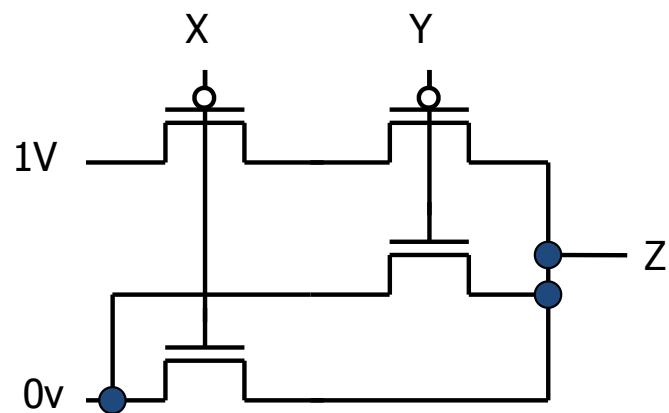


what is the  
relationship between x, y and z?

x	y	z
0 Volt	0 Volt	1 Volt
0 Volt	1 Volt	1 Volt
1 Volt	0 Volt	1 Volt
1 Volt	1 Volt	0 Volt

Called a *NAND gate*  
(*NOT AND*)

# Clickers/Peer Instruction

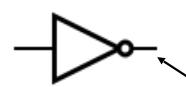
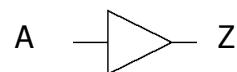


X	Y	Z	
0 Volt	0 Volt	A	Volts
0 Volt	1 Volt	B	Volts
1 Volt	0 Volt	C	Volts
1 Volt	1 Volt	D	Volts

# Combinational Logic Symbols

- Common combinational logic systems have standard symbols called logic gates

- Buffer, NOT



- AND, NAND



- OR, NOR



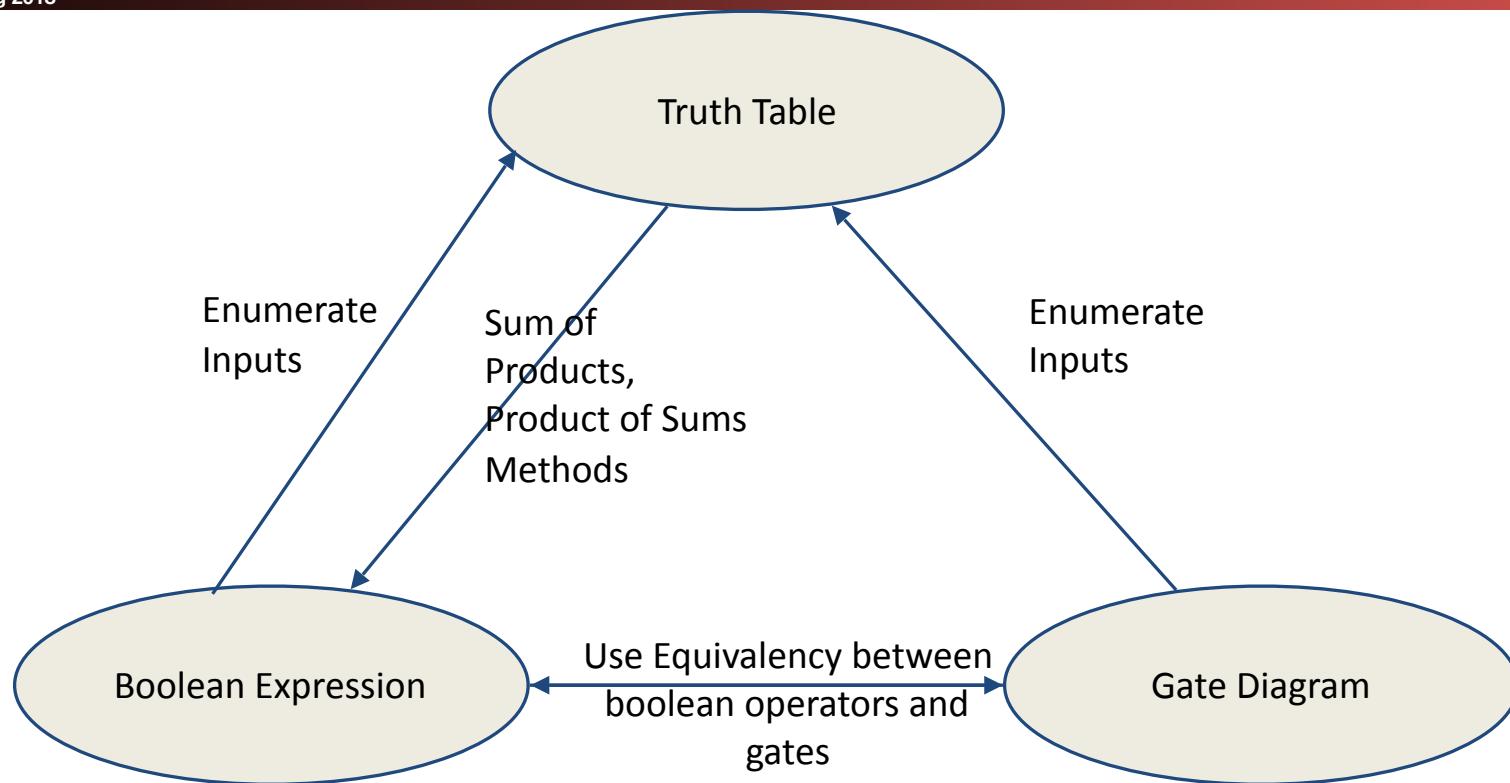
Inverting versions (NOT, NAND, NOR) easiest to implement with CMOS transistors (the switches we have available and use most)

# Boolean Algebra

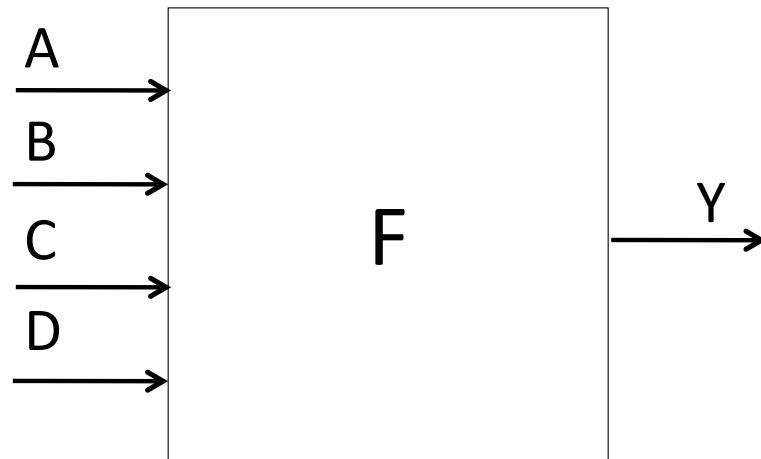
- Use plus “+” for OR
  - “logical sum”
- Use product for AND ( $a \cdot b$  or implied via  $ab$ )
  - “logical product”
- “Hat” to mean complement (NOT)
- Thus
- $ab + a + \bar{c}$   
=  $a \cdot b + a + \bar{c}$   
= (a AND b) OR a OR (NOT c )



# Representations of Combinational Logic (groups of logic gates)



# Truth Tables for Combinational Logic



Exhaustive list of the output value  
generated for each combination of inputs

How many logic functions can be defined  
with N inputs?

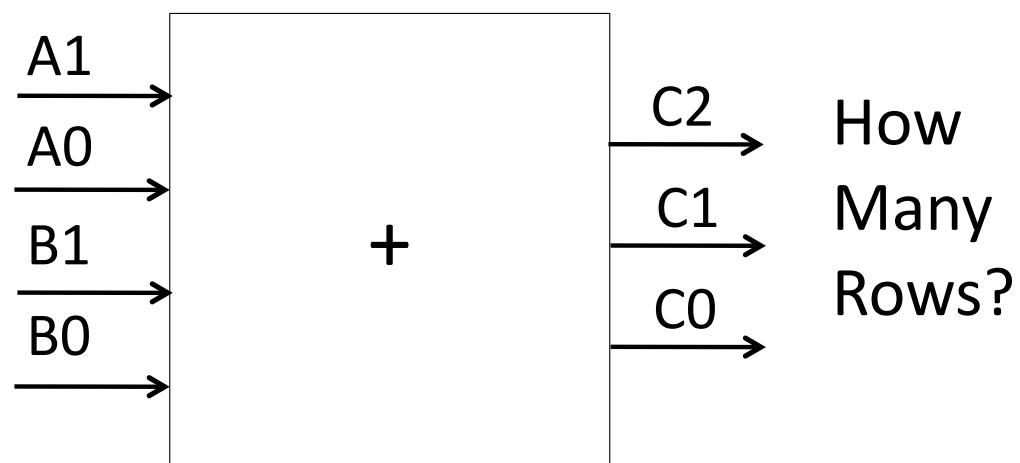
a	b	c	d	y
0	0	0	0	$F(0,0,0,0)$
0	0	0	1	$F(0,0,0,1)$
0	0	1	0	$F(0,0,1,0)$
0	0	1	1	$F(0,0,1,1)$
0	1	0	0	$F(0,1,0,0)$
0	1	0	1	$F(0,1,0,1)$
0	1	1	0	$F(0,1,1,0)$
1	1	1	1	$F(0,1,1,1)$
1	0	0	0	$F(1,0,0,0)$
1	0	0	1	$F(1,0,0,1)$
1	0	1	0	$F(1,0,1,0)$
1	0	1	1	$F(1,0,1,1)$
1	1	0	0	$F(1,1,0,0)$
1	1	0	1	$F(1,1,0,1)$
1	1	1	0	$F(1,1,1,0)$
1	1	1	1	$F(1,1,1,1)$

# Truth Table Example #1:

$y = F(a,b)$ : 1 iff  $a \neq b$

<b>a</b>	<b>b</b>	<b>y</b>
0	0	0
0	1	1
1	0	1
1	1	0

# Truth Table Example #2: 2-bit Adder



A	B	C
$a_1a_0$	$b_1b_0$	$c_2c_1c_0$
<hr/>		

# Truth Table Example #3: 32-bit Unsigned Adder

A	B	C
000 ... 0	000 ... 0	000 ... 00
000 ... 0	000 ... 1	000 ... 01
.	.	.
.	.	.
.	.	.
111 ... 1	111 ... 1	111 ... 10

How  
Many  
Rows?

# Truth Table Example #4: 3-input Majority Circuit

$Y =$

This is called *Sum of Products* form;  
Just another way to represent the TT  
as a logical expression

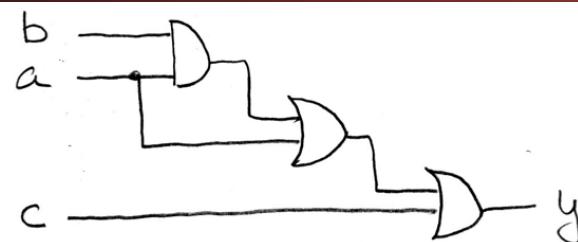
More simplified forms  
(fewer gates and wires)

a	b	c	y
0	0	0	0
0	0	1	0
0	1	0	0
0	1	1	1
1	0	0	0
1	0	1	1
1	1	0	1
1	1	1	1

# And in Conclusion, ...

- Multiple Hardware Representations
  - Analog voltages quantized to represent logic 0 and logic 1
  - Transistor switches form gates: AND, OR, NOT, NAND, NOR
  - Truth table mapped to gates for combinational logic design
  - Boolean algebra for gate minimization

# Boolean Algebra: Circuit & Algebraic Simplification



original circuit

$$y = ((ab) + a) + c$$

$$\begin{aligned} &= ab + a + c \\ &= a(b + 1) + c \\ &= a(1) + c \\ &= a + c \end{aligned}$$



equation derived from original circuit

algebraic simplification

simplified circuit

# Laws of Boolean Algebra

$X \bar{X} = 0$	$X + \bar{X} = 1$	Complementarity
$X 0 = 0$	$X + 1 = 1$	Laws of 0's and 1's
$X 1 = X$	$X + 0 = X$	Identities
$X X = X$	$X + X = X$	Idempotent Laws
$X Y = Y X$	$X + Y = Y + X$	Commutativity
$(X Y) Z = Z (Y Z)$	$(X + Y) + Z = Z + (Y + Z)$	Associativity
$X (Y + Z) = X Y + X Z$	$X + Y Z = (X + Y) (X + Z)$	Distribution
$\underline{\underline{X}} Y + X = X$	$\underline{(X + Y)} X = X$	Uniting Theorem
$\underline{\underline{X}} Y + X = \underline{\underline{X}} + Y$	$\underline{(X + Y)} X = \underline{X} Y$	Uniting Theorem v. 2
$\underline{\underline{X}} Y = X + Y$	$X + Y = X Y$	DeMorgan's Law

# Boolean Algebraic Simplification Example

$$y = ab + a + c$$

# Boolean Algebraic Simplification Example

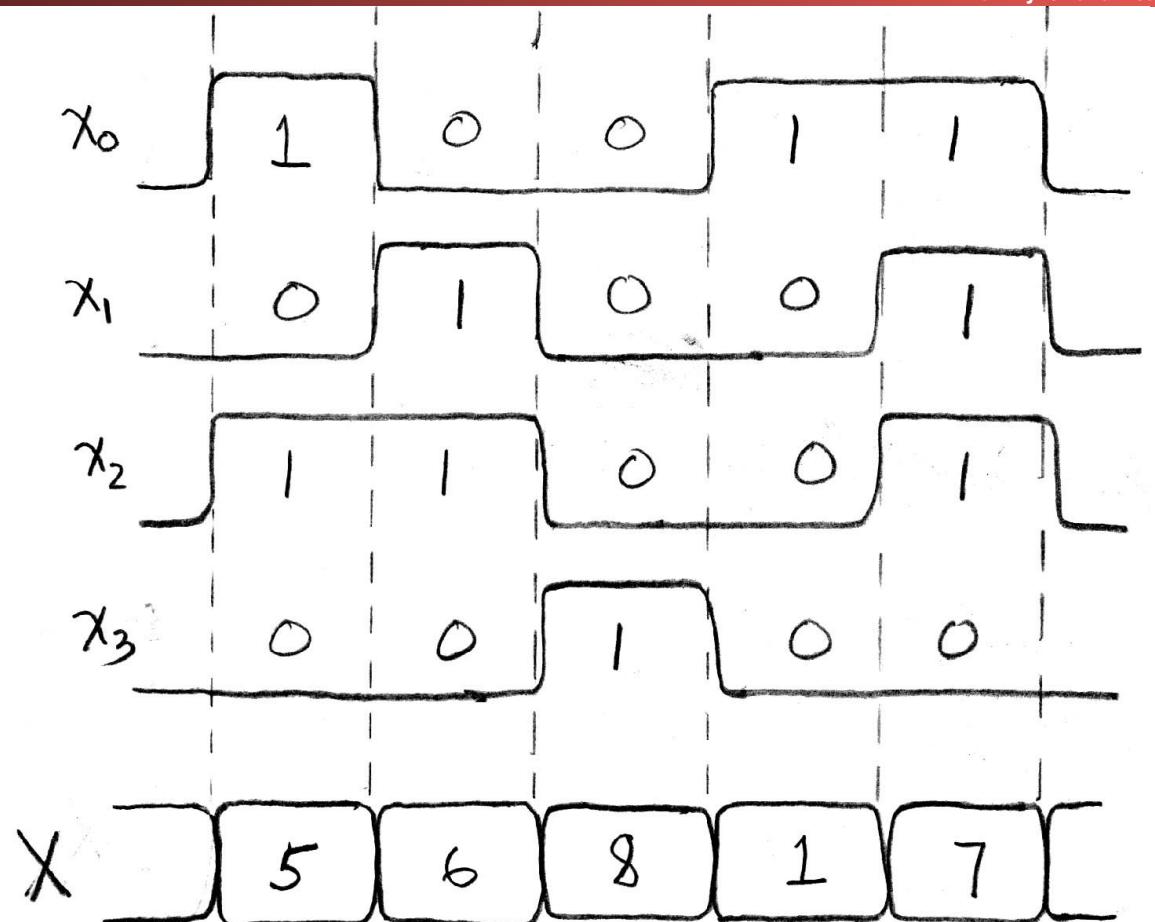
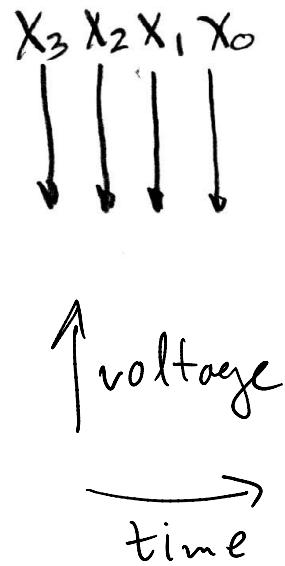
a b c	y	$= ab + a + c$	
0 0 0		$= a(b + 1) + c$	<i>distribution, identity</i>
0 0 1		$= a(1) + c$	<i>law of 1's</i>
0 1 0		$= a + c$	<i>identity</i>
0 1 1			
1 0 0			
1 0 1			
1 1 0			
1 1 1			

# Clickers/Peer Instruction

Simplify  $Z = A + BC + \bar{A}(\bar{B}\bar{C})$

- A:  $Z = 0$
- B:  $Z = A(1 + \bar{B}\bar{C})$
- C:  $Z = (A + BC)$
- D:  $Z = BC$
- E:  $Z = 1$

# Signals and Waveforms: Showing Time & Grouping



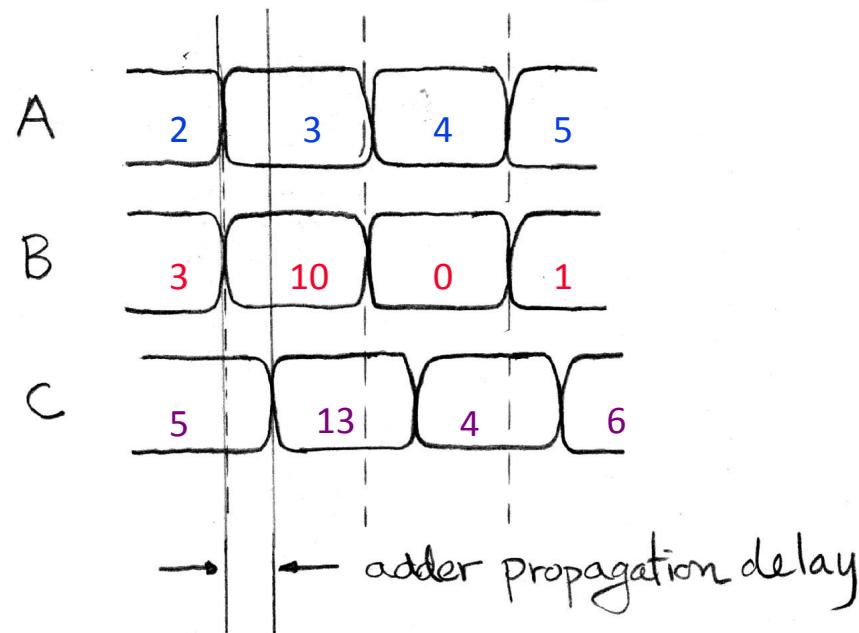
# Signals and Waveforms: Circuit Delay



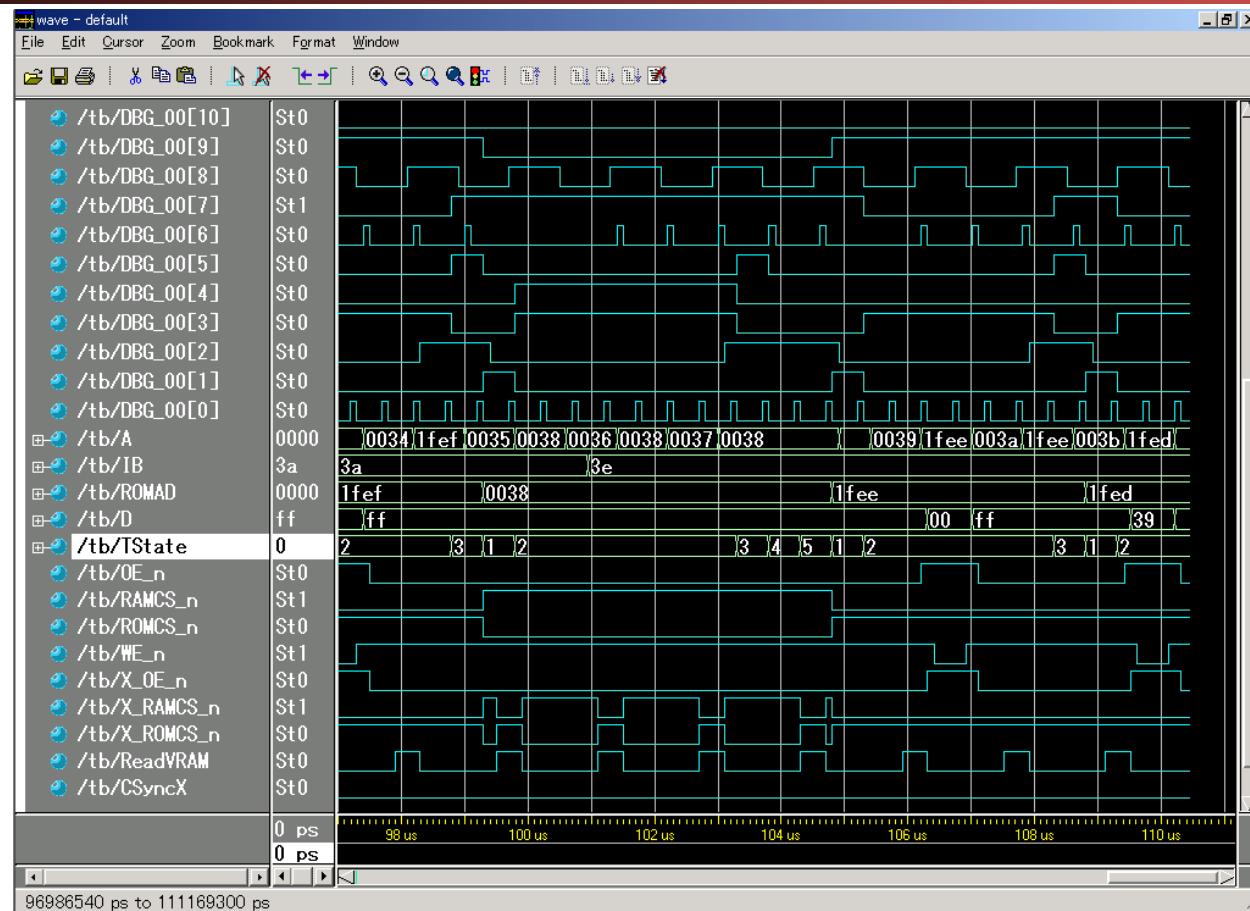
$$A = [a_3, a_2, a_1, a_0]$$

$$B = [b_3, b_2, b_1, b_0]$$

$$A \xrightarrow{4} \equiv \begin{matrix} a_0 & \longrightarrow \\ a_1 & \longrightarrow \\ a_2 & \longrightarrow \\ a_3 & \longrightarrow \end{matrix}$$



# Sample Debugging Waveform



# Type of Circuits

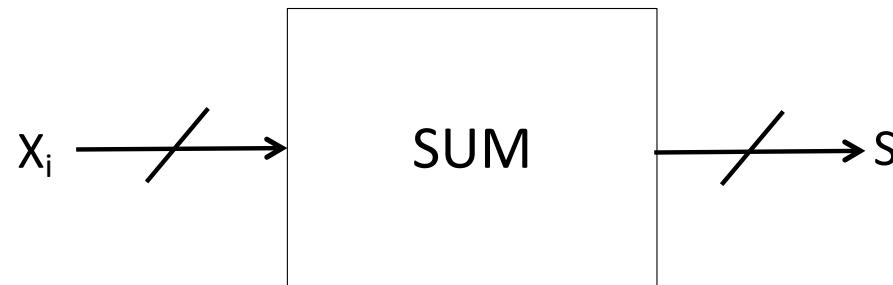
- *Synchronous Digital Systems* consist of two basic types of circuits:
  - Combinational Logic (CL) circuits
    - Output is a function of the inputs only, not the history of its execution
    - E.g., circuits to add A, B (ALUs)
  - Sequential Logic (SL)
    - Circuits that “remember” or store information
    - aka “State Elements”
    - E.g., memories and registers (Registers)

# Uses for State Elements

- Place to store values for later re-use:
  - Register files (like x1-x31 in RISC-V)
  - Memory (caches and main memory)
- *Help control flow of information between combinational logic blocks*
  - State elements hold up the movement of information at input to combinational logic blocks to allow for orderly passage

# Accumulator Example

Why do we need to control the flow of information?



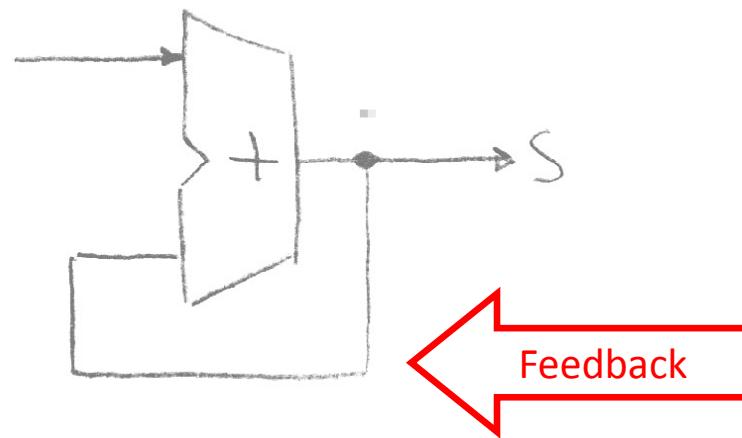
Want:

$$\begin{aligned} S &= 0; \\ \text{for } &(i=0; i < n; i++) \\ S &= S + X_i \end{aligned}$$

Assume:

- Each  $X$  value is applied in succession, one per cycle
- After  $n$  cycles the sum is present on  $S$

# First Try: Does this work?



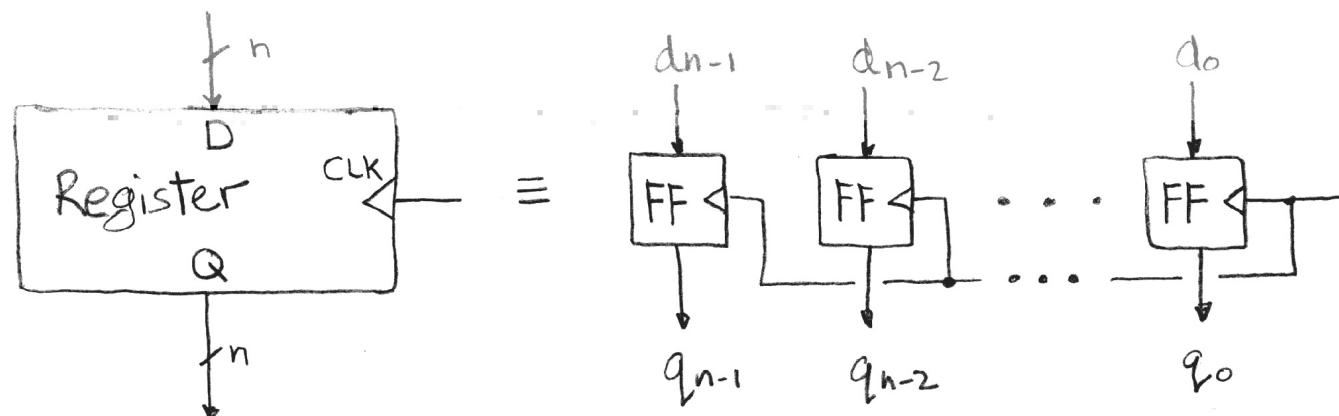
No!

Reason #1: How to control the next iteration of the 'for' loop?

Reason #2: How do we say: 'S=0'?

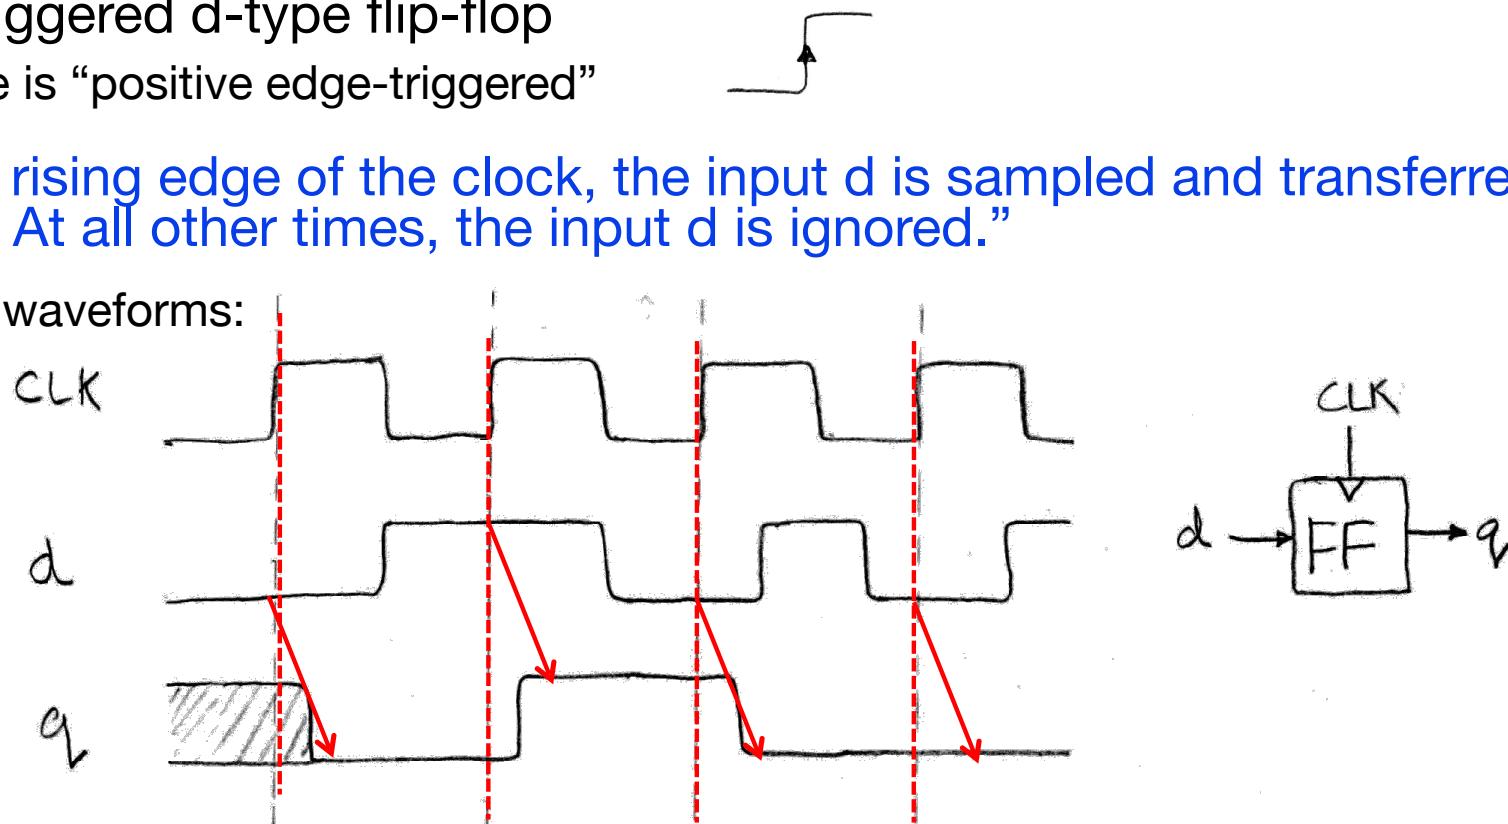
# Register Internals

- n instances of a “Flip-Flop”
- Flip-flop name because the output flips and flops between 0 and 1
- D is “data input”, Q is “data output”
- Also called “D-type Flip-Flop”

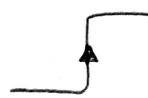


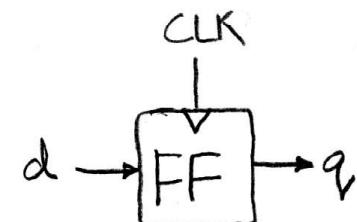
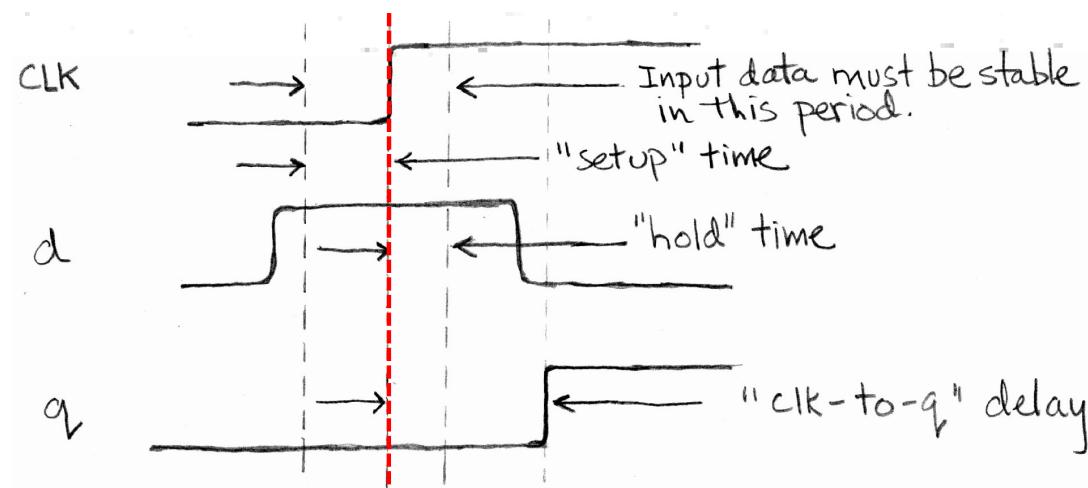
# Flip-Flop Operation

- Edge-triggered d-type flip-flop
  - This one is “positive edge-triggered”
- “On the rising edge of the clock, the input d is sampled and transferred to the output. At all other times, the input d is ignored.”
- Example waveforms:



# Flip-Flop Timing

- Edge-triggered d-type flip-flop
  - This one is “positive edge-triggered” 
- “On the rising edge of the clock, the input d is sampled and transferred to the output. At all other times, the input d is ignored.”
- Example waveforms (more detail):



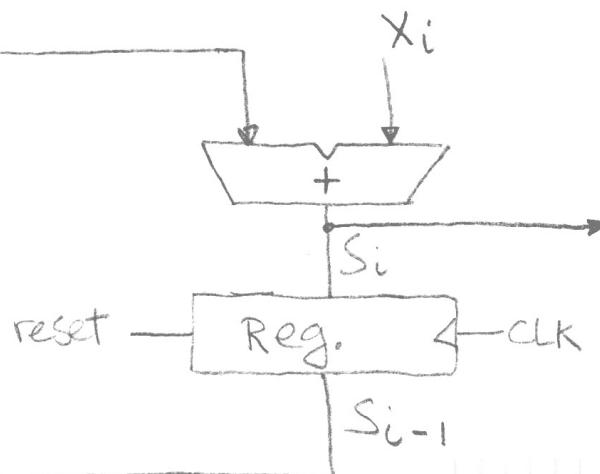
# Camera Analogy Timing Terms

- Want to take a portrait – timing right before and after taking picture
- *Set up time* – don't move since about to take picture (open camera shutter)
- *Hold time* – need to hold still after shutter opens until camera shutter closes
- *Time click to data* – time from open shutter until can see image on output (viewscreen)

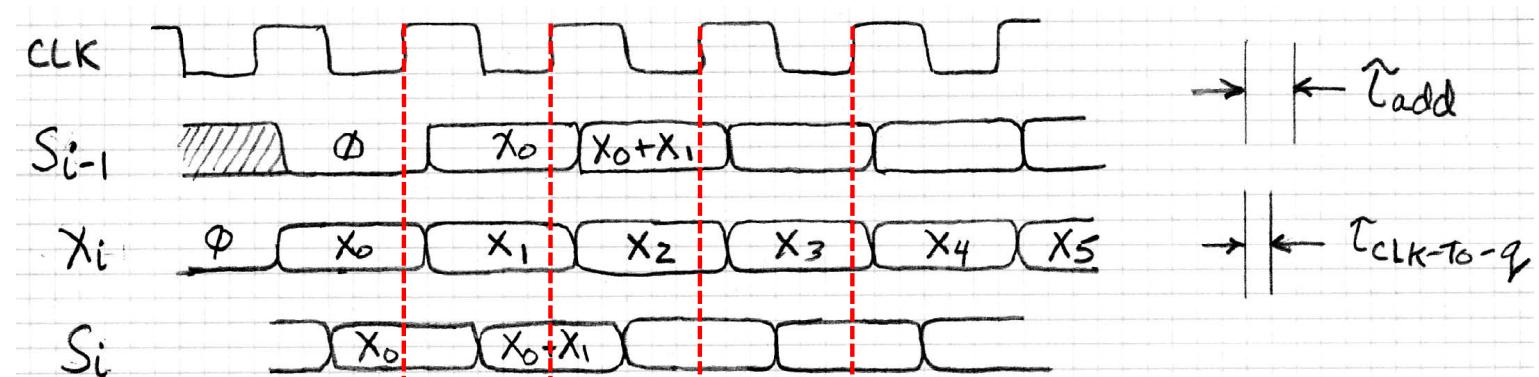
# Hardware Timing Terms

- **Setup Time:** when the input must be stable *before* the edge of the CLK
- **Hold Time:** when the input must be stable *after* the edge of the CLK
- **“CLK-to-Q” Delay:** how long it takes the output to change, measured from the edge of the CLK

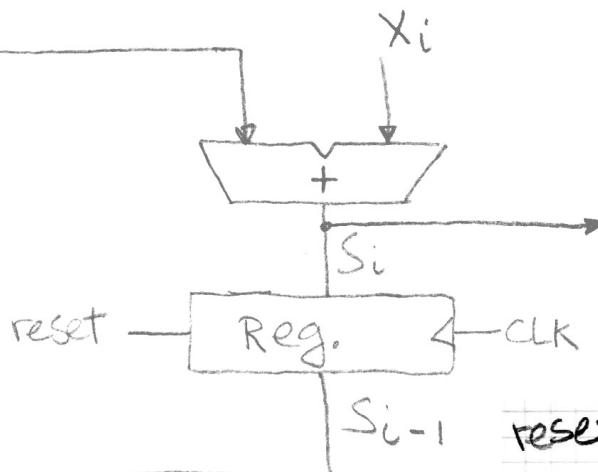
# Accumulator Timing 1/2



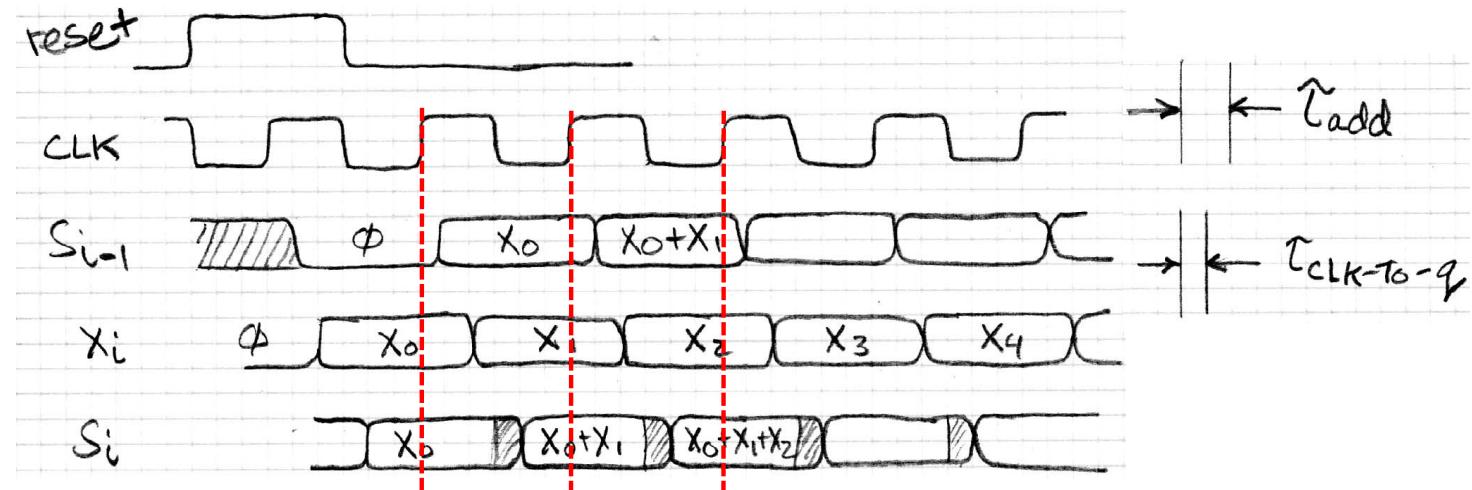
- Reset input to register is used to force it to all zeros (takes priority over D input).
- $S_{i-1}$  holds the result of the  $i^{\text{th}}-1$  iteration.
- Analyze circuit timing starting at the output of the register.



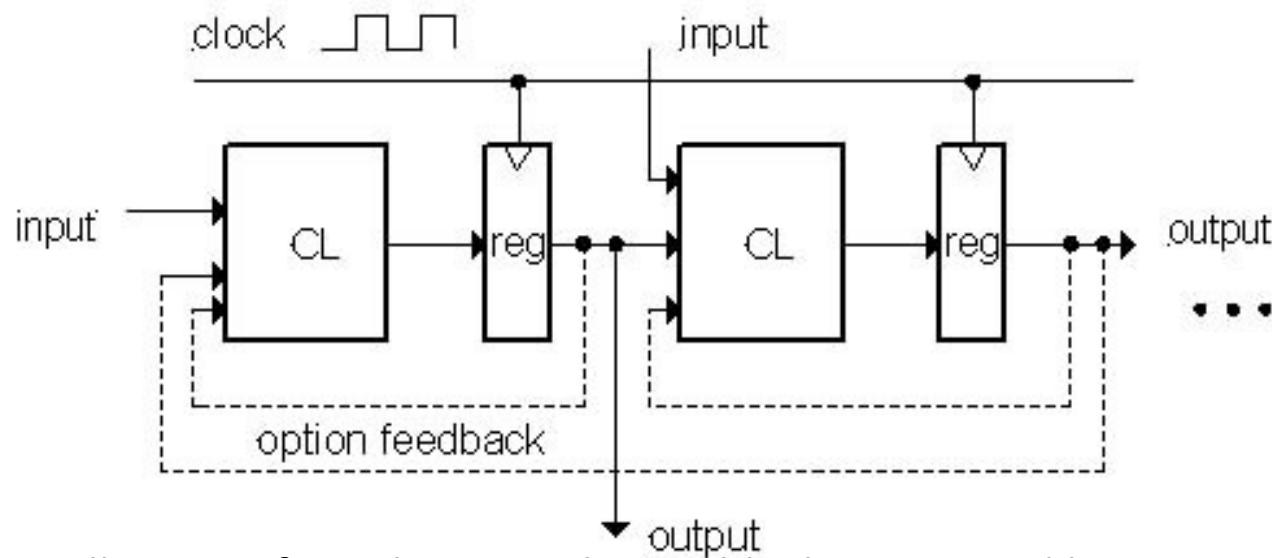
# Accumulator Timing 2/2



- reset signal shown.
- Also, in practice  $X$  might not arrive to the adder at the same time as  $S_{i-1}$
- $S_i$  temporarily is wrong, but register always captures correct value.
- In good circuits, instability never happens around rising edge of clk.



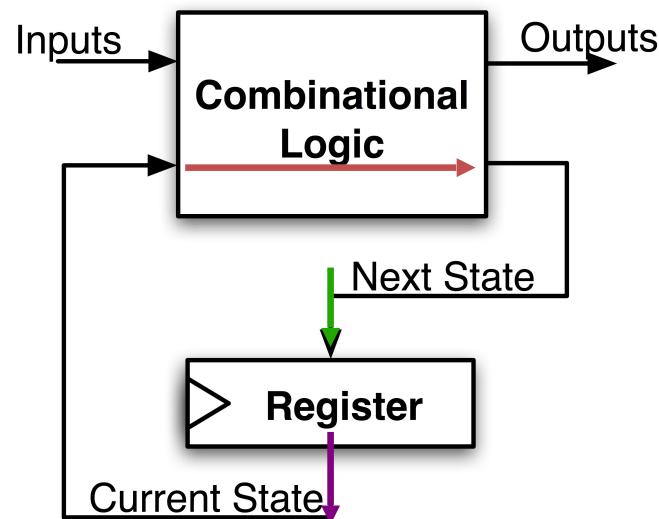
# Model for Synchronous Systems



- Collection of Combinational Logic blocks separated by registers
- Feedback is optional
- Clock signal(s) connects only to clock input of registers
- Clock (CLK): steady square wave that synchronizes the system
- Register: several bits of state that samples on rising edge of CLK (positive edge-triggered) or falling edge (negative edge-triggered)

# Maximum Clock Frequency

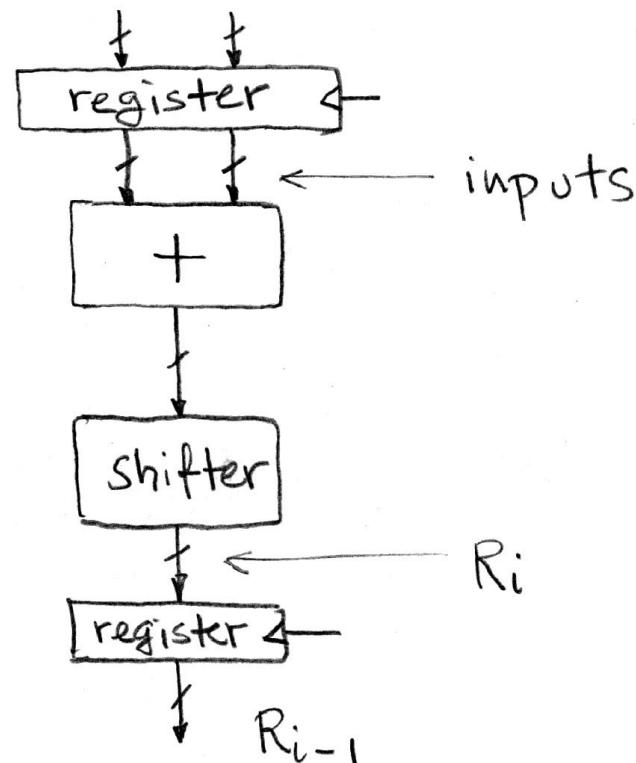
- What is the maximum frequency of this circuit?



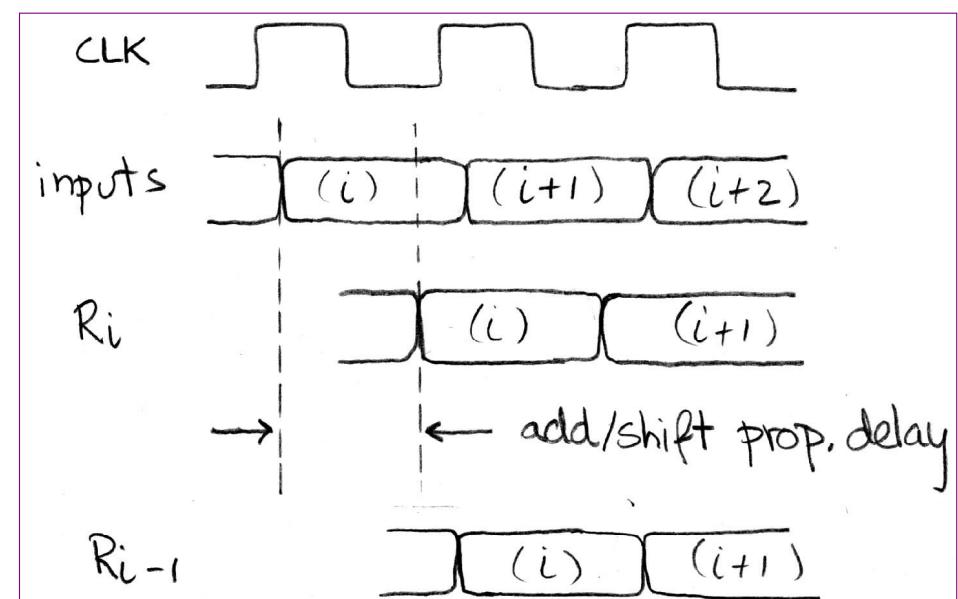
Hint:  
Frequency = 1/Period

$$\text{Period} = \text{Max Delay} = \text{CLK-to-Q Delay} + \text{CL Delay} + \text{Setup Time}$$

# Critical Paths



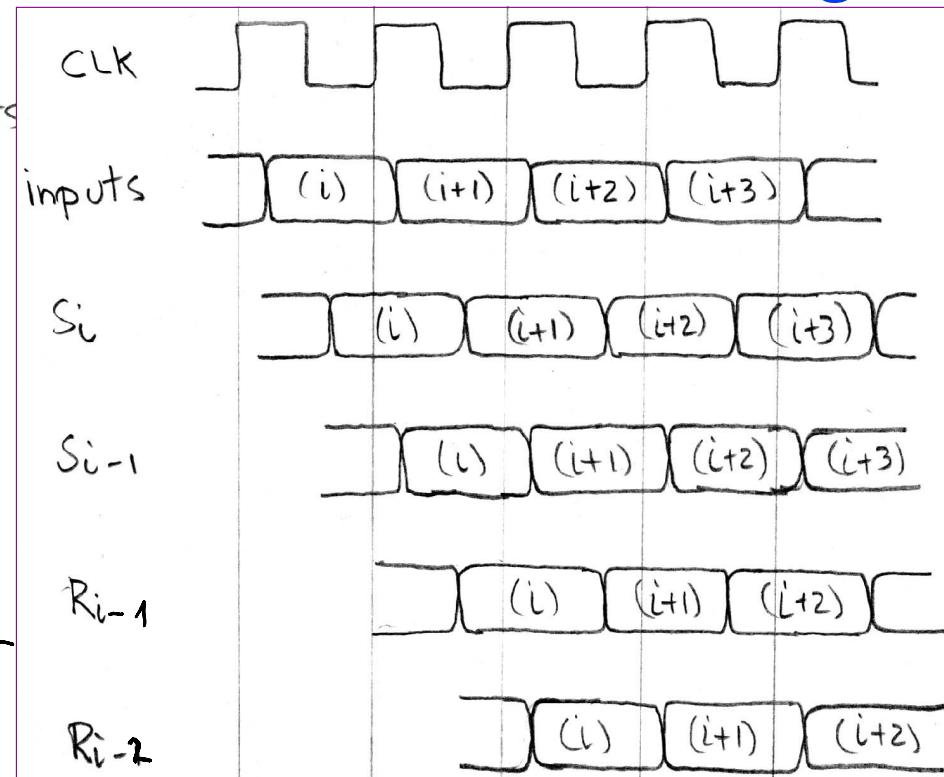
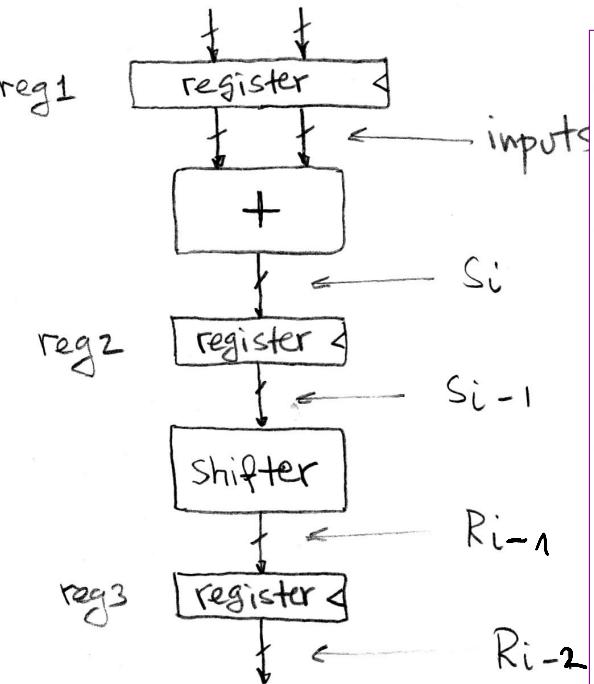
Timing...



Note: delay of 1 clock cycle from input to output.  
Clock period limited by propagation delay of adder/shifter.

# Pipelining to improve performance

Timing...



- Insertion of register allows higher clock frequency
- More outputs per second (higher bandwidth)
- But each individual result takes longer (greater latency)

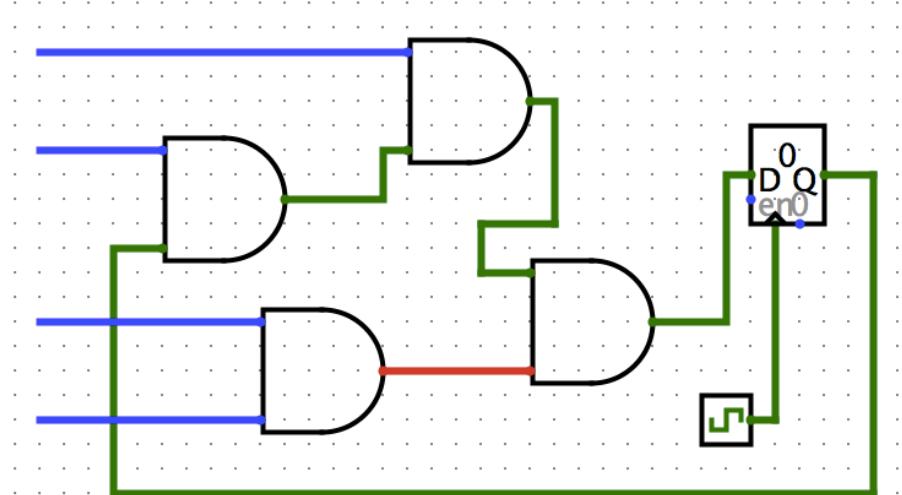
# Recap of Timing Terms

- **Clock (CLK)** - steady square wave that synchronizes system
- **Setup Time** - when the input must be stable before the rising edge of the CLK
- **Hold Time** - when the input must be stable after the rising edge of the CLK
- “**CLK-to-Q**” Delay - how long it takes the output to change, measured from the rising edge of the CLK
  
- **Flip-flop** - one bit of state that samples every rising edge of the CLK (positive edge-triggered)
- **Register** - several bits of state that samples on rising edge of CLK or on LOAD (positive edge-triggered)

# Clickers/Peer Instruction

What is maximum clock frequency? (assume all unconnected inputs come from some register)

- A: 5 GHz
- B: 200 MHz
- C: 500 MHz
- D: 1/7 GHz
- E: 1/6 GHz



Clock->Q 1ns  
Setup 1ns  
Hold 1ns  
AND delay 1ns