

More On Synchronization And Threading

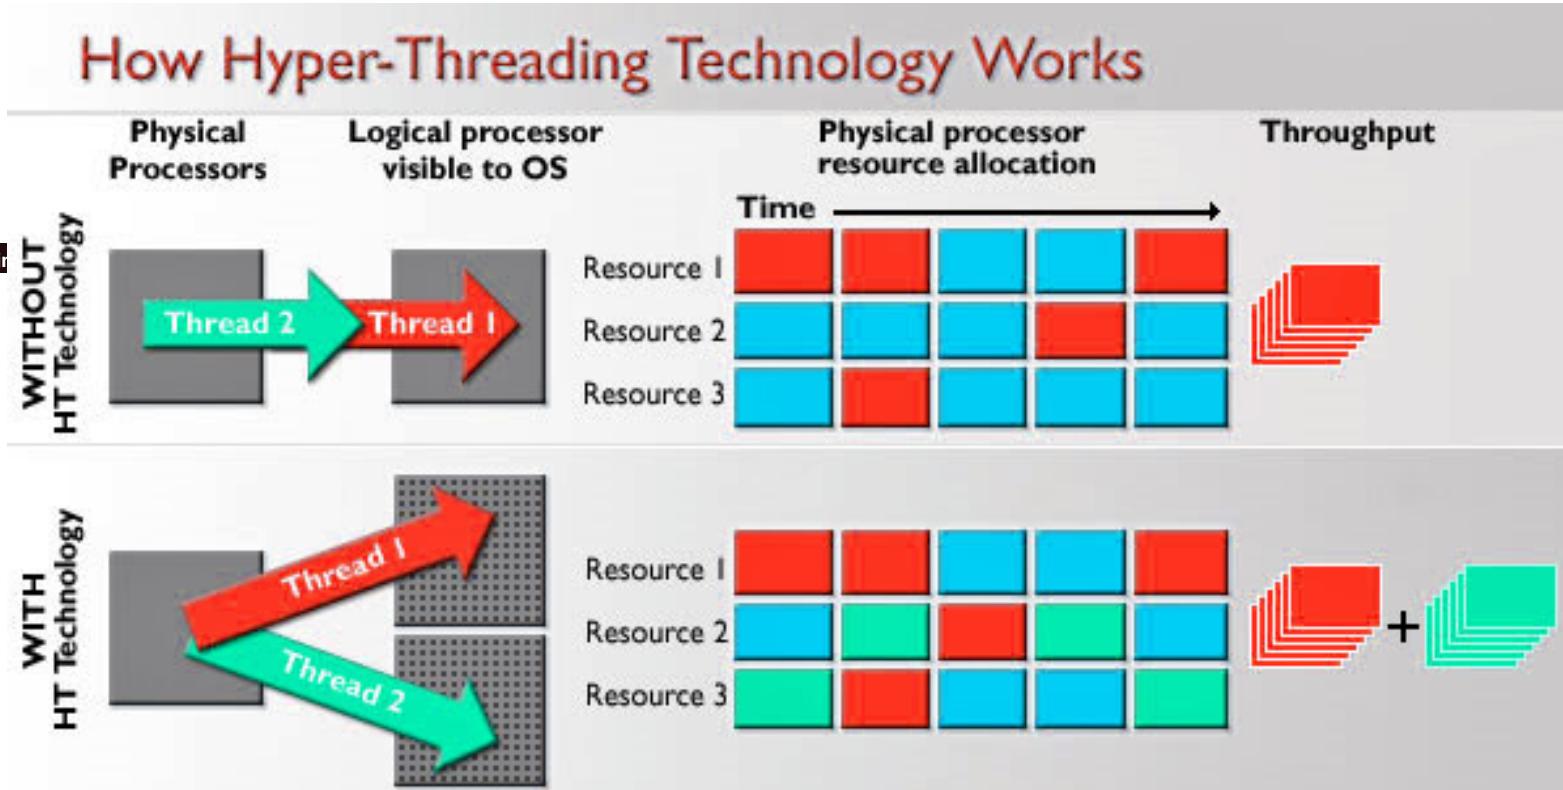
Reminder: ILP vs. TLP

- Instruction Level Parallelism
 - Multiple instructions in execution at the same time, e.g., instruction pipelining
 - Superscalar: launch more than one instruction at a time, typically from one instruction stream
 - ILP limited because of pipeline hazards

ILP vs. TLP

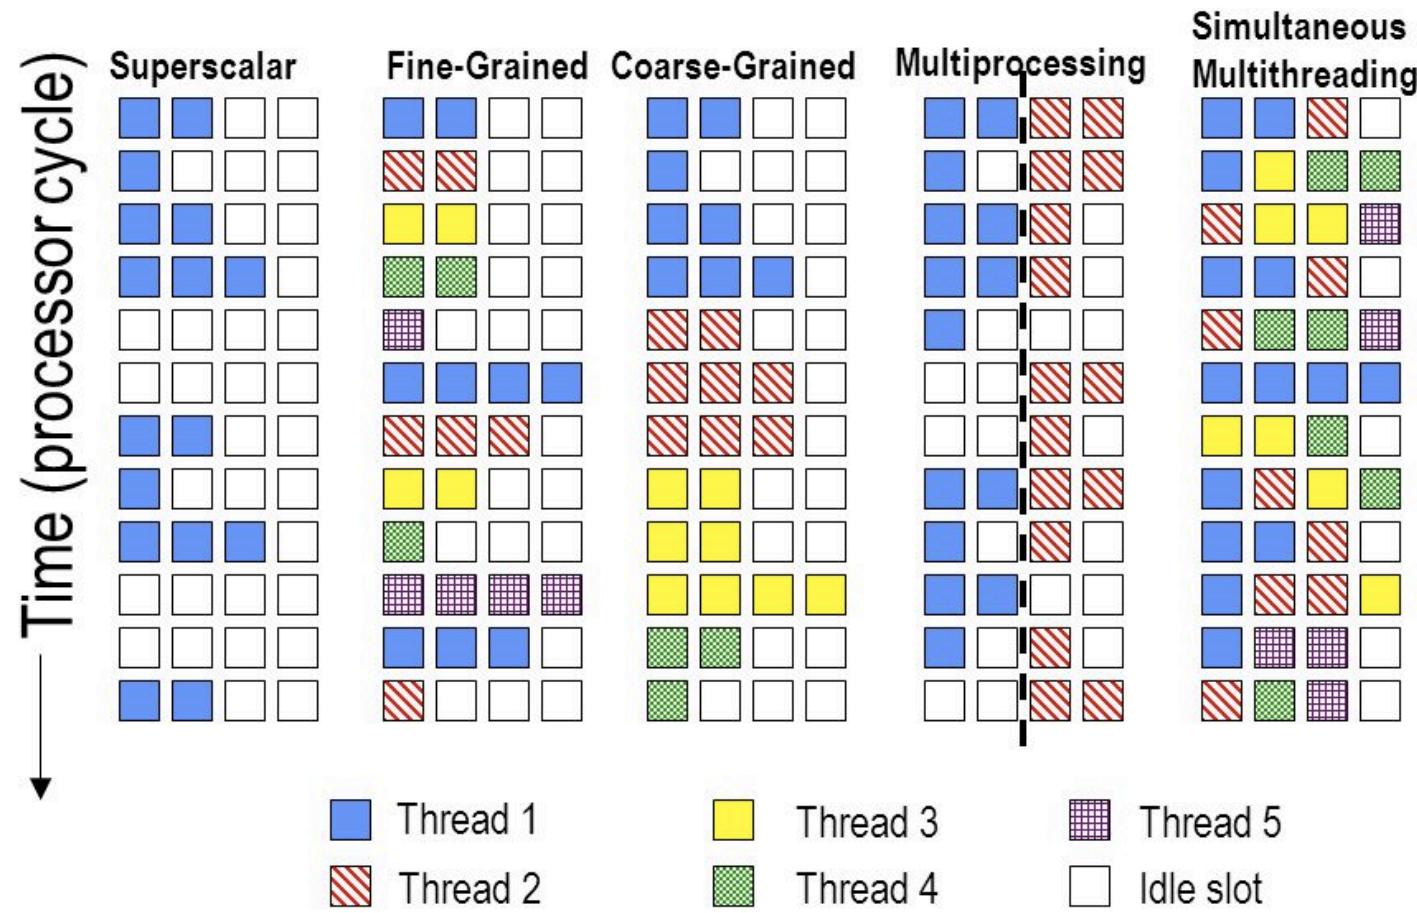
- Thread Level Parallelism
 - **Thread**: sequence of instructions, with own program counter and processor state (e.g., register file) but common memory within a process
 - **Process**: consists of at least one but could be arbitrary number of threads
 - Each process has its own memory space
 - **Multicore**:
 - Physical CPU: One thread (at a time) per CPU, in software OS switches threads typically in response to I/O events like disk read/write
 - Logical CPU: Fine-grain thread switching, in hardware, when thread blocks due to cache miss/memory access
 - Hyperthreading (aka Simultaneous Multithreading--SMT): Exploit superscalar architecture to launch instructions from different threads at the same time!

How Hyper-Threading Technology Works



- SMT (Symmetric Multithreading/Intel Hyperthreading): Logical CPUs > Physical CPUs
 - Run multiple threads at the same time per core
 - Each thread has own architectural state (PC, Registers, etc.)
 - Share resources (cache, instruction unit, execution units)
 - Improves Core CPI (clock ticks per instruction)
 - May degrade Thread CPI (Utilization/Bandwidth v. Latency)
 - See <http://dada.cs.washington.edu/smt/>

Summary: Multithreaded Categories

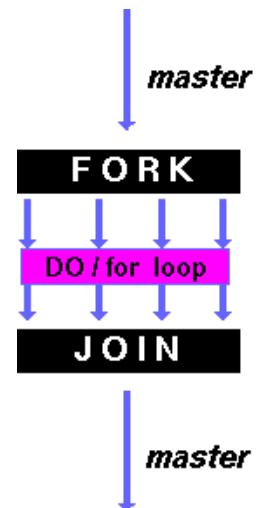


OpenMP Building Block: **for** loop rather than just the parallel block

- **for (i=0; i<max; i++) zero[i] = 0;**
- Breaks for loop into chunks, and allocate each to a separate thread
 - e.g. if max = 100 with 2 threads:
assign 0-49 to thread 0, and 50-99 to thread 1
- Must have relatively simple “shape” for an OpenMP-aware compiler to be able to parallelize it
 - Necessary for the run-time system to be able to determine how many of the loop iterations to assign to each thread:
Not a good idea to be changing the loop bounds in the loop itself
- No premature exits from the loop allowed
 - i.e. No break, return, exit, goto statements, just simple **for** and **while** loops

OpenMP Parallel for pragma

- ```
#pragma omp parallel for
for (i=0; i<max; i++) zero[i] = 0;
```
- Master thread creates additional threads, each with a separate execution context
- All variables declared outside for loop are shared by default, except for loop index which is ***implicitly*** private per thread
- Implicit “barrier” synchronization at end of for loop
- Divide index regions sequentially per thread
  - Thread 0 gets 0, 1, ..., (max/n)-1;
  - Thread 1 gets max/n, max/n+1, ..., 2\*(max/n)-1

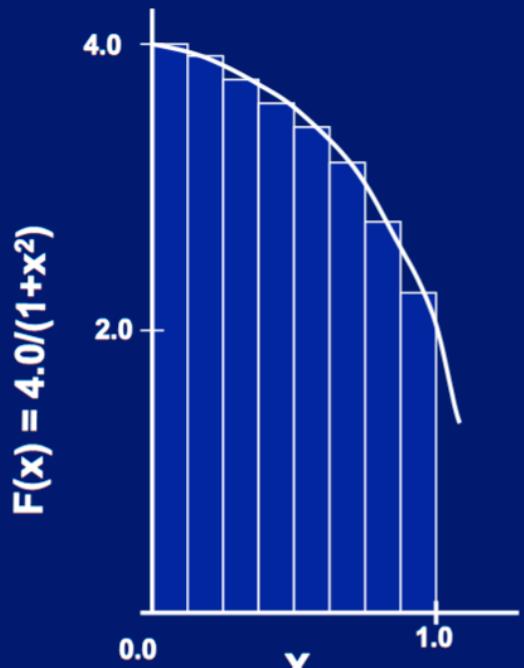


# Example 2: Computing $\pi$

## Numerical Integration

Computer Science 61C Spring 2018

Wawrynek and Weaver



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width  $\Delta x$  and height  $F(x_i)$  at the middle of interval  $i$ .

<http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>

# Working Parallel $\pi$ without a for loop

```
#include <stdio.h>
#include <omp.h>

Com Wawrynek and Weaver

void main () {
 const int NUM_THREADS = 4;
 const long num_steps = 10;
 double step = 1.0/((double)num_steps);
 double sum[NUM_THREADS];
 for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
 omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
 int id = omp_get_thread_num();
 for (int i=id; i<num_steps; i+=NUM_THREADS) {
 double x = (i+0.5) *step;
 sum[id] += 4.0*step/(1.0+x*x);
 printf("i =%3d, id =%3d\n", i, id);
 }
}
double pi = 0;
for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
printf ("pi = %6.12f\n", pi);
```

# Trial Run

```
#include <stdio.h>
#include <omp.h>

Com

void main () {
 const int NUM_THREADS = 4;
 const long num_steps = 10;
 double step = 1.0/((double)num_steps);
 double sum[NUM_THREADS];
 for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
 omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
 int id = omp_get_thread_num();
 for (int i=id; i<num_steps; i+=NUM_THREADS) {
 double x = (i+0.5) *step;
 sum[id] += 4.0*step/(1.0+x*x);
 printf("i =%3d, id =%3d\n", i, id);
 }
}
 double pi = 0;
 for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
 printf ("pi = %6.12f\n", pi);
```

|                     |
|---------------------|
| i = 1, id = 1       |
| i = 0, id = 0       |
| i = 2, id = 2       |
| i = 3, id = 3       |
| i = 5, id = 1       |
| i = 4, id = 0       |
| i = 6, id = 2       |
| i = 7, id = 3       |
| i = 9, id = 1       |
| i = 8, id = 0       |
| pi = 3.142425985001 |

# Scale up: num\_steps = $10^6$

```
#include <stdio.h>
#include <omp.h>

Comp

void main () {
 const int NUM_THREADS = 4;
 const long num_steps = 1000000;
 double step = 1.0/((double)num_steps);
 double sum[NUM_THREADS];
 for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
 omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
 int id = omp_get_thread_num();
 for (int i=id; i<num_steps; i+=NUM_THREADS) {
 double x = (i+0.5) *step;
 sum[id] += 4.0*step/(1.0+x*x);
 // printf("i =%3d, id =%3d\n", i, id);
 }
}
 double pi = 0;
 for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
 printf ("pi = %6.12f\n", pi);
```

Wawrynek and Weaver

pi = 3.141592653590

# Can We Parallelize Computing sum?

```
#include <stdio.h>
#include <omp.h>

void main () {
 const int NUM_THREADS = 1000;
 const long num_steps = 100000;
 double step = 1.0/((double)num_steps);
 double sum[NUM_THREADS];
 for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
 double pi = 0;
 omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
 int id = omp_get_thread_num();
 for (int i=id; i<num_steps; i+=NUM_THREADS) {
 double x = (i+0.5) *step;
 sum[id] += 4.0*step/(1.0+x*x);
 }
 pi += sum[id];
}
printf ("pi = %6.12f\n", pi);
```

Cor

Wawrynek and Weaver

Always looking for ways to beat Amdahl's Law ...

## Summation inside parallel section

- Insignificant speedup in this example, but ...
- **pi = 3.138450662641**
- Wrong! And value changes between runs?!
- What's going on?

# What's Going On?

```
#include <stdio.h>
#include <omp.h>

void main () {
 const int NUM_THREADS = 1000;
 const long num_steps = 100000;
 double step = 1.0/((double)num_steps);
 double sum[NUM_THREADS];
 for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
 double pi = 0;
 omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
 int id = omp_get_thread_num();
 for (int i=id; i<num_steps; i+=NUM_THREADS) {
 double x = (i+0.5) *step;
 sum[id] += 4.0*step/(1.0+x*x);
 }
 pi += sum[id]; ←
}
printf ("pi = %6.12f\n", pi);
```

Co

Wawrynek and Weaver

- Operation is really  $\mathbf{pi} = \mathbf{pi} + \mathbf{sum[id]}$
- What if >1 threads reads current (same) value of **pi**, computes the sum, stores the result back to **pi**?
- Each processor reads same intermediate value of **pi**!
- Result depends on who gets there when
  - A “race” → result is not deterministic but if we locked this we'd lose almost all speedup

# OpenMP Reduction

- ```
double avg, sum=0.0, A[MAX]; int i;
#pragma omp parallel for private ( sum )
for (i = 0; i <= MAX ; i++) {sum += A[i];}
avg = sum/MAX; // bug, we only get the master thread's sum
```
- Problem is that we really want sum over all threads!
- **Reduction:** specifies that 1 or more variables that are private to each thread are subject of reduction operation at end of parallel region:
reduction(operation:var) where
 - Operation: operator to perform on the variables (var) at the end of the parallel region
 - Var: One or more variables on which to perform scalar reduction: private than combined
- ```
double avg, sum=0.0, A[MAX]; int i;
#pragma omp for reduction(+ : sum)
for (i = 0; i <= MAX ; i++) {sum += A[i];}
avg = sum/MAX;
```

# Calculating $\pi$ Simple Version

```
#include <omp.h>
#define NUM_THREADS 4
static long num_steps = 100000; double step;

void main () {
 int i; double x, pi, sum[NUM_THREADS];
 step = 1.0/(double) num_steps;
#pragma omp parallel private (i, x)
{
 int id = omp_get_thread_num();
 for (i=id, sum[id]=0.0; i<num_steps; i=i+NUM_THREADS)
 {
 x = (i+0.5)*step;
 sum[id] += 4.0/(1.0+x*x);
 }
}
for(i=1; i<NUM_THREADS; i++)
 sum[0] += sum[i]; pi = sum[0];
printf ("pi = %6.12f\n", pi);
}
```

# Version 2: parallel for, reduction

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000;
double step;
void main ()
{ int i; double x, pi, sum = 0.0;
 step = 1.0/(double) num_steps;
#pragma omp parallel for private(x) reduction(+:sum)
 for (i=1; i<= num_steps; i++){
 x = (i-0.5)*step;
 sum = sum + 4.0/(1.0+x*x);
 }
 pi = sum;
 printf ("pi = %6.8f\n", pi);
}
```

# Reduction Options...

- Arithmetic
  - + \* -
- Comparison
  - min max
- Logical
  - & && | || ^

# OpenMP Timing

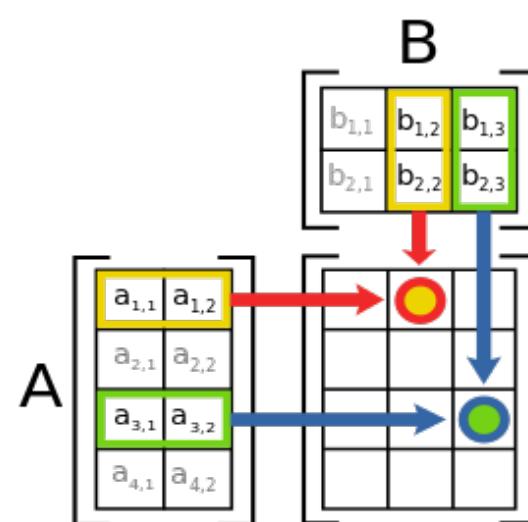
- Elapsed wall clock time:

```
double omp_get_wtime(void);
```

- Returns elapsed wall clock time in seconds
- Time is measured per thread, no guarantee can be made that two distinct threads measure the same time
- Time is measured from “some time in the past,” so subtract results of two calls to `omp_get_wtime` to get elapsed time

# Matrix Multiply in OpenMP

```
// C[M][N] = A[M][P] × B[P][N]
start_time = omp_get_wtime();
#pragma omp parallel for private(tmp, j, k)
for (i=0; i<M; i++) {
 for (j=0; j<N; j++) {
 tmp = 0.0;
 for (k=0; k<P; k++) {
 /* C(i,j) = sum(over k) A(i,k) * B(k,j) */
 tmp += A[i][k] * B[k][j];
 }
 C[i][j] = tmp; // doing this to tmp prevents
 // potential cache issues from multiple
 // writers & "false sharing"
 }
}
run_time = omp_get_wtime() - start_time;
```

B  


Outer loop spread across N threads;  
inner loops inside a single thread

# Matrix Multiply in Open MP

- More performance optimizations available:
  - Higher compiler optimization (-O2, -O3) to reduce number of instructions executed
  - Cache blocking to improve memory performance
  - Using SIMD AVX instructions to raise floating point computation rate (DLP)

# Data Races and Synchronization

- Two memory accesses form a *data race* if from different threads access same location, at least one is a write, and they occur one after another
- If there is a data race, result of program varies depending on chance (which thread first?)
- Avoid data races by synchronizing writing and reading to get *deterministic* behavior
- Synchronization done by user-level routines that rely on hardware synchronization instructions

# Reminder: Locks

- Computers use locks to control access to shared resources
  - Serves purpose of microphone in example
  - Also referred to as “semaphore”
    - Although “semaphores” have slightly different semantics
- Usually implemented with a variable
  - In most object-oriented languages its an object that you can call
  - Under the hood its usually just an integer that has atomic updates:  
0 == unlocked, 1 == locked
- Two operations:
  - lock.acquire() // blocks this thread until the lock is unlocked
  - lock.release() // Unlocks the lock
- OpenMP also has a lock
  - `#pragma omp critical`  
{  
... Only one thread is running at a time here  
}

# Hardware Synchronization (aka building atomicity)

- Hardware support required to prevent an interloper (another thread) from changing the value
  - **Atomic** read/write memory operation
    - No other access to the location allowed between the read and write
  - How best to implement in software?
    - Single instr? Atomic swap of register  $\leftrightarrow$  memory
    - Pair of instr? One for read, one for write

# Synchronization in RISC-V option one: Read/Write Pairs

- Load reserved: **lr rd, rs1**
  - Load the word pointed to by **rs** into **rd**, and add a reservation: this is **my memory!!!**
- Store conditional: **sc rd, rs1, rs2**
  - Store the value in **rs2** into the memory location pointed to by **rs1**, if the reservation is still valid and write the status in **rd**
  - Returns 0 (success) if location has not been updated by anybody else since the **lr**
  - Returns nonzero (failure) if location has been updated by anyone else, write doesn't happen

# Synchronization in RISC-V Example

- Atomic swap (to test/set lock variable)

Exchange contents of register and memory:  
 $s4 \leftrightarrow \text{Mem}(s1)$

**try:**

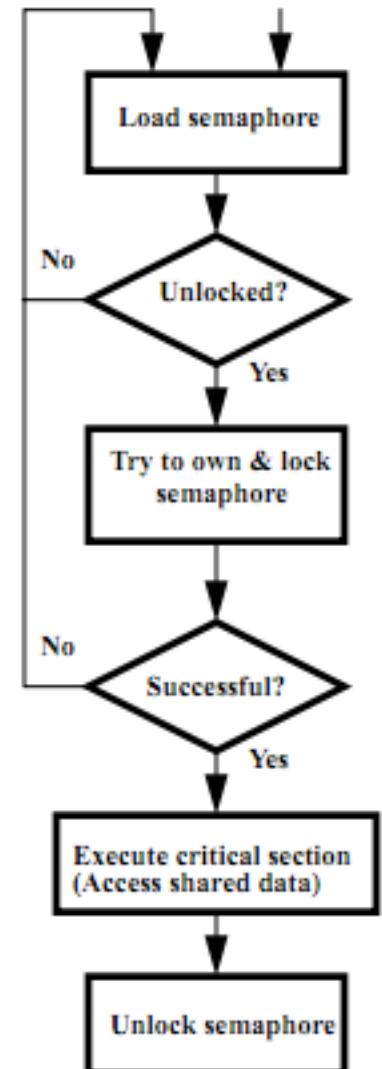
```
lr t1, s1 #load reserved
sc t0, s1, s4 #store conditional
bne t0,x0,try #loop if sc fails
add s4,x0,t1 #load value in s4
```

sc would fail if another threads executes a store here

# Test-and-Set

Computer Science 61C Spring 2018

- In a single atomic operation:
  - **Test** to see if a memory location is set (contains a 1)
  - **Set** it (to 1) if it isn't (it contained a zero when tested)
    - Otherwise indicate that the Set failed, so the program can try again
  - While accessing, no other instruction can modify the memory location, including other Test-and-Set instructions
- Useful for implementing lock operations



# Test-and-Set in RISC-V using lr/sc

- Example: RISC-V sequence for implementing a T&S at  $s_1$

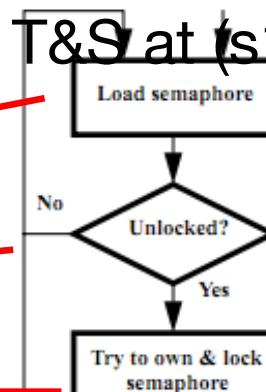
```
 li t2, 1
Try: lr t1, s1
 bne t1, x0, Try
 sc t0, s1, t2
 bne t0, x0, Try
```

Locked:

```
critical
```

Unlock:

```
sw x0, 0(s)
```



Idea is that not for programmers to use this directly, but as a tool for enabling implementation of parallel libraries

**Clickers:** Consider the following code when executed *concurrently* by two threads.

What possible values can result in  $*(s0)$ ?

```
* (s0) = 2
lw t0, 0(s0)
addi t0, t0, 1
add t0, t0, t0
sw t0, 0(s0)
```

- A: 6
- B: 14
- C: 6 or 14
- D: 6 or 14 or 30

# Why else use **lr/sc**?

- It actually allows some lock-free consistent structures
  - So although RISC-V has two **separate** atomic operation options, both have their uses
- EG...
  - **fail:**    **lr**    **t0**, **s0**,  
             **addi** **t0**, **t0**, 1  
             **add**   **t0**, **t0**, **t0**  
             **sc**    **t0**, **s0**, **t0**  
             **bne**   **t0**, **s0**, **fail**
- Allows the clicker question to be consistent without having a full lock, instead just 1 additional instruction to check for success on store conditional
  - Additionally if the code in between is short and restricted in instructions, this is **guaranteed** to work always

# RISC-V Alternative: Atomic Memory Operations

- Three instruction rtype instructions
  - Swap, and, add, or, xor, max, min
- Take the value ***pointed to*** by rs1
  - Load it into rd
  - Apply the operation to that value with rs2
  - store the result back to where rs1 is pointed to
- Note, annoyingly obscure memory semantics

# More On These Alternatives...

- Atomic swap is easy
  - It allows locks (s0 is a pointer to the lock, if that contains 0 its unlocked, nonzero if locked):

```
li t0 1
loop: amoswap t0 s0 t0
 bne t0 x0 loop
```
- Atomic associative operations are slightly different
  - Goal is to enable merging elements:  
Represents a common programming paradigm of the reduction in e.g. OpenMP
- And this is why RISC-V has **two** sets of atomic operations:
  - Atomic register+memory enables common reductions in single instructions, and is guaranteed to **always** work
  - **lrcsc** enables multiple-instructions for some level of lock-free manipulation

# Deadlock

- Deadlock: a system state in which no progress is possible because everything is locked waiting for something else
- Dining ~~Philosopher's~~ Lawyers Problem:
  - Pontificate until the left fork is available; when it is, pick it up
  - Pontificate until the right fork is available; when it is, pick it up
  - When both forks are held, eat for a fixed amount of time
  - Then, put the right fork down
  - Then, put the left fork down
  - Repeat from the beginning
- Solution?

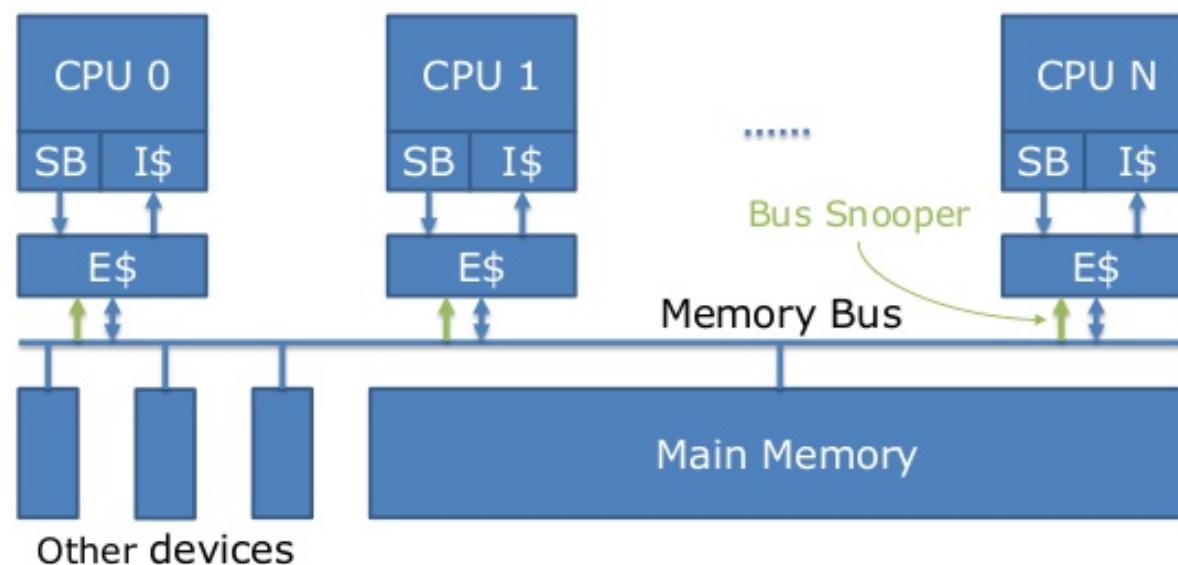


# Limiting Parallelism

- Locks act to inhibit parallelism
  - Innately sequential regions -> Amdahl's law problem...
  - Python is "threaded" but not really:  
A global interpreter lock means other threads can be waiting on I/O, but (mostly) only one compute thread at a time
- Can try to limit the locks
  - Rather than locking everything have a different lock for each region...
    - But be careful, then its much easier to get into deadlock situations
- Or try to eliminate locks altogether
  - Thus, e.g. why Go's parallelism is not focused around locks, and RISC-V **1r/sc** instructions

# (Chip) Multicore Multiprocessor

- SMP: (Shared Memory) Symmetric Multiprocessor
  - Two or more identical CPUs/Cores
  - Single shared coherent memory



# Multiprocessor Key Questions

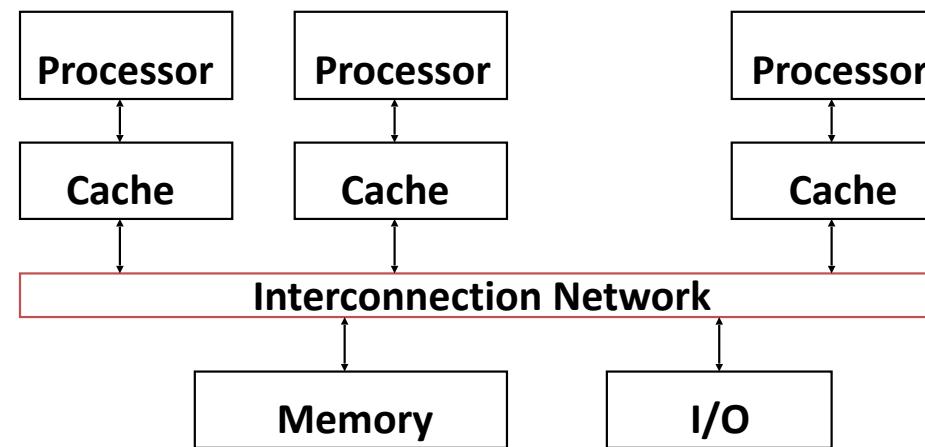
- Q1 – How do they share data?
- Q2 – How do they coordinate?
- Q3 – How many processors can be supported?

# Shared Memory Multiprocessor (SMP)

- Q1 – Single address space shared by all processors/cores
- Q2 – Processors coordinate/communicate through shared variables in memory (via loads and stores)
  - Use of shared data must be coordinated via synchronization primitives (locks) that allow access to data to only one processor at a time
- Effectively all multicore computers today are SMP
- Q3 - Depends on the workload!
  - Most systems go with "Best available single core within constraints, duplicate that"

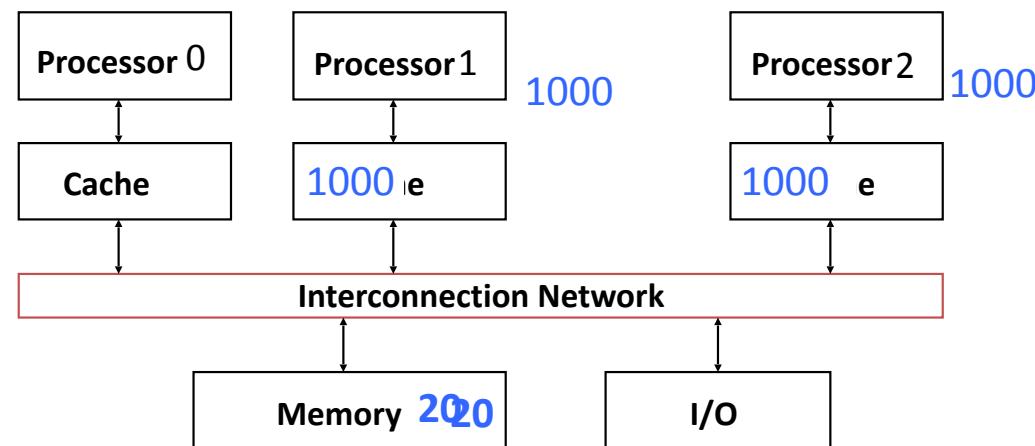
# Multiprocessor Caches

- Memory is a performance bottleneck even with one processor
- Use caches to reduce bandwidth demands on main memory
- Each core has a local private cache holding data it has accessed recently
- Only cache misses have to access the shared common memory



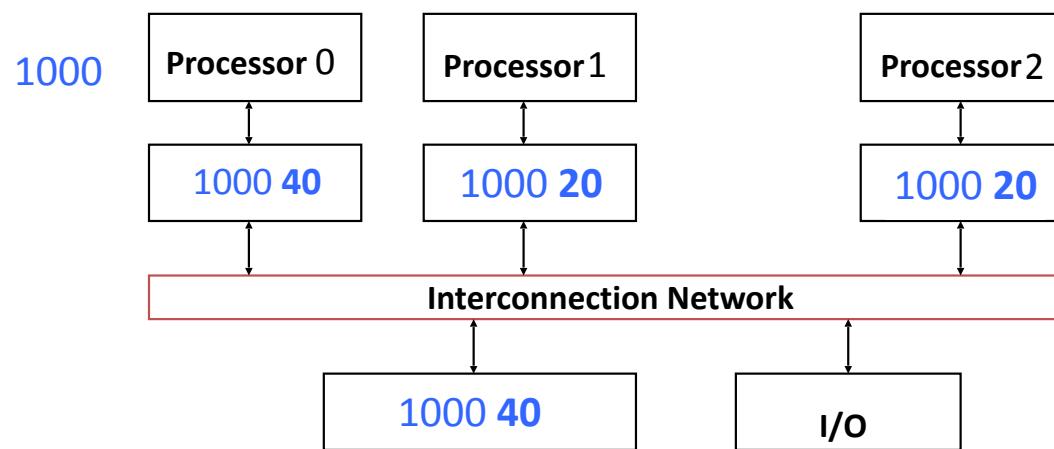
# Shared Memory and Caches

- What if?
  - Processors 1 and 2 read Memory[1000] (value 20)



# Shared Memory and Caches

- Now:
  - Processor 0 writes Memory[1000] with 40



Problem?

# Keeping Multiple Caches Coherent

- Architect's job: shared memory  
=> keep cache values *coherent*
- Idea: When any processor has cache miss or writes, notify other processors via interconnection network
  - If only reading, many processors can have copies
  - If a processor writes, invalidate any other copies
- Write transactions from one processor, other caches “snoop” the common interconnect checking for tags they hold
  - Invalidate any copies of same address modified in other cache

# How Does HW Keep \$ Coherent?

- We already saw how to do this with just ***valid*** and ***dirty***, but we can also think of it this way...
- Each cache tracks state of each ***block*** in cache:
  1. ***Shared***: up-to-date data, other caches may have a copy (Valid Bit only is set)
  2. ***Modified***: up-to-date data, changed (dirty), no other cache has a copy, OK to write, memory out-of-date (Dirty bit is also set)

# Two Optional Performance Optimizations of Cache Coherency via New States

- Each cache tracks state of each *block* in cache:
3. *Exclusive*: up-to-date data, no other cache has a copy, OK to write, memory up-to-date
    - If any other cache reads this line the state becomes *shared*
      - Can provide the line if I'm faster than main memory
      - If this writes to this line, state becomes *modified* but I don't need to broadcast this when I do the write
  4. *Owner*: up-to-date data, other caches may have a copy (they must be in Shared state)
    - But this cache now supplies data on read instead of going to memory:  
Saves the need for a write back when somebody else reads the page

# Name of This Common Cache Coherency Protocol: MOESI

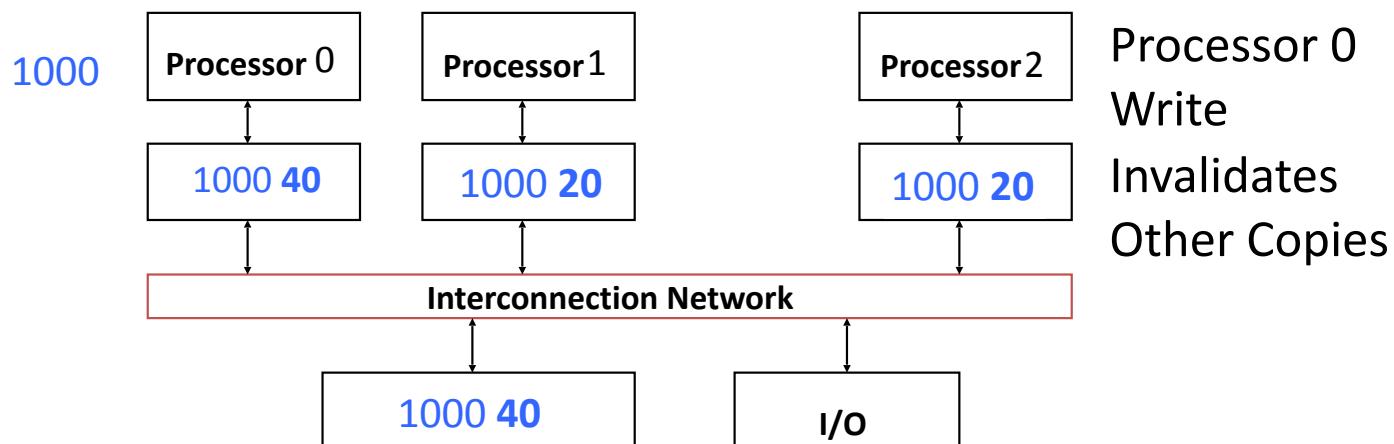
- Memory access to cache is either  
Modified (in cache)  
Owned (in cache)  
Exclusive (in cache)  
Shared (in cache)  
Invalid (not in cache)



Snooping/Snoopy Protocols  
e.g., the Berkeley Ownership Protocol  
See [http://en.wikipedia.org/wiki/Cache\\_coherence](http://en.wikipedia.org/wiki/Cache_coherence)  
Berkeley Protocol is a wikipedia stub!

# Shared Memory and Caches

- Example, now with cache coherence
  - Processors 1 and 2 read Memory[1000]
  - Processor 0 writes Memory[1000] with 40



# But An Alternate // Programming Paradigm... Communicating Sequential Processes

- OpenMP has **very restrictive** parallelism
  - Really only good for parallelizing loops with an optional reduction step
    - And Amdahl's law therefore quickly rears its ugly head
- Raw threads is **very easy** to get wrong
  - Deadlock situations
- Enter CSP:
  - A way for different threads to efficiently communicate
  - A good CSP language: Go (golang)
  - (Only going to cover this if there is time)

# What is Go

- Language created at Google starting in 2007
  - Primarily by a bunch of old Unix hands: Robert Griesemer, Rob Pike, and Ken Thompson
  - 1.0 released in March 2012
- Language continues to evolve, but a commitment to backwards compatibility (so far)
  - A correct program written today will still work tomorrow
    - I'm looking at you, **python 3....**
- Mostly C-ish but...
  - Strong typing, no pointer arithmetic, lambdas, interfaces, garbage collection and...
  - Strong emphasis on concurrent computation

# Good Go Resources

- The Go website:
  - <https://golang.org/>
- Especially useful: Effective Go:
  - A cheatsheet of programming idioms. Several example from this lecture stolen from there
  - [https://golang.org/doc/effective\\_go.html](https://golang.org/doc/effective_go.html)
- When searching Google, ask for **golang**, not go
  - The language may be Go, but golang refers to the language too
- If I'm starting new code and I need to care about performance, scalability, or maintainability, I use Go

# So Think of Go as:

- C's general structure & concepts
  - But with implied ;s and a garbage collector
- A better typing system with interfaces, slices, and maps
  - No class inheritance, however
- Much more symmetric functions
  - Can return multiple values
- Scheme-like lexical scope
  - Lambdas and interior function declarations
- Communicating Synchronous Processes (CSP) concurrency
  - Multiple things at once in the same shared memory space:  
quite suitable for MIMD

# Coroutines, err, goroutine

- Conceptually, a goroutine is just a thread...
  - `go fubar() // Executes fubar as a new coroutine`
- But in practice it is designed to be ***much*** lighter weight
  - Threads (e.g. in OpenMP) relatively expensive to create:  
Operating System involvement is never cheap
- Go's runtime instead pre-creates a series of threads
  - And then schedules the active coroutines itself to the available threads
- Result is goroutines are ***cheap***
  - It is only slightly more expensive than a plain function call:  
new goroutine just has a small independent stack  
context switching between goroutines is very cheap

# Channels

- Channels are the primary synchronization mechanism
  - You have locks but why bother?
  - A typed and (optionally) buffered communication channel
  - `c := make(chan int)`  
`e := make(chan fubar, 100)`
- Writing data to a channel:
  - `c <- 32 // Writing blocks if unbuffered or full`
- Reading from a channel:
  - `var f = <- e // Reading blocks if no data`

# Channels as Synchronization Barriers

- Any writes in the code before writing to the channel complete first
  - `globalA = ...`
  - `d <- 1 // The write to globalA must complete just before this`
- Any reads in the code after reading from the channel do not start until channel-read takes place
  - `<- d`
  - `fubar(globalA) // Won't read globalA before channel read`
- Otherwise, compiler can reorder **however it wants** as long as sequential semantics are preserved for the sequential function
  - Go may have removed a lot of ways to shoot yourself in the foot...  
but unless you use channels you will easily blow it off with race conditions

# Without Synchronization Barriers, The Compiler Can Go To Town

Computer Science 61C Spring 2018

Wawrynek and Weaver

- This doesn't work!
  - Compiler can reorder the writes between x and done safely
  - Similar variants also possible
- This is one of the two biggest pitfalls of Go:
  - Unless you explicitly synchronize, multiple processes can write in "weird" ways
  - The other is the abysmal error/exception handling mechanism

```
var x : string
var done : bool
func foo() {
 ...
 x = "something"
 done = true
}

func bar() {
 go foo()
 for !done { ... }
 y := x
}
```

# Using goroutines

- For things which may block or wait
  - EG, on input/output, waiting for stuff to happen, etc
  - Just create as many as you want!
    - Its cheap so why not: let the scheduler do useful work when another one is waiting
    - EG, if building a webserver, you launch a goroutine to handle each communication stream
- For performance tasks
  - Only create as many as there are CPU cores
    - Otherwise you are wasting resources
  - But it should be more efficient than OpenMP:
    - Thread creation is significantly more overhead than go

# Select

- Select allows you to wait on multiple channels

- ```
select {
    case c <- x:
        x, y = y, x+y
    case d := <- e:
        fmt.Println("Got %v", d)
    default:
        time.Sleep(50 * time.Millisecond)
}
```

- If can write or read to a channel, do so and **then** execute the associated case
 - If multiple cases are valid, chose one *at random*
- If nothing is available, execute default (if any)
 - If no default, just block until you can write or read
 - Default enables non-blocking read & write

Lots of other features for code correctness

- Compiler is, umm, persnickety
 - It is an error to declare but not use a variable or include but not use a package
- Designed to turn comments into documents
 - Including examples
- Libraries for building example and test routines
- Built in package management
- “Single workspace” notion
 - Use common modules to prevent code drift