

# CS162

## Operating Systems and Systems Programming

### Lecture 7

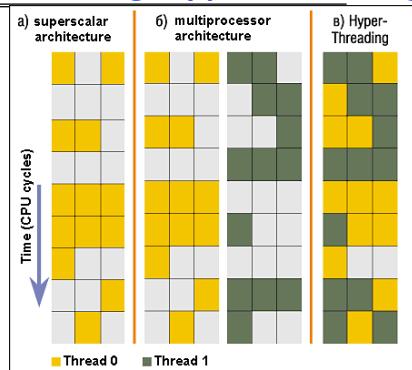
#### Concurrency (Continued), Synchronization

February 7<sup>th</sup>, 2018

Profs. Anthony D. Joseph and Jonathan Ragan-Kelley  
<http://cs162.eecs.Berkeley.edu>

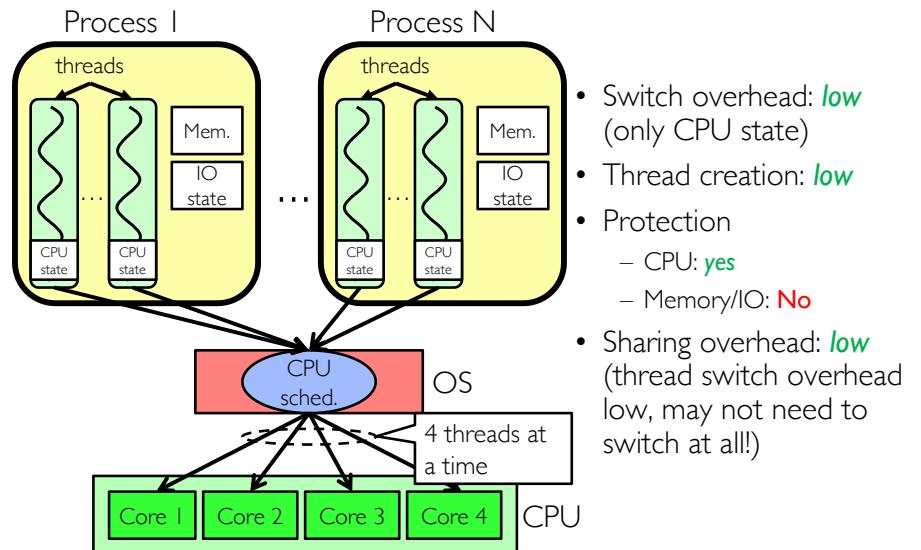
#### Recall: Simultaneous MultiThreading/Hyperthreading

- Hardware technique
  - Superscalar processors can execute multiple instructions that are independent
  - Hyperthreading duplicates register state to make a second “thread,” allowing more instructions to run
- Can schedule each thread as if were separate CPU
  - But, sub-linear speedup!
- Original called “Simultaneous Multithreading”
  - <http://www.cs.washington.edu/research/smt/index.html>
  - Intel, SPARC, Power (IBM)
  - A virtual core on AWS’ EC2 is basically a hyperthread

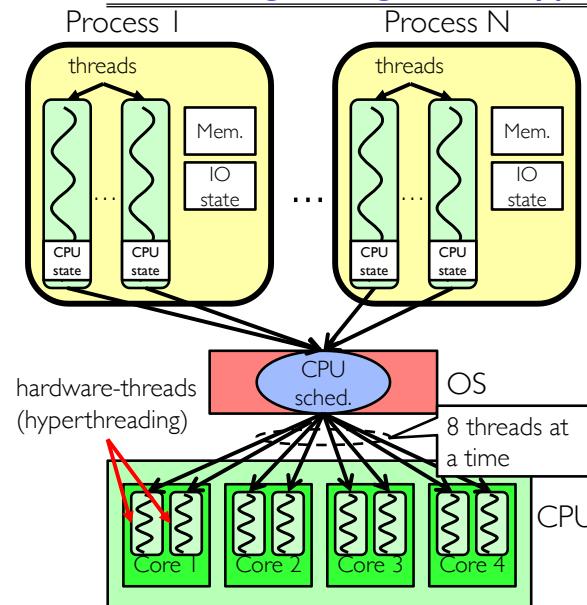


Colored blocks show instructions executed

#### Recall: Putting it Together: Multi-Cores



#### Putting it Together: Hyper-Threading



- Switch overhead between hardware-threads: **very-low** (done in hardware)
- Contention for ALUs/FPUs may hurt performance

## Classification

# threads Per AS:	# of addr spaces:	One	Many
One	MS/DOS, early Macintosh	Traditional UNIX	
Many	Embedded systems (Geoworks, VxWorks, JavaOS, etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 10 Win NT to XP, Solaris, HP-UX, OS X	

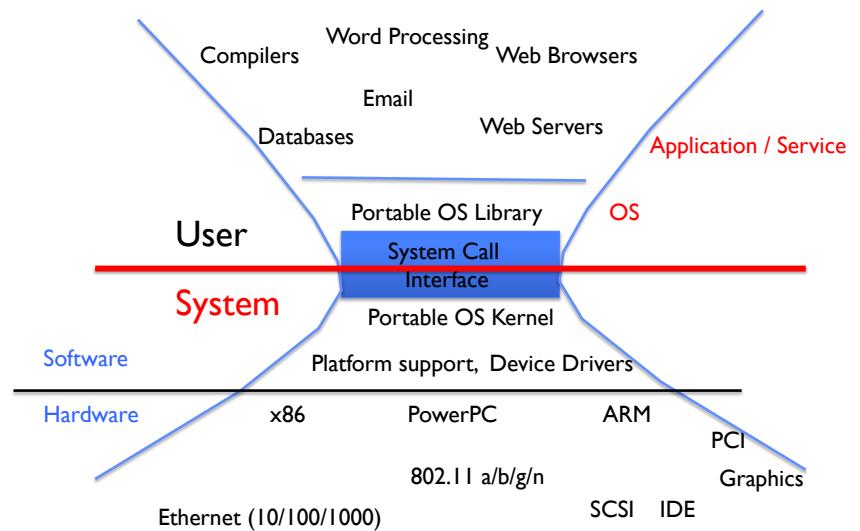
- Most operating systems have either
  - One or many address spaces
  - One or many threads per address space

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.5

## Operating System as Design

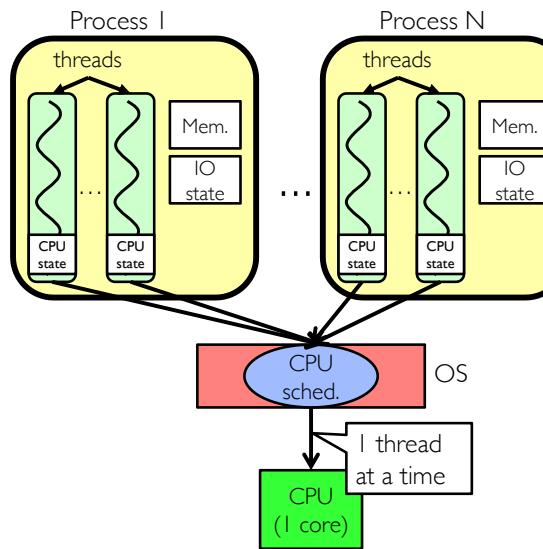


2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.6

## Conceptual Framework



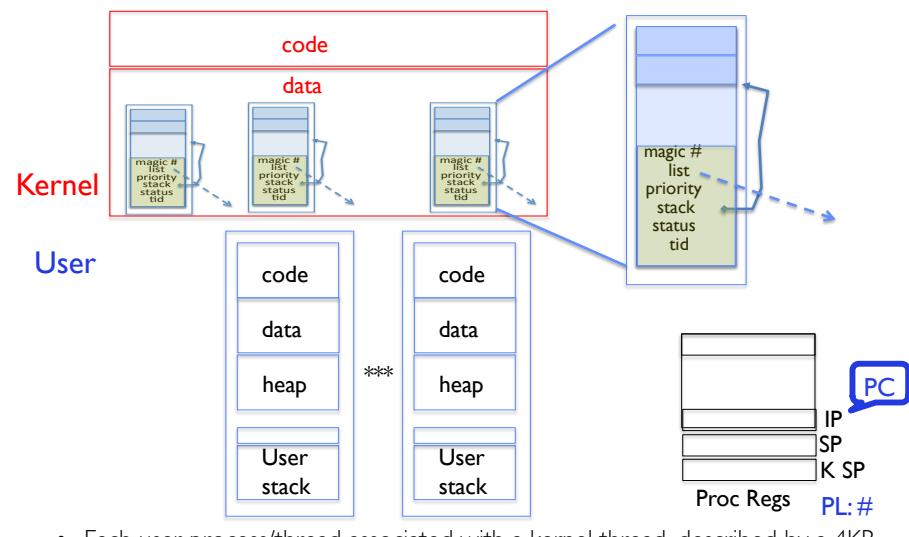
- Physical Addresses Shared
  - So: Processes and Address Translation
- Single CPU Must Be Shared
  - So: Threads
- Processes Aren't Trusted
  - So: Kernel/Userspace Split
- Threads Might Not Cooperate
  - So: Use timer interrupts to context switch ("preemption")

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.7

## Recall: MT Kernel IT Process ala Pintos/x86



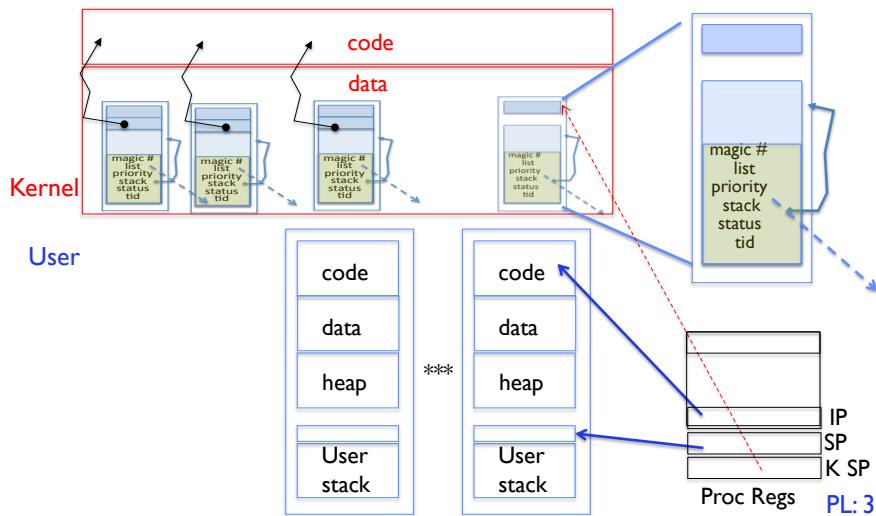
- Each user process/thread associated with a kernel thread, described by a 4KB page object containing TCB and kernel stack for the kernel thread

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.8

## In User thread, w/ Kernel thread waiting



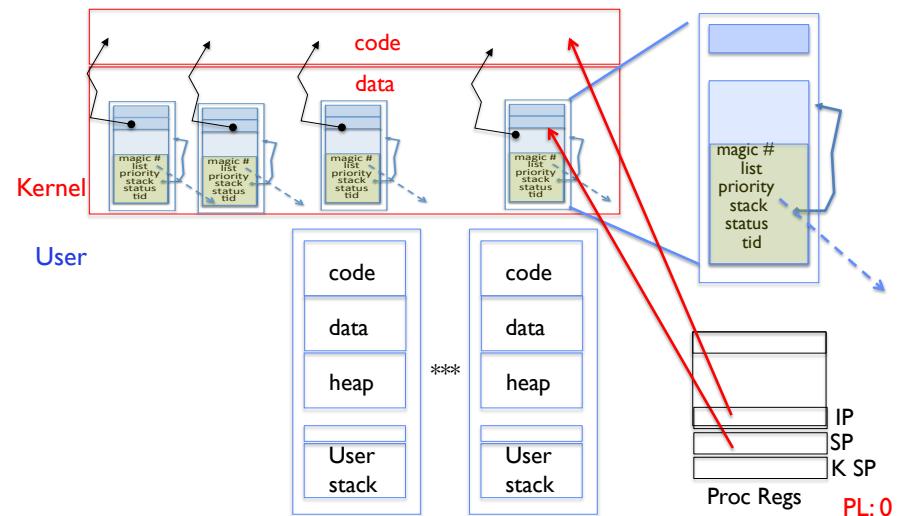
- During user thread execution, associated kernel thread is “standing by”
- x86 CPU holds interrupt SP in register

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.9

## User → Kernel



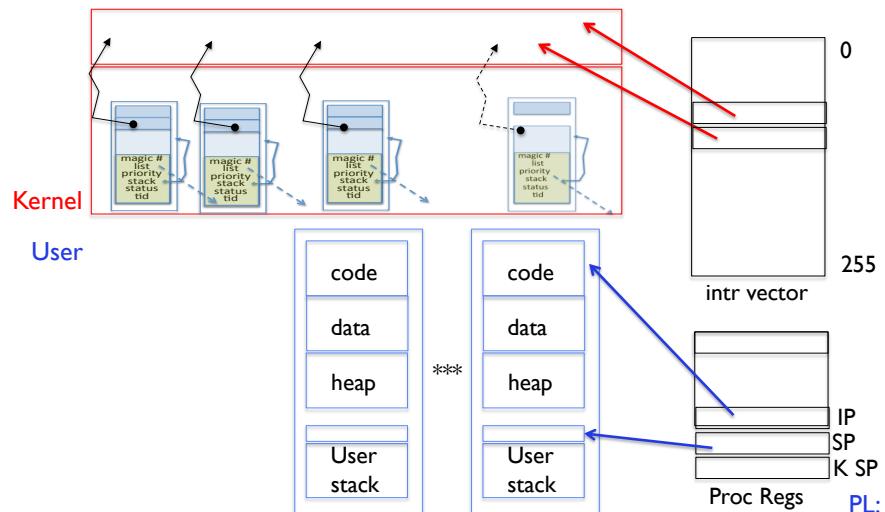
- Mechanism to resume k-thread goes through interrupt vector

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.10

## User → Kernel via interrupt vector



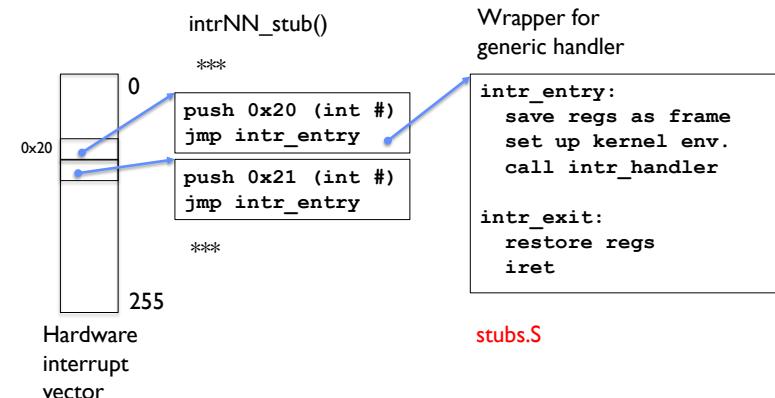
- Interrupt transfers control through the Interrupt Vector (IDT in x86)
- iret restores user stack and PL

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.11

## Pintos Interrupt Processing

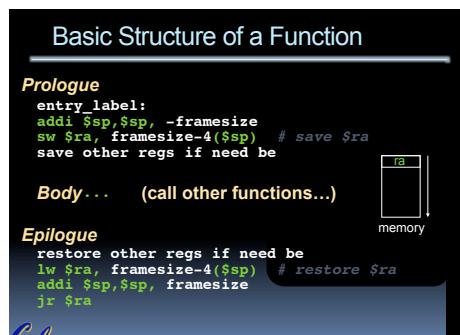


2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

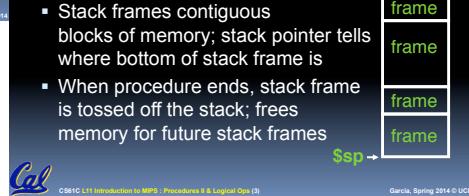
Lec 7.12

## Recall: cs61C THE STACK FRAME



## The Stack (review)

- Stack frame includes:
    - Return “instruction” address
    - Parameters
    - Space for other local variables `0xFFFFFFFF`
  - Stack frames contiguous blocks of memory; stack pointer tells where bottom of stack frame is
  - When procedure ends, stack frame is tossed off the stack; frees memory for future stack frames

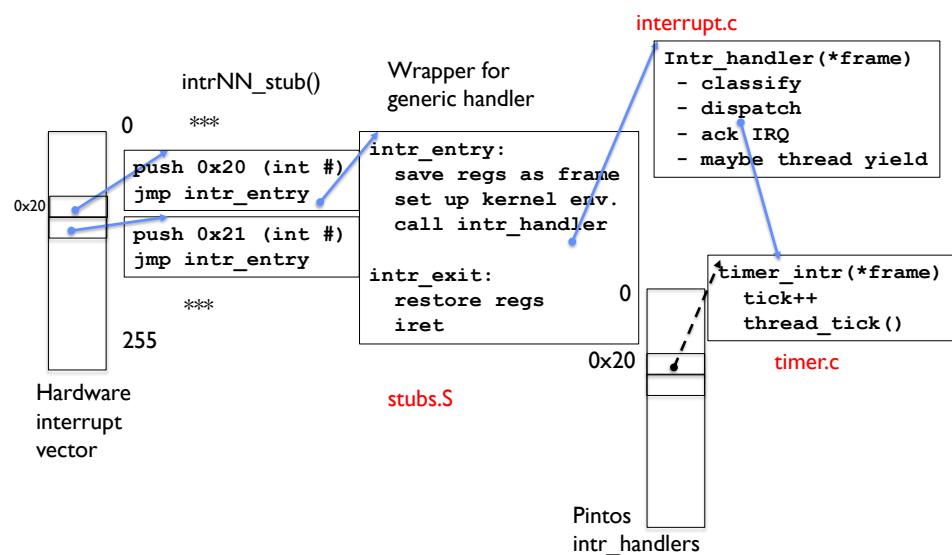


2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.13

## Pintos Interrupt Processing

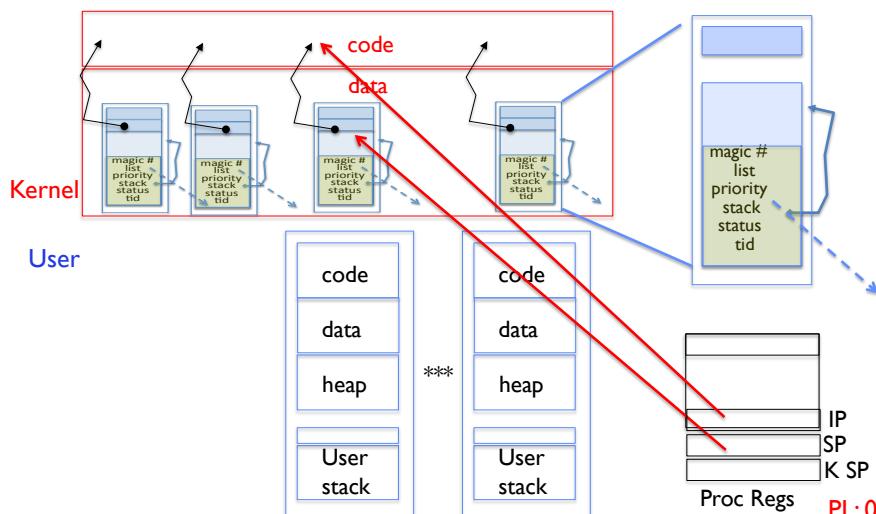


2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

2/7/18 Joseph and Ragan-Kelley CS162 © UCB Spring 2018 Lec 7.13 2/7/18 Joseph and Ragan-Kelley CS162 © UCB Spring 2018 Lec 7.14

## In Kernel thread



- Kernel threads execute with small stack in thread structure
  - Scheduler selects among ready kernel and user threads

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.15

Timer may trigger thread switch

- thread\_tick
    - Updates thread counters
    - If quanta exhausted, sets yield flag
  - thread\_yield
    - On path to rtn from interrupt
    - Sets current thread back to READY
    - Pushes it back on ready\_list
    - Calls schedule to select next thread to run upon iret
  - Schedule
    - Selects next thread to run
    - Calls switch\_threads to change regs to point to stack for thread to resume
    - Sets its status to RUNNING
    - If user thread, activates the process
    - Returns back to intr handler

2/7/18

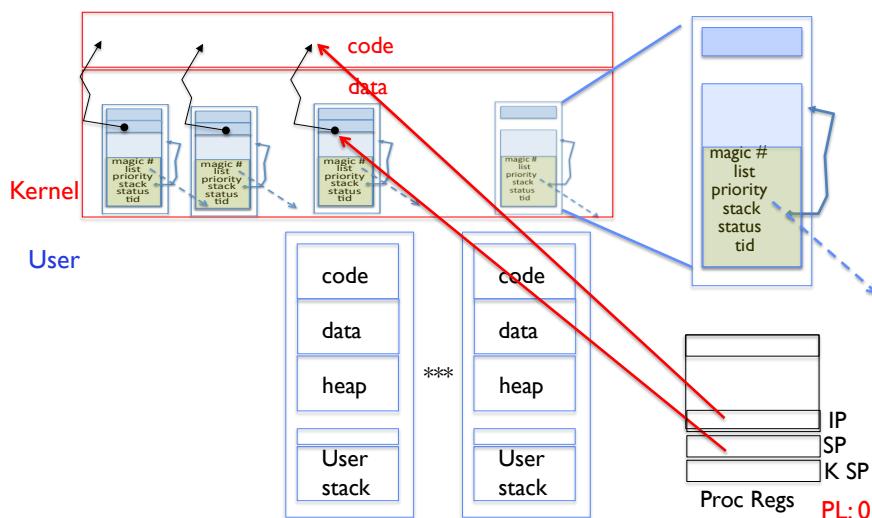
Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Joseph and Raegan-Kelley CS162 @ UICB Spring 2018 Lec 7.16

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7 | 6

## Thread Switch (switch.S)



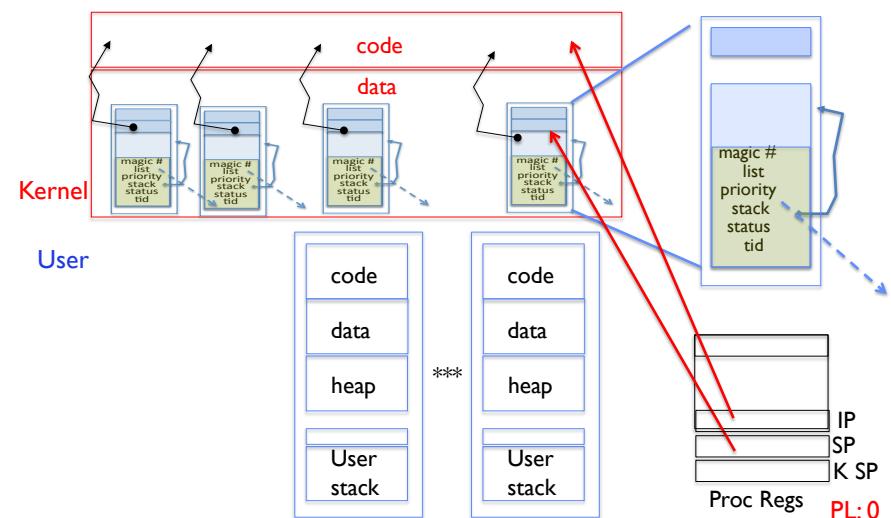
- switch\_threads: save regs on current small stack, change SP, return from destination thread's call to switch\_threads

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.17

## Switch to Kernel Thread for Process

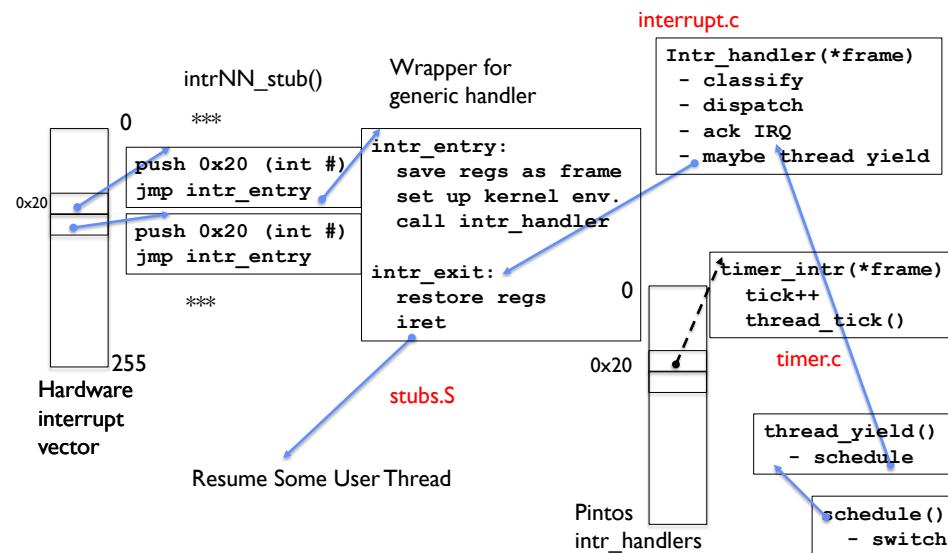


2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.18

## Pintos Return from Processing

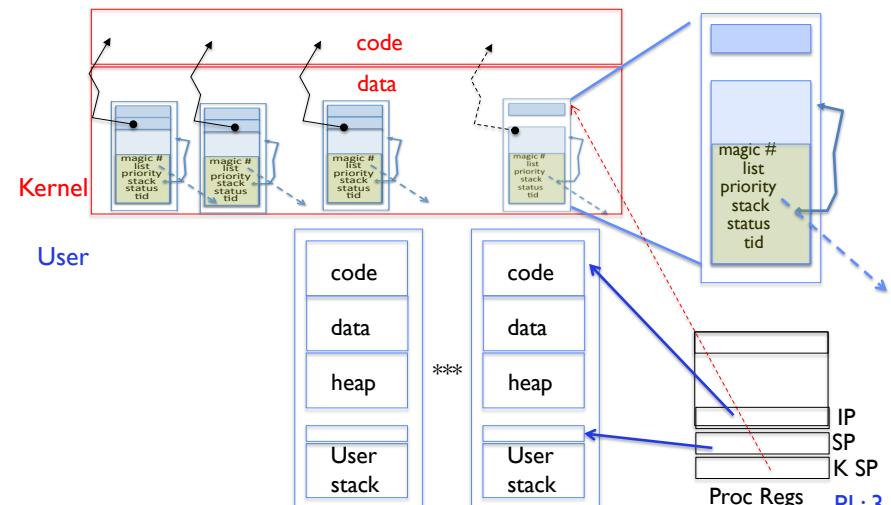


2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.19

## Kernel → User



- Interrupt return (iret) restores user stack and PL

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.20

## Rest of Today's Lecture

- The Concurrency Problem
- Synchronization Operations
- Higher-level Synchronization Abstractions
  - Semaphores, monitors, and condition variables

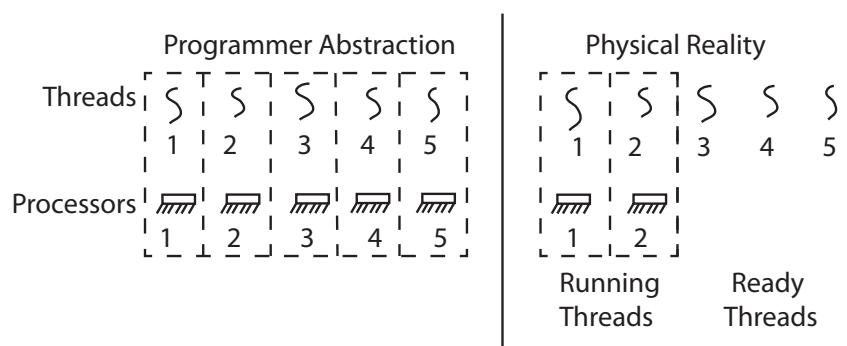


2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.21

## Recall: Thread Abstraction



- Infinite number of processors
- Threads execute with variable speed
  - Programs must be designed to work with any schedule

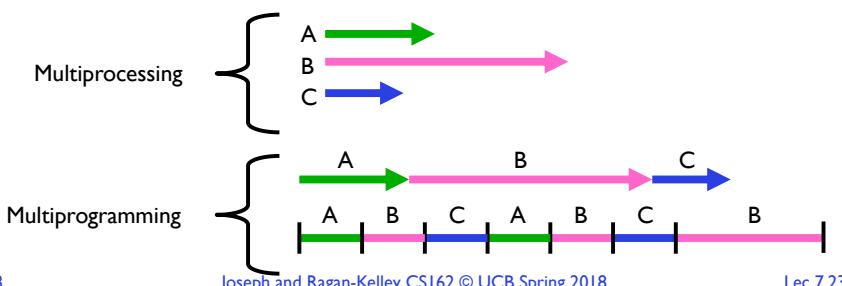
2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.22

## Multiprocessing vs Multiprogramming

- Remember Definitions:
  - Multiprocessing = Multiple CPUs or cores or hyperthreads (HW per-instruction interleaving)
  - Multiprogramming = Multiple Jobs or Processes
  - Multithreading = Multiple threads per Process
- What does it mean to run two threads “concurrently”?
  - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...



2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.23

## Correctness for systems with concurrent threads

- If dispatcher can schedule threads in any way, programs must work under all circumstances
  - Can you test for this?
  - How can you know if your program works?
- **Independent Threads:**
  - No state shared with other threads
  - Deterministic  $\Rightarrow$  Input state determines results
  - Reproducible  $\Rightarrow$  Can recreate Starting Conditions, I/O
  - Scheduling order doesn't matter (if `switch()` works!!!)
- **Cooperating Threads:**
  - Shared State between multiple threads
  - Non-deterministic
  - Non-reproducible
- Non-deterministic and Non-reproducible means that bugs can be intermittent
  - Sometimes called “Heisenbugs”

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.24

## Interactions Complicate Debugging

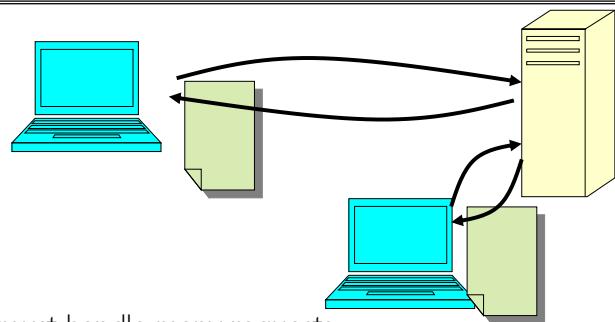
- Is any program truly independent?
  - Every process shares the file system, OS resources, network, etc.
  - Extreme example: buggy device driver causes thread A to crash “independent thread” B
- Non-deterministic errors are really difficult to find
  - Example: Memory layout of kernel+user programs
    - » depends on scheduling, which depends on timer/other things
    - » Original UNIX had a bunch of non-deterministic errors
  - Example: Something which does interesting I/O
    - » User typing of letters used to help generate secure keys

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.25

## High-level Example: Web Server



- Server must handle many requests
- Non-cooperating version:

```
serverLoop() {
    connection = AcceptCon();
    ProcessFork(ServiceWebPage(), connection);
}
```
- What are some disadvantages of this technique?

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.27

## Why allow cooperating threads?

- Advantage 1: Share resources
  - One computer, many users
  - One bank balance, many ATMs
    - » What if ATMs were only updated at night?
  - Embedded systems (robot control: coordinate arm & hand)
- Advantage 2: Speedup
  - Overlap I/O and computation
    - » Many different file systems do read-ahead
  - Multiprocessors – chop up program into parallel pieces
- Advantage 3: Modularity
  - More important than you might think
  - Chop large problem up into simpler pieces
    - » To compile, for instance, `gcc` calls `cpp` | `cc1` | `cc2` | `as` | `ld`
    - » Makes system easier to extend

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.26

## Threaded Web Server

- Instead, use a single process
- Multithreaded (cooperating) version:

```
serverLoop() {
    connection = AcceptCon();
    ThreadFork(ServiceWebPage(), connection);
}
```
- Looks almost the same, but has many advantages:
  - Can share file caches kept in memory, results of CGI scripts, other things
  - Threads are *much* cheaper to create than processes, so this has a lower per-request overhead
- What about Denial of Service attacks or digg / Slashdot effects?

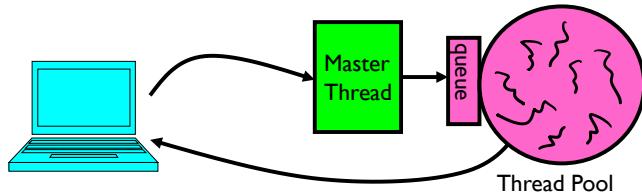
2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.28

## Thread Pools

- Problem with previous version: Unbounded Threads
  - When web-site becomes too popular – throughput sinks
- Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprogramming



```

master() {
    allocThreads(worker, queue);
    while(TRUE) {
        con=AcceptCon();
        Enqueue(queue, con);
        wakeUp(queue);
    }
}

worker(queue) {
    while(TRUE) {
        con=Dequeue(queue);
        if (con==null)
            sleepOn(queue);
        else
            ServiceWebPage(con);
    }
}

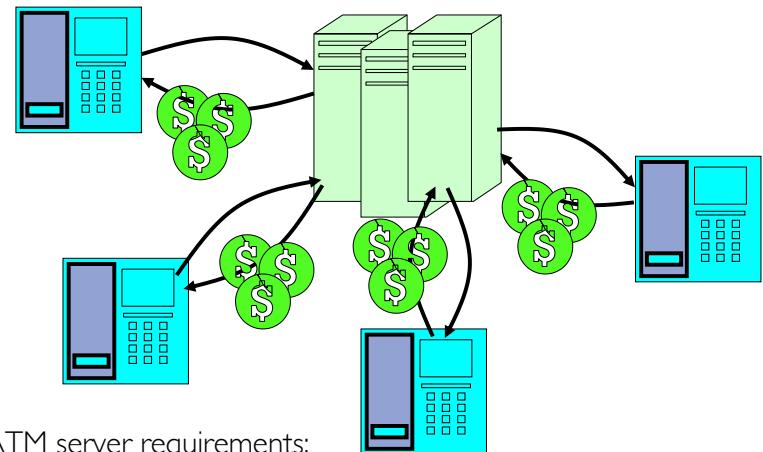
```

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.29

## ATM Bank Server



- ATM server requirements:

- Service a set of requests
- Do so without corrupting database
- Don't hand out too much money

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.30

## ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```

BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}

ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}

Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* May use disk I/O */
    acct->balance += amount;
    StoreAccount(acct);      /* Involves disk I/O */
}

```

- How could we speed this up?
  - More than one request being processed at once
  - Event driven (overlap computation and I/O)
  - Multiple threads (multi-proc, or overlap comp and I/O)

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.31

## Event Driven Version of ATM server

- Suppose we only had one CPU
  - Still like to overlap I/O with computation
  - Without threads, we would have to rewrite in event-driven style
- Example

```

BankServer() {
    while(TRUE) {
        event = WaitForNextEvent();
        if (event == ATMRequest)
            StartOnRequest();
        else if (event == AcctAvail)
            ContinueRequest();
        else if (event == AcctStored)
            FinishRequest();
    }
}

```

- What if we missed a blocking I/O step?
- What if we have to split code into hundreds of pieces which could be blocking?
- This technique is used for programming GPUs (Graphics Processing Unit)

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.32

## Can Threads Make This Easier?

- Threads yield overlapped I/O and computation without having to “deconstruct” code into non-blocking fragments
  - One thread per request
- Requests proceed to completion, blocking as required:

```
Deposit(acctId, amount) {  
    acct = GetAccount(actId); /* May use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct); /* Involves disk I/O */  
}
```

- Unfortunately, shared state can get corrupted:

<u>Thread 1</u> load r1, acct->balance  add r1, amount1 store r1, acct->balance	<u>Thread 2</u> load r1, acct->balance add r1, amount2 store r1, acct->balance
---	---

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.33

## Atomic Operations

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- Atomic Operation:** an operation that always runs to completion or not at all
  - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
  - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
  - Consequently – weird example that produces “3” on previous slide can’t happen
- Many instructions are not atomic
  - Double-precision floating point store often not atomic
  - VAX and IBM 360 had an instruction to copy a whole array

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.35

## Problem is at the Lowest Level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A $x = 1;$ Thread B $y = 2;$ 

- However, what about (Initially,  $y = 12$ ):

Thread A $x = 1;$  $x = y + 1;$ Thread B $y = 2;$  $y = y * 2;$ 

- What are the possible values of  $x$ ?

- Or, what are the possible values of  $x$  below?

Thread A $x = 1;$ Thread B $x = 2;$ 

- $X$  could be 1 or 2 (non-deterministic!)

- Could even be 3 for serial processors:

» Thread A writes 0001, B writes 0010 → scheduling order ABABABBA yields 3!

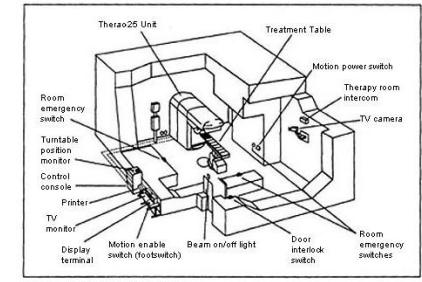
2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.34

## Correctness Requirements

- Threaded programs must work for all interleavings of thread instruction sequences
  - Cooperating threads inherently non-deterministic and non-reproducible
  - Really hard to debug unless carefully designed!
- Examples:



2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.36

## Administrivia

- Group/Section assignments finalized!
  - If you are not in group, talk to us immediately!
- Attend assigned sections
  - Need to know your TA!
    - » Participation is 8% of your grade
    - » Should attend section with your TA
- First design doc due next **Wednesday**
  - This means you should be well on your way with Project 1
  - Watch for notification from your TA to sign up for design review
- Basic semaphores work in PintOS!
  - However, you will need to implement priority scheduling behavior both in semaphore and ready queue

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.37

## BREAK

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.38

## Motivation: “Too Much Milk”

- Great thing about OS's – analogy between problems in OS and problems in real life
  - Help you understand real life problems better
  - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.39

## Definitions

- **Synchronization:** using atomic operations to ensure cooperation between threads
  - For now, only loads and stores are atomic
  - We are going to show that its hard to build anything useful with only reads and writes
- **Mutual Exclusion:** ensuring that only one thread does a particular thing at a time
  - One thread excludes the other while doing its task
- **Critical Section:** piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code
  - Critical section is the result of mutual exclusion
  - Critical section and mutual exclusion are two ways of describing the same thing

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.40

## More Definitions

- **Lock:** prevents someone from doing something
  - Lock before entering critical section and before accessing shared data
  - Unlock when leaving, after accessing shared data
  - Wait if locked
- » Important idea: all synchronization involves waiting
- For example: fix the milk problem by putting a key on the refrigerator
  - Lock it and take key if you are going to go buy milk
  - Fixes too much: roommate angry if only wants OJ



– Of Course – We don't know how to make a lock yet

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.41

## Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of "lock")
  - Remove note after buying (kind of "unlock")
  - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy Milk;  
        remove note;  
    }  
}
```



2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.43

## Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
  - Impulse is to start coding first, then when it doesn't work, pull hair out
  - Instead, think first, then code
  - Always write down behavior first
- What are the correctness properties for the "Too much milk" problem???
  - Never more than one person buys
  - Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.42

## Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of "lock")
  - Remove note after buying (kind of "unlock")
  - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
Thread A  
if (noMilk) {  
    context switch  
    if (noNote) {  
        if (noMilk) {  
            if (noNote) {  
                leave Note;  
                buy Milk;  
                remove Note;  
            }  
        }  
    }  
}  
  
leave Note;  
buy Milk;  
remove Note;
```

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.44

## Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
    - Leave a note before buying (kind of “lock”)
    - Remove note after buying (kind of “unlock”)
    - Don’t buy if note (wait)
  - Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



- Result?
    - Still too much milk **but only occasionally!**
    - Thread can get context switched after checking milk and note but before buying milk!
  - Solution makes problem worse since fails **intermittently**
    - Makes it really hard to debug...
    - Must work despite what the dispatcher does!

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.45

## Too Much Milk Solution #2

- How about labeled notes?
    - Now we can leave note before checking
  - Algorithm looks like this:

```
Thread A
leave note A;
if (noNote B) {
    if (noMilk) {
        buy Milk;
    }
}
remove note A;
```

```
Thread B  
leave note B;  
if (noNoteA) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note B;
```

- Does this work?
  - Possible for neither thread to buy milk
    - Context switches at exactly the wrong times can lead each to think that the other is going to buy
  - Really insidious:
    - **Extremely unlikely** that this would happen, but will at worse possible time
    - Probably something like this in UNIX

2/7/18

Joseph and Bagan-Kelley CSI62 @ UCB Spring 2018

| ec 7.47

## Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
    - Let's try to fix this by placing note first
  - Another try at previous solution:

```
leave Note;  
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
    }  
}  
remove note;
```



- What happens here?
    - Well, with human, probably nothing bad
    - With computer: no one ever buys milk

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.46

## Too Much Milk Solution #2: Problem!



- I thought you had the milk! But I thought you had the milk!
  - This kind of lockup is called “starvation!”

2/7/18

Joseph and Bagan-Kelley CS162 © UCB Spring 2018

Lec 7.48

## Too Much Milk Solution #3

- Here is a possible two-note solution:

```
Thread A                                Thread B
leave note A;
while (note B) {\\"X
    do nothing;
}
if (noMilk) {
    buy milk;
}
remove note A;

leave note B;
if (noNote A) {\\"Y
    if (noMilk) {
        buy milk;
    }
remove note B;
```

- Does this work? Yes. Both can guarantee that:

- It is safe to buy, or
- Other will buy, ok to quit

- At X:

- If no note B, safe for A to buy,
- Otherwise wait to find out what will happen

- At Y:

- If no note A, safe for B to buy
- Otherwise, A is either buying or waiting for B to quit

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.49

## Case I

- "leave note A" happens before "if (noNote A)"

```
leave note A;                                leave note B;
while (note B) {\\"X                         if (noNote A) {\\"Y
    do nothing;                               if (noMilk) {
}
};                                         buy milk;
}                                           remove note A;
```

happened before

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.50

## Case I

- "leave note A" happens before "if (noNote A)"

```
leave note A;                                leave note B;
while (note B) {\\"X                         if (noNote A) {\\"Y
    do nothing;                               if (noMilk) {
}
};                                         buy milk;
}                                           remove note B;
```

happened before

```
if (noMilk) {
    buy milk;
}
remove note A;
```

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.51

## Case I

- "leave note A" happens before "if (noNote A)"

```
leave note A;                                leave note B;
while (note B) {\\"X                         if (noNote A) {\\"Y
    do nothing;                               if (noMilk) {
}
};                                         buy milk;
}                                           remove note A;
```

happened before

```
Wait for
note B to
be removed
```

```
if (noMilk) {
    buy milk;
}
remove note B;
```

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.52

## Case 2

- "if (noNote A)" happens before "leave note A"

```
leave note A; while (note B) {\\"X  
    do nothing;  
};  
  
if (noMilk) {  
    buy milk;  
}  
remove note B;  
  
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.53

## Case 2

- "if (noNote A)" happens before "leave note A"

```
leave note A; while (note B) {\\"X  
    do nothing;  
};  
  
if (noMilk) {  
    buy milk;  
}  
remove note A;  
  
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.54

## Case 2

- "if (noNote A)" happens before "leave note A"

```
leave note A; while (note B) {\\"X  
    do nothing;  
};  
  
if (noMilk) {  
    buy milk;  
}  
remove note A;  
  
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.55

## Review: Solution #3 discussion

- Our solution protects a single "CriticalSection" piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
  - Really complex – even for this simple an example
    - » Hard to convince yourself that this really works
  - A's code is different from B's – what if lots of threads?
    - » Code would have to be slightly different for each thread
  - While A is waiting, it is consuming CPU time
    - » This is called "busy-waiting"
- There's a better way
  - Have hardware provide higher-level primitives than atomic load & store
  - Build even higher-level programming abstractions on this hardware support

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.56

## Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock
  - `lock.Acquire()` – wait until lock is free, then grab
  - `lock.Release()` – Unlock, waking up anyone waiting
  - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
millock.Acquire();
if (nomilk)
    buy milk;
millock.Release();
```
- Once again, section of code between `Acquire()` and `Release()` called a “**Critical Section**”
- Of course, you can make this even simpler: suppose you are out of ice cream instead of milk
  - Skip the test since you always need more ice cream ;-)

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.57

## Summary

- Concurrent threads are a very useful abstraction
  - Allow transparent overlapping of computation and I/O
  - Allow use of parallel processing when available
- Concurrent threads introduce problems when accessing shared data
  - Programs must be insensitive to arbitrary interleavings
  - Without careful design, shared variables can become completely inconsistent
- Important concept: Atomic Operations
  - An operation that runs to completion or not at all
  - These are the primitives on which to construct various synchronization primitives

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.59

## Where are we going with synchronization?

Programs	Shared Programs			
Higher-level API	Locks	Semaphores	Monitors	Send/Receive
Hardware	Load/Store	Disable Ints	Test&Set	Compare&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
  - Everything is pretty painful if only atomic primitives are load and store
  - Need to provide primitives useful at user-level

2/7/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 7.58