

CS162

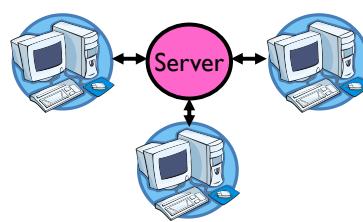
Operating Systems and Systems Programming

Lecture 21

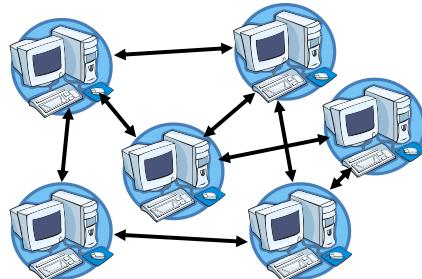
Layering, E2E Argument

April 11th, 2018

Profs. Anthony D. Joseph & Jonathan Ragan-Kelley
<http://cs162.eecs.Berkeley.edu>



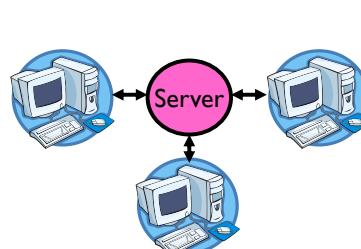
Client/Server Model



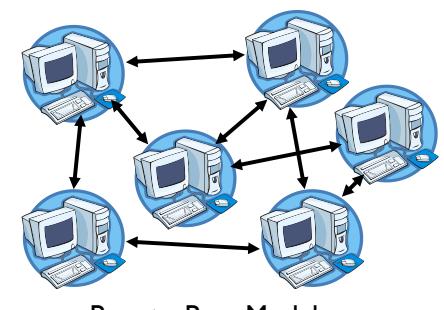
Peer-to-Peer Model

- **Distributed System:** physically separate computers working together on some task
 - Early model: multiple servers working together
 - » Probably in the same room or building
 - » Often called a “cluster”
 - Later models: peer-to-peer/wide-spread collaboration

Centralized vs Distributed Systems



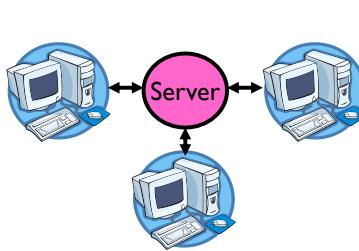
Client/Server Model



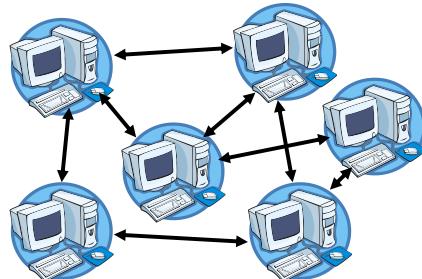
Peer-to-Peer Model

- **Centralized System:** System in which major functions are performed by a single physical computer
 - Originally, everything on single computer
 - Later: client/server model

Centralized vs Distributed Systems



Client/Server Model



Peer-to-Peer Model

- **Distributed System:** physically separate computers working together on some task
 - Early model: multiple servers working together
 - » Probably in the same room or building
 - » Often called a “cluster”
 - Later models: peer-to-peer/wide-spread collaboration

Distributed Systems: Motivation/Issues/Promise

- Why do we want distributed systems?
 - Cheaper and easier to build lots of simple computers
 - Easier to add power incrementally
 - Users can have complete control over some components
 - Collaboration: much easier for users to collaborate through network resources (such as network file systems)
- The *promise* of distributed systems:
 - Higher availability: one machine goes down, use another
 - Better durability: store data in multiple locations
 - More security: each piece easier to make secure

Distributed Systems: Reality

- Reality has been disappointing
 - Worse availability: depend on every machine being up
 - » Lamport: “**a distributed system is one where I can't do work because some machine I've never heard of isn't working!**”
 - Worse reliability: can lose data if any machine crashes
 - Worse security: anyone in world can break into system
 - Coordination is more difficult
 - Must coordinate multiple copies of shared state information (using only a network)
 - What would be easy in a centralized system becomes a lot more difficult

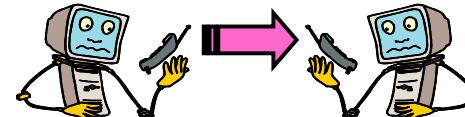
4/11/18

CS162 ©UCB Spring 2018

Lec 21.5

Distributed Systems: Goals/Requirements

- **Transparency:** the ability of the system to mask its complexity behind a simple interface
 - Possible transparencies:
 - **Location:** Can't tell where resources are located
 - **Migration:** Resources may move without the user knowing
 - **Replication:** Can't tell how many copies of resource exist
 - **Concurrency:** Can't tell how many users there are
 - **Parallelism:** System may speed up large jobs by splitting them into smaller pieces
 - **Fault Tolerance:** System may hide various things that go wrong
 - Transparency and collaboration require some way for different processors to communicate with one another

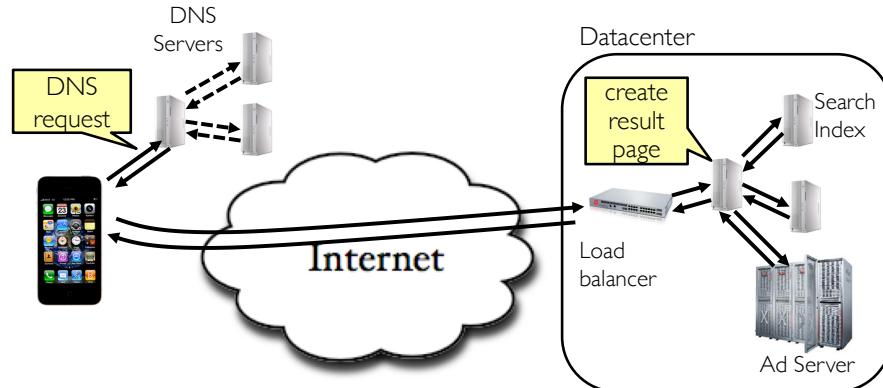


4/11/18

CS162 ©UCB Spring 2018

Lec 21.6

Example: What's in a Search Query?



- Complex interaction of multiple components in multiple administrative domains

4/11/18

CS162 ©UCB Spring 2018

Lec 21.7

Goals for Today

- Layering
 - End-to-end arguments

Some slides generated from Vern Paxson and Scott Shenker lecture notes

4/11/18

CS162 ©UCB Spring 2018

Lec 21.8

Why is Networking Important?

- Virtually all apps you use communicate over network
 - Many times main functionality is implemented remotely (e.g., Google services, Amazon, Facebook, Twitter, ...)
- Thus, connectivity is key service provided by an OS
 - Many times, connectivity issues → among top complaints



4/11/18

CS162 @UCB Spring 2018

Lec 21.9

Why is Networking Important?

- Virtually all apps you use communicate over network
 - Many times main functionality is implemented remotely (e.g., Google services, Amazon, Facebook, Twitter, ...)
- Thus, connectivity is key service provided by an OS
 - Many times, connectivity issues → among top complaints

- Some of the hottest opportunities in the OS space:
 - Optimize OS for network elements (e.g., intrusion detection, firewalls)
 - OSes for Software Defined Networks (SDNs)

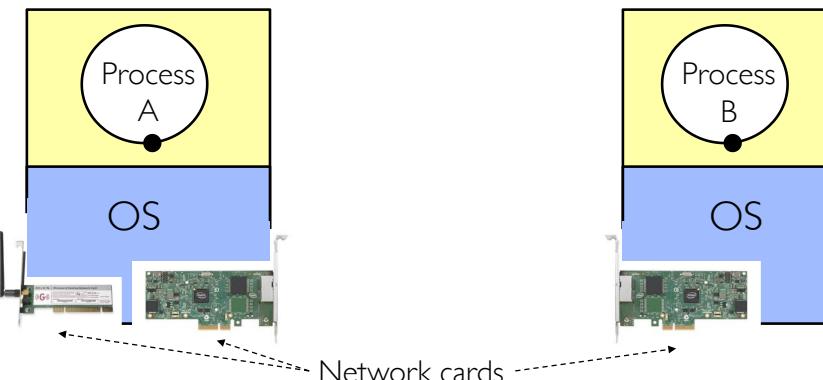
4/11/18

CS162 @UCB Spring 2018

Lec 21.10

Network Concepts

- **Network (interface) card/controller:** hardware that physically connects a computer to the network
- A computer can have more than one networking cards
 - E.g., one card for wired network, and one for wireless network



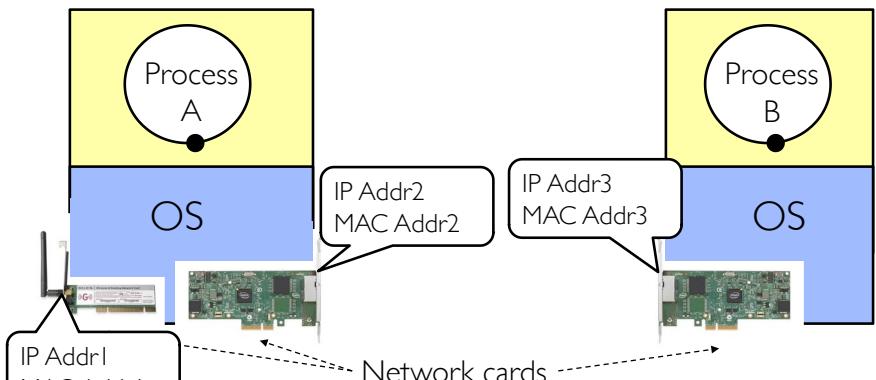
4/11/18

CS162 @UCB Spring 2018

Lec 21.11

Network Concepts (cont'd)

- Typically, each network card is associated two addresses:
 - Media Access Control (MAC), or physical, address
 - IP, or network, address (can be shared by network cards on same host)



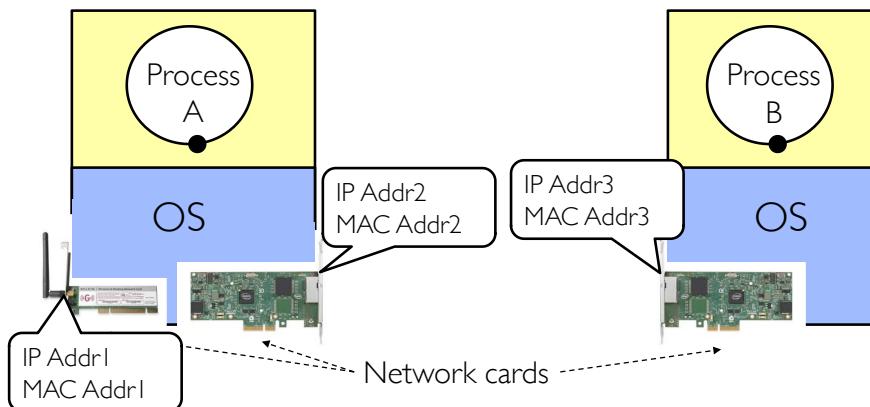
4/11/18

CS162 @UCB Spring 2018

Lec 21.12

Network Concepts (cont' d)

- **MAC address:** 48-bit unique identifier assigned by card vendor
- **IP Address:** 32-bit (or 128-bit for IPv6) address assigned by network administrator or dynamically when computer connects to network



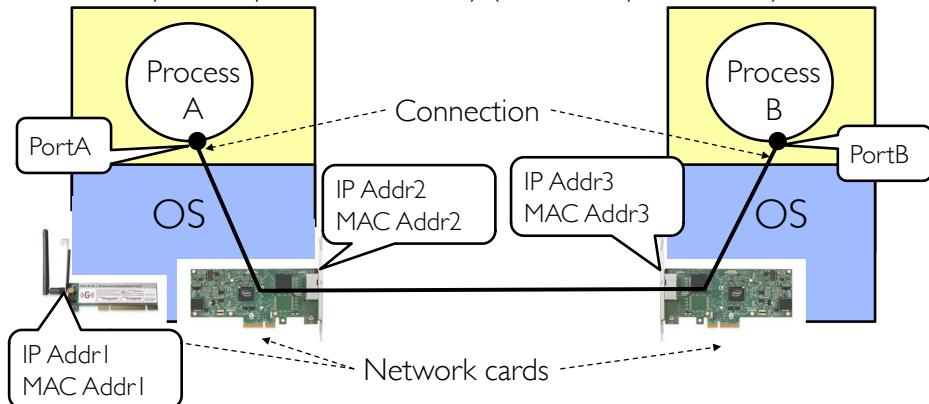
4/11/18

CS162 @UCB Spring 2018

Lec 21.13

Network Concepts (cont' d)

- **Connection:** communication channel between two processes
- Each endpoint is identified by a **port number**
 - **Port number:** 16-bit identifier assigned by app or OS
 - Globally, an endpoint is identified by (IP address, port number)



4/11/18

CS162 @UCB Spring 2018

Lec 21.14

Main Network Functionalities

- **Delivery:** deliver packets between any two hosts in the Internet
 - E.g., how do you deliver a packet from a host in Berkeley to a host in Tokyo?
- **Reliability:** tolerate packet losses
 - E.g., how do you ensure all bits of a file are delivered in the presence of packet losses?
- **Flow control:** avoid overflowing the receiver buffer
 - Recall our bounded buffer example: stop sender from overflowing buffer
 - E.g., how do you ensure that a sever that can send at 10Gbps doesn't overwhelm a 4G phone?
- **Congestion control:** avoid overflowing the buffer of a router along the path
 - What happens if we don't do it?

4/11/18

CS162 @UCB Spring 2018

Lec 21.15

Protocol Standardization

- Ensure communicating hosts speak the same protocol
 - Standardization to enable multiple implementations
 - Or, the same folks have to write all the software
- Standardization: Internet Engineering Task Force
 - Based on working groups that focus on specific issues
 - Produces "Request For Comments" (RFCs)
 - » Promoted to standards via rough consensus and running code
 - IETF Web site is <http://www.ietf.org>
 - RFCs archived at <http://www.rfc-editor.org>
- De facto standards: same folks writing the code
 - P2P file sharing, Skype, <your protocol here>...

4/11/18

CS162 @UCB Spring 2018

Lec 21.16

Layering: The Problem

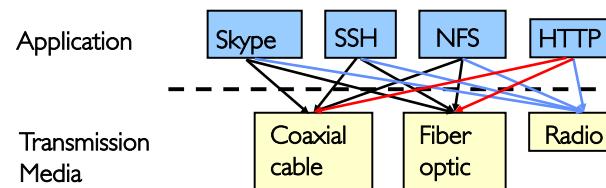
- Many different applications
 - email, web, P2P, etc.
- Many different network styles and technologies
 - Circuit-switched vs packet-switched, etc.
 - Wireless vs. wired vs optical, etc.
- How do we organize this mess?

4/11/18

CS162 ©UCB Spring 2018

Lec 21.17

The Problem (cont' d)



- Re-implement every application for every technology?
- No! But how does the Internet design avoid this?

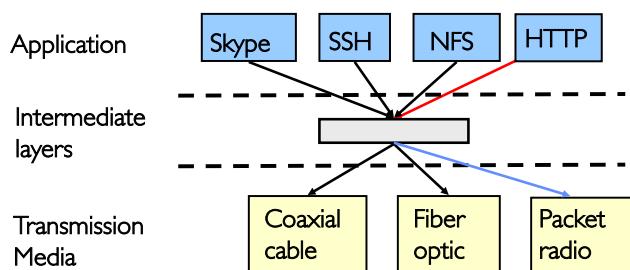
4/11/18

CS162 ©UCB Spring 2018

Lec 21.18

Solution: Intermediate Layers

- Introduce intermediate layers that provide **set of abstractions** for various network functionality & technologies
 - A new app/media implemented only once
 - Variation on “add another level of indirection”



4/11/18

CS162 ©UCB Spring 2018

Lec 21.19

Software System Modularity

Partition system into modules & abstractions:

- Well-defined interfaces give flexibility
 - **Hides** implementation - thus, it can be freely changed
 - Extend functionality of system by adding new modules
- E.g., libraries encapsulating set of functionality
- E.g., programming language + compiler abstracts away not only how the particular CPU works ...
 - ... but also the **basic computational model**
- Well-defined interfaces hide information
 - Isolate **assumptions**
 - Present high-level **abstractions**
 - But can impair performance

4/11/18

CS162 ©UCB Spring 2018

Lec 21.20

Network System Modularity

Like software modularity, but:

- Implementation distributed across many machines (routers and hosts)
- Must decide:
 - How to break system into modules
 - » Layering
 - What functionality does each module implement
 - » End-to-End Principle
 - Where state is stored
 - » Fate-sharing
- We will address these choices in turn

4/11/18

CS162 ©UCB Spring 2018

Lec 21.21

Properties of Layers (OSI Model)

- **Service:** what a layer does
- **Service interface:** how to access the service
 - Interface for layer above
- **Protocol (peer interface):** how peers communicate to implement the service
 - Set of rules and formats that specify the communication between network elements
 - Does **not** specify the implementation on a single machine, but how the layer is implemented **between** machines

4/11/18

CS162 ©UCB Spring 2018

Lec 21.23

Layering: A Modular Approach

- Partition the system
 - Each layer **solely** relies on services from layer below
 - Each layer **solely** exports services to layer above
- Interface between layers defines interaction
 - Hides implementation details
 - Layers can change without disturbing other layers

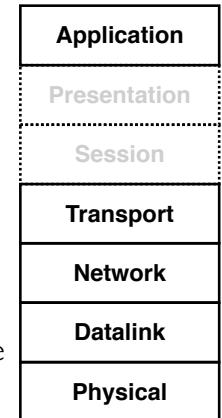
4/11/18

CS162 ©UCB Spring 2018

Lec 21.22

OSI Layering Model

- Open Systems Interconnection (OSI) model
 - Developed by International Organization for Standardization (ISO) in 1984
 - **Seven** layers
- Internet Protocol (IP)
 - Only **five** layers
 - The functionalities of the missing layers (i.e., Presentation and Session) are provided by the Application layer



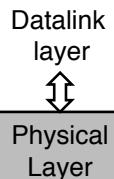
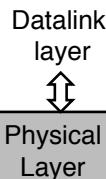
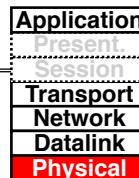
4/11/18

CS162 ©UCB Spring 2018

Lec 21.24

Physical Layer (1)

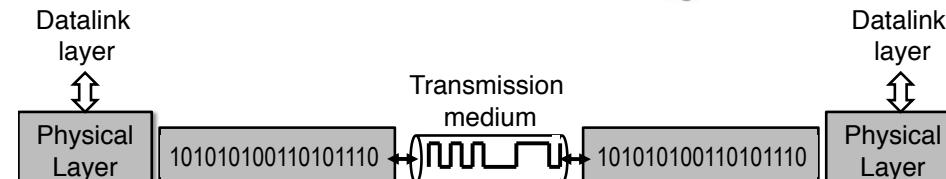
- Service:** move information between two systems connected by a physical link
- Interface:** specifies how to send and receive bits
- Protocol:** **coding scheme** used to represent a bit, voltage levels, duration of a bit
- Examples: coaxial cable, optical fiber links; transmitters, receivers



4/11/18

CS162 ©UCB Spring 2018

Lec 21.25

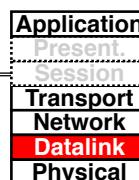


Datalink Layer (2)

- Service:**

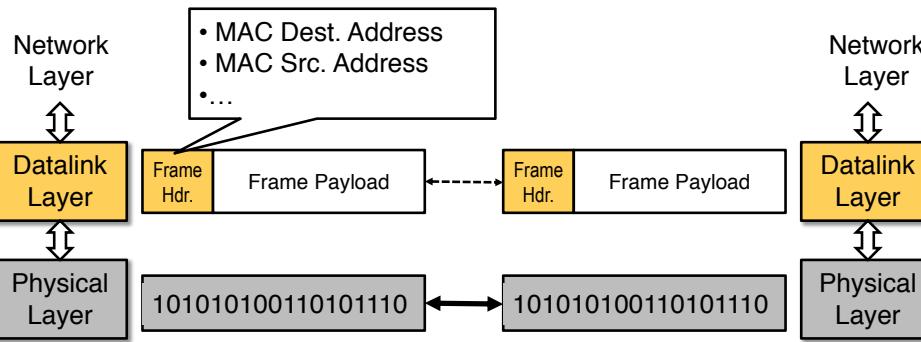
- Enable end hosts to exchange frames (atomic messages) on the same physical line or wireless link
- Possible other services:
 - » **Arbitrate access** to common physical media
 - » May provide **reliable transmission, flow control**

- Interface:** send *frames* to other end hosts; receive *frames* addressed to end host
- Protocols:** addressing, Media Access Control (MAC) (e.g., CSMA/CD)
 - *Carrier Sense Multiple Access / Collision Detection*



Datalink Layer (2)

- Each frame has a header which contains a source and a destination MAC address
- MAC (Media Access Control) address
 - Uniquely identifies a network interface
 - 48-bit, assigned by the device manufacturer



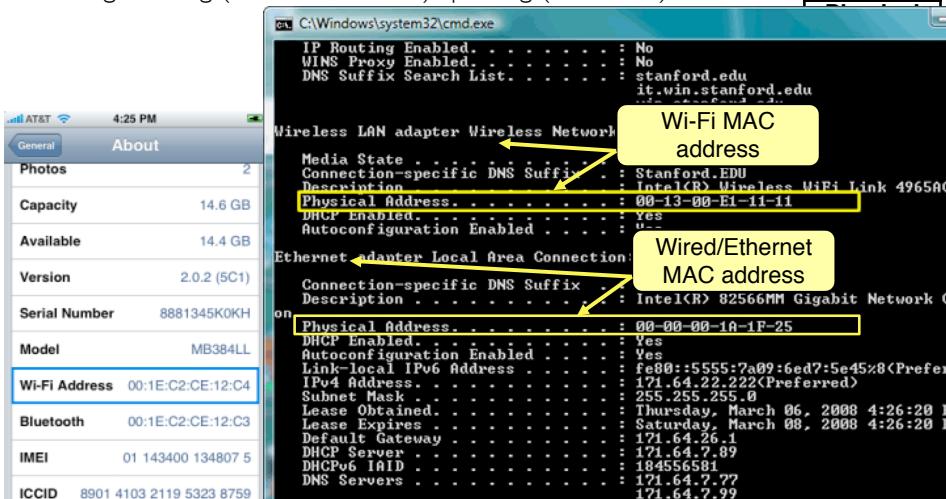
4/11/18

CS162 ©UCB Spring 2018

Lec 21.27

MAC Address Examples

- Can easily find MAC addr. on your machine/device:
 - E.g., ifconfig (Linux, Mac OS X), ipconfig (Windows)



4/11/18

CS162 ©UCB Spring 2018

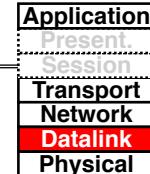
Lec 21.28

Local Area Networks (LANs)

- LAN: group of hosts/devices that
 - are in the same geographical proximity (e.g., same building, room)
 - use same physical communication technology
- Examples:
 - all laptops connected wirelessly at a Starbucks café
 - all devices and computers at home
 - all hosts connected to wired Ethernet in an office



Ethernet cable and port



4/11/18

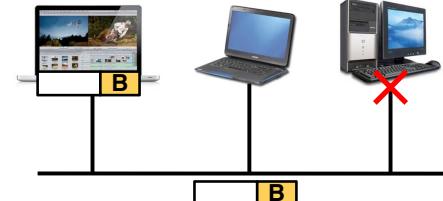
CS162 ©UCB Spring 2018

Lec 21.29

LANs

- All hosts in a LAN can share same physical communication media
 - Also called, broadcast channel
- Each frame is delivered to every host
- If a host is not the intended recipient, it drops the frame

MAC Addr: A MAC Addr: B MAC Addr: C



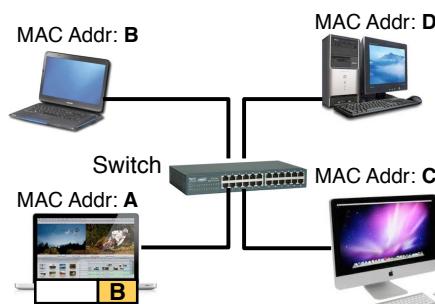
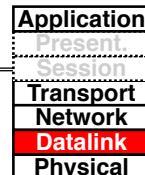
4/11/18

CS162 ©UCB Spring 2018

Lec 21.30

Switches

- Hosts in same LAN can be also connected by switches
- A switch forwards frames only to intended recipients
 - Far more efficient than broadcast channel



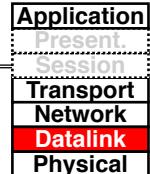
4/11/18

CS162 ©UCB Spring 2018

Lec 21.31

Media Access Control (MAC) Protocols

- Problem:
 - How do hosts access a broadcast media?
 - How do they avoid collisions?
- Three solutions:
 - Channel partition
 - “Taking turns”
 - Random access

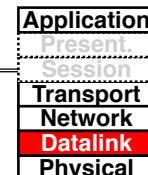


4/11/18

CS162 ©UCB Spring 2018

Lec 21.32

MAC Protocols



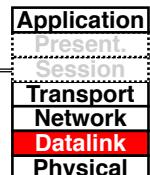
- Channel partitioning protocols:
 - Allocate I/N bandwidth to every host
 - Share channel efficiently and fairly at high load
 - Inefficient at low load (where load = # senders):
 - » I/N bandwidth allocated even if only 1 active node!
 - E.g., Frequency Division Multiple Access (FDMA); optical networks
- “Taking turns” protocols:
 - Pass a token around active hosts
 - A host can only send data if it has the token
 - More efficient at low loads: single node can use >> I/N bandwidth
 - Overhead in acquiring the token
 - Vulnerable to failures (e.g., failed node or lost token)
 - E.g., Token ring

4/11/18

CS162 ©UCB Spring 2018

Lec 21.33

MAC Protocols



- Random Access
 - Efficient at low load: single node can fully utilize channel
 - High load: collision overhead
- Key ideas of random access:
 - Carrier sense (CS)
 - » Listen before speaking, and don't interrupt
 - » Checking if someone else is already sending data
 - » ... and waiting till the other node is done
 - Collision detection (CD)
 - » If someone else starts talking at the same time, stop
 - » Realizing when two nodes are transmitting at once
 - » ...by detecting that the data on the wire is garbled
 - Randomness
 - » Don't start talking again right away
 - » Waiting for a random time before trying again
- Examples: CSMA/CD, Ethernet, best known implementation

4/11/18

CS162 ©UCB Spring 2018

Lec 21.34

Administrivia

- Project 3: initial design doc due today at 11:59PM
- Midterm 3 coming up on **Wed 4/25 6:30-8PM**
 - All topics up to and including Lecture 23
 - » Focus will be on Lectures 17 – 23 and associated readings, and Projects 3
 - » But expect 20-30% questions from materials from Lectures 1-16
 - LKS 245, Hearst Field Annex A1, VLSB 2060, Barrows 20, Wurster 102 ([see Piazza for your room assignment](#))
 - Closed book
 - 2 pages hand-written notes both sides

4/11/18

CS162 ©UCB Spring 2018

Lec 21.35

BREAK

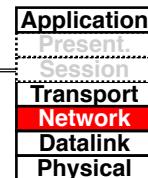
4/11/18

CS162 ©UCB Spring 2018

Lec 21.36

(Inter) Network Layer (3)

- Service:**
 - Deliver packets to specified **network addresses** across multiple datalink layer networks
 - Possible other services:
 - Packet scheduling/priority
 - Buffer management
- Interface:** send *packets* to specified network address destination; receive packets destined for end host
- Protocols:** define network addresses (globally unique); construct forwarding tables; packet forwarding



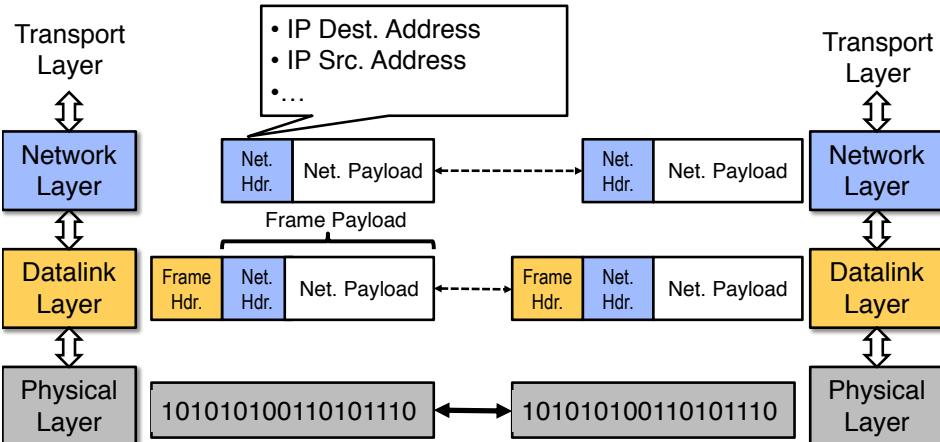
4/11/18

CS162 ©UCB Spring 2018

Lec 21.37

(Inter) Network Layer (3)

- IP address:** unique addr. assigned to network device
- Assigned by network administrator or dynamically when host connects to network



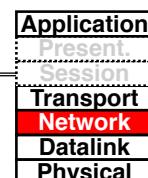
4/11/18

CS162 ©UCB Spring 2018

Lec 21.38

Wide Area Network

- Wide Area Network (WAN):** network that covers a broad area (e.g., city, state, country, entire world)
 - E.g., Internet is a WAN
- WAN connects multiple datalink layer networks (LANs)
- Datalink layer networks are connected by **routers**
 - Different LANs can use different communication technology (e.g., wireless, cellular, optics, wired)



Host A
(IP A)



R2

R4

R3

Host B
(IP B)



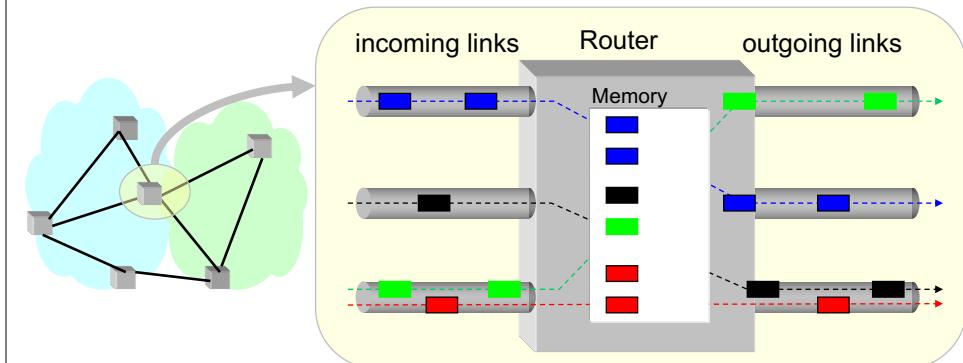
4/11/18

CS162 ©UCB Spring 2018

Lec 21.39

Routers

- Forward each packet received on an **incoming link** to an **outgoing link** based on packet's destination IP address (towards its destination)
- Store & forward:** packets are buffered before being forwarded
- Forwarding table:** mapping between IP address and the output link



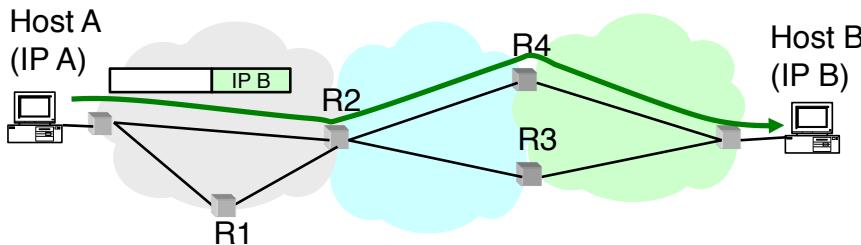
4/11/18

CS162 ©UCB Spring 2018

Lec 21.40

Packet Forwarding

- Upon receiving a packet, a router
 - read the IP destination address of the packet
 - consults its forwarding table → output port
 - fowards packet to corresponding output port



4/11/18

CS162 ©UCB Spring 2018

Lec 21.41

IP Addresses vs. MAC Addresses

- Why does packet forwarding using IP addr. scale?
- Because IP addresses can be aggregated
 - E.g., all IP addresses at UC Berkeley start with **0xA9E5**, i.e., any address of form **0xA9E5******* belongs to Berkeley
 - Thus, a router in NY needs to keep a **single** entry for **all** hosts at Berkeley
 - If we were using MAC addresses the NY router would need to maintain **an entry for every** Berkeley host!!
- Analogy:
 - Give this letter to person with SSN: 123-45-6789 vs.
 - Give this letter to “John Smith, 123 First Street, LA, US”



4/11/18

CS162 ©UCB Spring 2018

Lec 21.43

IP Addresses vs. MAC Addresses

- Why not use MAC addresses for routing?
 - Doesn't scale
- Analogy
 - MAC address → SSN
 - IP address → (unreadable) home address
- MAC address: uniquely associated to the device for the entire lifetime of the device
- IP address: changes as the device location changes
 - Your notebook IP address at school is different from home



4/11/18

Lec 21.42

The Internet Protocol (IP)

- Internet Protocol: Internet's network layer
- Service it provides: “Best-Effort” Packet Delivery
 - Tries its “best” to deliver packet to its destination
 - Packets may be lost
 - Packets may be corrupted
 - Packets may be delivered out of order



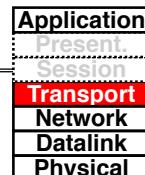
4/11/18

CS162 ©UCB Spring 2018

Lec 21.44

Transport Layer (4)

- Service:**
 - Provide end-to-end communication between processes
 - Demultiplexing of communication between hosts
 - Possible other services:
 - Reliability in the presence of errors
 - Timing properties
 - Rate adaptation (flow-control, congestion control)
- Interface:** send message to specific process at given destination; local process receives messages sent to it
- Protocol:** port numbers, perhaps implement reliability, flow control, packetization of large messages, framing
- Examples: TCP and UDP



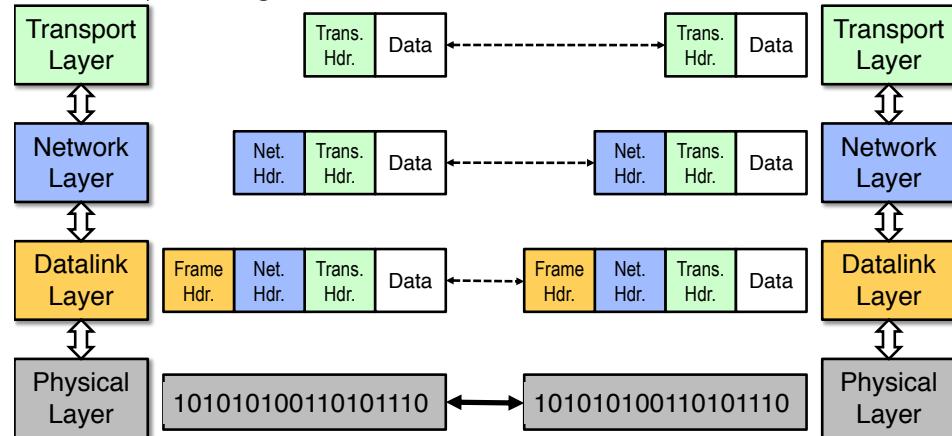
4/11/18

CS162 ©UCB Spring 2018

Lec 21.45

Port Numbers

- Port number: 16-bit number identifying the end-point of a transport connection
 - E.g., 80 identifies the port on which a processing implementing HTTP server can be connected



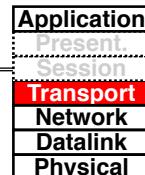
4/11/18

CS162 ©UCB Spring 2018

Lec 21.46

Internet Transport Protocols

- Datagram service (UDP)
 - No-frills extension of “best-effort” IP
 - Multiplexing/Demultiplexing among processes
- Reliable, in-order delivery (TCP)
 - Connection set-up & tear-down
 - Discarding corrupted packets (segments)
 - Retransmission of lost packets (segments)
 - Flow control
 - Congestion control
- Services **not available**
 - Delay and/or bandwidth guarantees
 - Sessions that survive change-of-IP-address



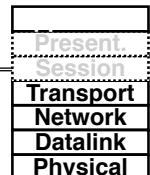
4/11/18

CS162 ©UCB Spring 2018

Lec 21.47

Application Layer (7 - not 5!)

- Service:** any service provided to the end user
- Interface:** depends on the application
- Protocol:** depends on the application
- Examples: Skype, SMTP (email), HTTP (Web), Halo, BitTorrent ...
- What happened to layers 5 & 6?
 - “Session” and “Presentation” layers
 - Part of **OSI** architecture, but not Internet architecture
 - Their functionality is provided by application layer

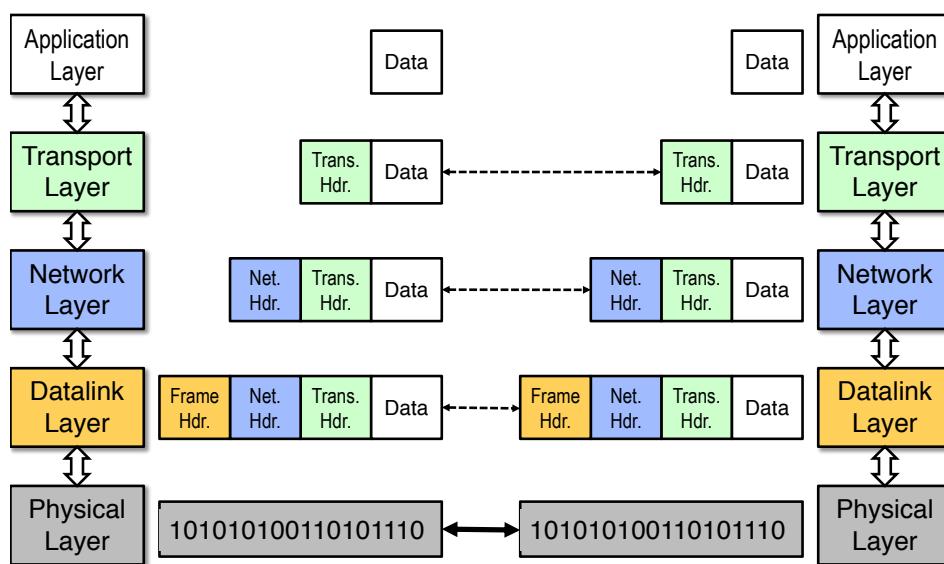


4/11/18

CS162 ©UCB Spring 2018

Lec 21.48

Application Layer (5)



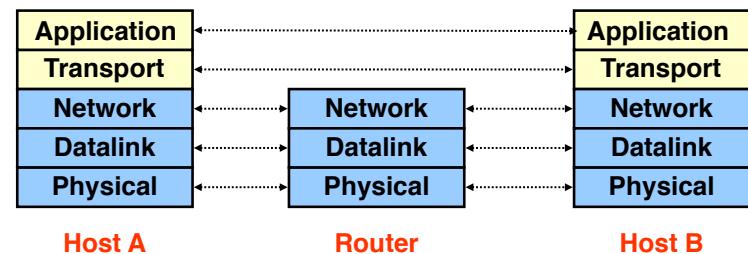
4/11/18

CS162 ©UCB Spring 2018

Lec 21.49

Five Layers Summary

- Lower three layers implemented everywhere
- Top two layers implemented only at hosts
- Logically, layers interact with peer's corresponding layer



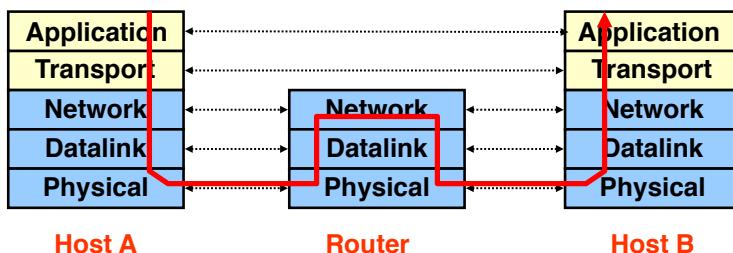
4/11/18

CS162 ©UCB Spring 2018

Lec 21.50

Physical Communication

- Communication goes down to physical network
- Then from network peer to peer
- Then up to relevant layer

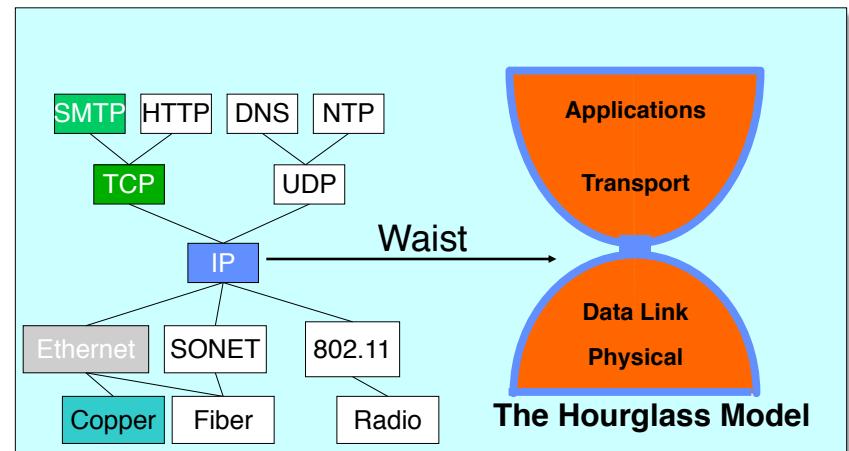


4/11/18

CS162 ©UCB Spring 2018

Lec 21.51

The Internet Hourglass



There is just **one** network-layer protocol, IP.
The “narrow waist” facilitates **interoperability**.

4/11/18

CS162 ©UCB Spring 2018

Lec 21.52

Implications of Hourglass

Single Internet-layer module (**IP**):

- Allows arbitrary networks to interoperate
 - Any network technology that supports IP can exchange packets
- Allows applications to function on all networks
 - Applications that can run on IP can **use any network**
- Supports simultaneous innovations above and below IP
 - But changing IP itself, i.e., **IPv6**, very involved

4/11/18

CS162 ©UCB Spring 2018

Lec 21.53

Placing Network Functionality

- Hugely influential paper: “End-to-End Arguments in System Design” by Saltzer, Reed, and Clark (‘84)
- “Sacred Text” of the Internet
 - Endless disputes about what it means
 - Everyone cites it as supporting their position

4/11/18

CS162 ©UCB Spring 2018

Lec 21.55

Drawbacks of Layering

- Layer N may duplicate layer N-1 functionality
 - E.g., error recovery to retransmit lost data
- Layers may need same information
 - E.g., timestamps, maximum transmission unit size
- Layering can hurt performance
 - E.g., hiding details about what is really going on
- Some layers are not always cleanly separated
 - Inter-layer dependencies for performance reasons
 - Some dependencies in standards (header checksums)
- Headers start to get really big
 - Sometimes header bytes >> actual content

4/11/18

CS162 ©UCB Spring 2018

Lec 21.54

Basic Observation

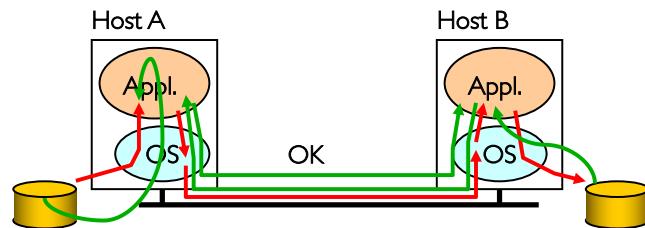
- Some types of network functionality can only be correctly implemented **end-to-end**
 - Reliability, security, etc
- Because of this, end hosts:
 - Can satisfy the requirement without network's help
 - Will/must do so, since can't **rely** on network's help
- Therefore **don't** go out of your way to implement them in the network

4/11/18

CS162 ©UCB Spring 2018

Lec 21.56

Example: Reliable File Transfer



- Solution 1: make each step reliable, and then **concatenate** them
- Solution 2: end-to-end **check** and try again if necessary

4/11/18

CS162 ©UCB Spring 2018

Lec 21.57

Discussion

- Solution 1 is **incomplete**
 - What happens if memory is corrupted?
 - Receiver has to do the check anyway!
- Solution 2 is **complete**
 - Full functionality can be entirely implemented at application layer with **no** need for reliability from lower layers
- *Is there any need to implement reliability at lower layers?*
 - Well, it could be **more efficient**

4/11/18

CS162 ©UCB Spring 2018

Lec 21.58

End-to-End Principle

Implementing this functionality in the network:

- Doesn't reduce host implementation complexity
- Does increase network complexity
- Probably imposes delay and overhead on all applications, **even if they don't need functionality**
- However, implementing in network **can** enhance performance in some cases
 - E.g., very lossy link

4/11/18

CS162 ©UCB Spring 2018

Lec 21.59

Conservative Interpretation of E2E

- Don't implement a function at the lower levels of the system unless it can be completely implemented at this level
- Unless you can relieve the burden from hosts, don't bother

4/11/18

CS162 ©UCB Spring 2018

Lec 21.60

Moderate Interpretation

- Think twice before implementing functionality in the network
- If hosts can implement functionality correctly, implement it in a lower layer **only** as a performance enhancement
- But do so only if it **does not impose burden** on applications that do not require that functionality
- This is the interpretation we are using

4/11/18

CS162 ©UCB Spring 2018

Lec 21.61

Summary (1/2)

- Layered architecture powerful abstraction for organizing complex networks
- Internet: 5 layers
 - Physical: send bits
 - Datalink: Connect two hosts on same physical media
 - Network: Connect two hosts in a wide area network
 - Transport: Connect two processes on (remote) hosts
 - Applications: Enable applications running on remote hosts to interact
- Unified Internet layering (Application/Transport/Internet/Link/Physical) decouples apps from networking technologies

4/11/18

CS162 ©UCB Spring 2018

Lec 21.62

Summary (2/2)

- E2E argument encourages us to keep IP simple
- If higher layer can implement functionality correctly, implement it in a lower layer **only** if
 - it improves the performance significantly for application that need that functionality, and
 - it **does not impose burden** on applications that do not require that functionality

4/11/18

CS162 ©UCB Spring 2018

Lec 21.63