

CS162

Operating Systems and Systems Programming

Lecture 23

Key-Value Stores, Security

April 18th, 2018

Profs. Anthony D. Joseph & Jonathan Ragan-Kelley
<http://cs162.eecs.Berkeley.edu>

Key Values: Examples

- Amazon:
 - Key: customerID
 - Value: customer profile (e.g., buying history, credit card, ..)
- Facebook, Twitter:
 - Key: UserID
 - Value: user profile (e.g., posting history, photos, friends, ...)
- iCloud/iTunes:
 - Key: Movie/song name
 - Value: Movie, Song



Key Value Storage

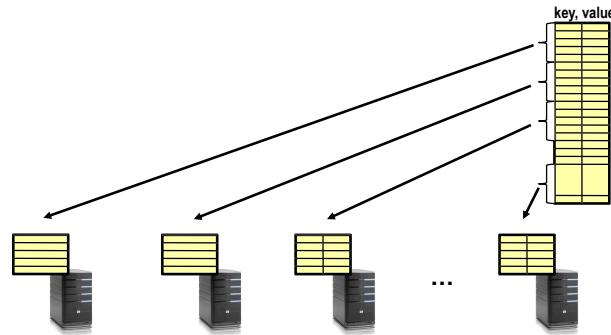
- Can handle huge volumes of data, e.g., Petabytes!
 - Store (key, value) tuples
- Simple interface
 - `put(key, value);` // insert/write “value” associated with “key”
 - `value = get(key);` // get/read data associated with “key”
- Used sometimes as a simpler but more scalable “database”

Key-Value Storage Systems in Real Life

- **Amazon**
 - DynamoDB: internal key value store used for Amazon.com (cart)
 - Simple Storage System (S3)
- **BigTable/HBase/Hypertable:** distributed, scalable data store
- **Cassandra:** “distributed data management system” (developed by Facebook)
- **Memcached:** in-memory key-value store for small chunks of arbitrary data (strings, objects)
- **BitTorrent distributed file location:** peer-to-peer sharing system
- ...

Key Value Store

- Also called Distributed Hash Tables (DHT)
- Main idea: partition set of key-values across many machines



4/18/18

CS162 ©UCB Spring 2018

Lec 23.5

Challenges



- **Fault Tolerance:** handle machine failures without losing data and without degradation in performance
- **Scalability:**
 - Need to scale to thousands of machines
 - Need to allow easy addition of new machines
- **Consistency:** maintain data consistency in face of node failures and message losses
- **Heterogeneity** (if deployed as peer-to-peer systems):
 - Latency: 1ms to 1000ms
 - Bandwidth: 32Gb/s to 100Mb/s

4/18/18

CS162 ©UCB Spring 2018

Lec 23.6

Key Questions

- **put(key, value):** where to store a new (key, value) tuple?
- **get(key):** where is the value associated with a given "key" stored?
- And, do the above while providing
 - Fault Tolerance
 - Scalability
 - Consistency

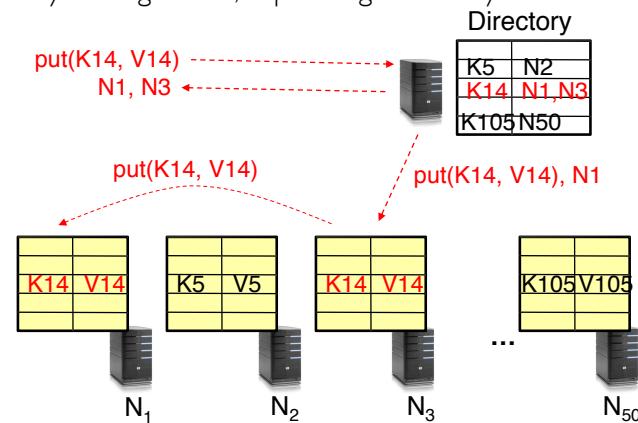
4/18/18

CS162 ©UCB Spring 2018

Lec 23.7

Directory-Based Architecture

- Directory maintains mapping between **keys** and **machines (nodes)** that store the **values** associated with the **keys**
- Replicate value on several nodes for fault-tolerance
- Scale by adding nodes, replicating Directory



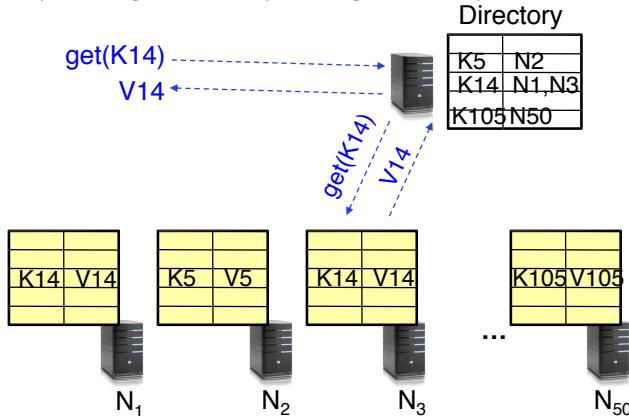
4/18/18

CS162 ©UCB Spring 2018

Lec 23.8

Directory-Based Architecture

- Directory maintains mapping between **keys** and **machines (nodes)** that store the **values** associated with the **keys**
- Replicate value on several nodes for fault-tolerance
- Scale by adding nodes, replicating Directory



4/18/18

CS162 ©UCB Spring 2018

Lec 23.9

Large Variety of Consistency Models

- Atomic consistency (linearizability): reads/writes (gets/puts) to replicas appear as if there was a single underlying replica (single system image)
 - Think “one updated at a time”
 - Transactions
- Eventual consistency: given enough time all updates will propagate through the system
 - One of the weakest form of consistency; used by many systems in practice
 - Must eventually converge on single value/key (coherence)
- And many others: causal consistency, sequential consistency, strong consistency, ...

4/18/18

CS162 ©UCB Spring 2018

Lec 23.11

Consistency

- Need to make sure that a value is replicated correctly
- How do you know a value has been replicated on every node?
 - Wait for acknowledgements from every node
- What happens if a node fails during replication?
 - Pick another node and try again
- What happens if a node is slow?
 - Slow down the entire put()? Pick another node?
- In general, with multiple replicas
 - Slow puts and fast gets

4/18/18

CS162 ©UCB Spring 2018

Lec 23.10

Quorum Consensus

- Improve put() and get() operation performance
- Define a replica set of size N
 - put() waits for acknowledgements from at least W replicas
 - get() waits for responses from at least R replicas
 - $W+R > N$
- Why does it work?
 - There is at least one node that contains the update
- Why might you use $W+R > N+1$?

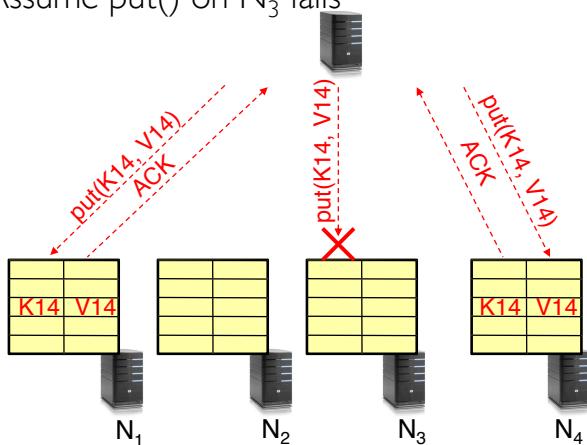
4/18/18

CS162 ©UCB Spring 2018

Lec 23.12

Quorum Consensus Example

- $N=3, W=2, R=2$
- Replica set for $K14$: $\{N_1, N_3, N_4\}$
- Assume $\text{put}()$ on N_3 fails



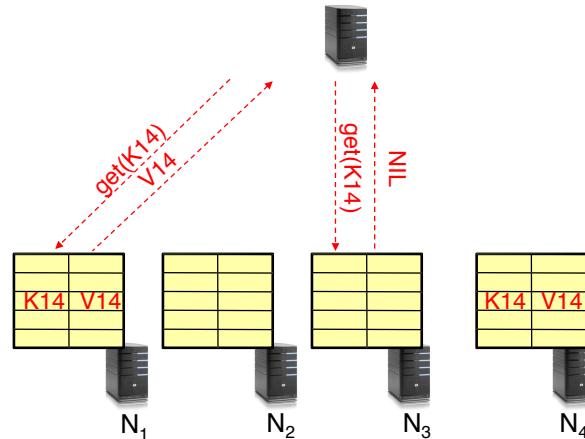
4/18/18

CS162 ©UCB Spring 2018

Lec 23.13

Quorum Consensus Example

- Now, issuing $\text{get}()$ to any two nodes out of three will return the answer



4/18/18

CS162 ©UCB Spring 2018

Lec 23.14

What is Computer Security Today?

- Computing in the presence of an adversary!
 - Adversary is the security field's defining characteristic
- Reliability, robustness, and fault tolerance
 - Dealing with Mother Nature (random failures)
- Security
 - Dealing with actions of a knowledgeable attacker dedicated to causing harm
 - Surviving malice, and not just mishance
- Wherever there is an adversary, there is a computer security problem!



CIMPILITY®
BlackEnergy
SCADA malware

Mirai IoT botnet

4/18/18

CS162 ©UCB Spring 2018

Lec 23.15

Protection vs. Security

- **Protection:** mechanisms for controlling access of programs, processes, or users to resources
 - Page table mechanism
 - Round-robin schedule
 - Data encryption
- **Security:** use of protection mech. to prevent misuse of resources
 - Misuse defined with respect to policy
 - » E.g.: prevent exposure of certain sensitive information
 - » E.g.: prevent unauthorized modification/deletion of data
 - Need to consider external environment the system operates in
 - » Most well-constructed system cannot protect information if user accidentally reveals password – social engineering challenge

4/18/18

CS162 ©UCB Spring 2018

Lec 23.16

Security Requirements

- Authentication
 - Ensures that a user is who is claiming to be
- Data integrity
 - Ensure that data is not changed from source to destination or after being written on a storage device
- Confidentiality
 - Ensures that data is read only by authorized users
- Non-repudiation
 - Sender/client can't later claim didn't send/write data
 - Receiver/server can't claim didn't receive/write data

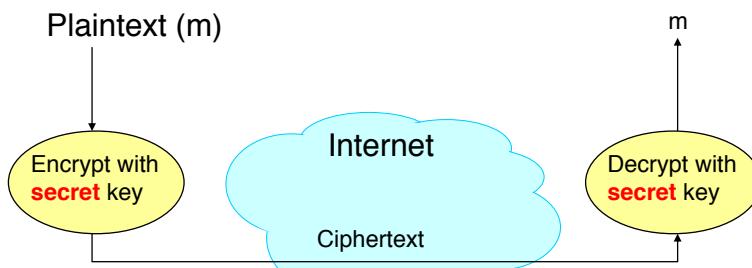
4/18/18

CS162 ©UCB Spring 2018

Lec 23.17

Using Symmetric Keys

- Same key for encryption and decryption
- Achieves confidentiality
- Vulnerable to tampering and replay attacks



4/18/18

CS162 ©UCB Spring 2018

Lec 23.19

Securing Communication: Cryptography

- Cryptography: communication in the presence of adversaries
- Studied for thousands of years
 - See the Simon Singh's **The Code Book** for an excellent, highly readable history
- Central goal: confidentiality
 - How to encode information so that an adversary can't extract it, but a friend can
- General premise: there is a key, possession of which allows decoding, but without which decoding is infeasible
 - Thus, key must be kept secret and not guessable

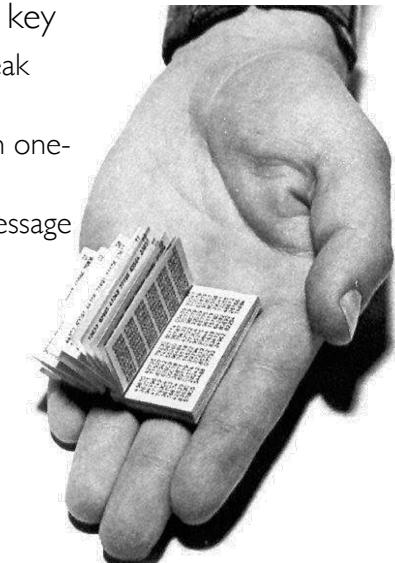
4/18/18

CS162 ©UCB Spring 2018

Lec 23.18

Symmetric Keys

- Can just XOR plaintext with the key
 - Easy to implement, but easy to break using frequency analysis
 - Unbreakable alternative: XOR with one-time pad
 - » Use a different key for each message



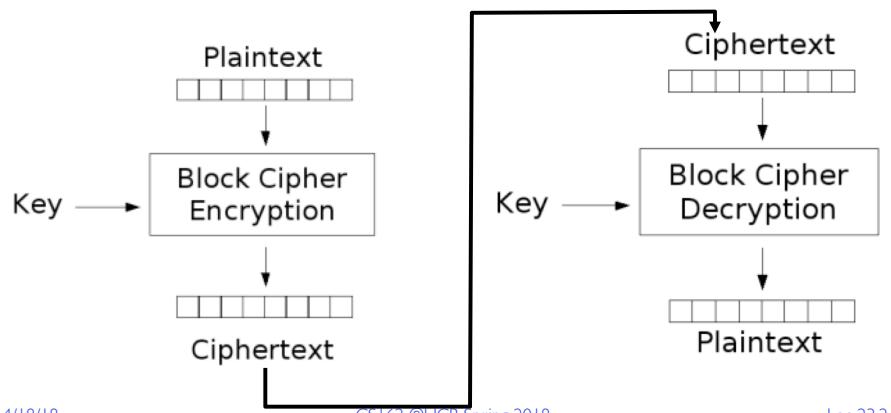
4/18/18

CS162 ©UCB Spring 2018

Lec 23.20

Block Ciphers with Symmetric Keys

- More sophisticated (e.g., block cipher) algorithms
 - Works with a block size (e.g., 64 bits)
- Can encrypt blocks separately:
 - Same plaintext \Rightarrow same ciphertext
- Much better:
 - Add in counter and/or link ciphertext of previous block



Symmetric Key Ciphers - DES & AES

- Data Encryption Standard (DES)
 - Developed by IBM in 1970s, standardized by NBS/NIST
 - 56-bit key (decreased from 64 bits at NSA's request)
 - Still fairly strong other than brute-forcing the key space
 - But custom hardware can crack a key in < 24 hours
 - Today many financial institutions use Triple DES
 - DES applied 3 times, with 3 keys totaling 168 bits
- Advanced Encryption Standard (AES)
 - Replacement for DES standardized in 2002
 - Key size: 128, 192 or 256 bits
- How fundamentally strong are they?
 - No one knows (no proofs exist)

4/18/18 CS162 ©UCB Spring 2018 Lec 23.22

Authentication in Distributed Systems

- What if identity must be established across network?

A cartoon illustration shows a woman with blonde hair and a man with brown hair. They are separated by a blue horizontal bar labeled 'Network'. The woman is holding a telephone receiver and the man is holding a coffee cup. The word 'PASS:ing.' is written vertically next to the man's arm.

 - Need way to prevent exposure of information while still proving identity to remote system
 - Many of the original UNIX tools sent passwords over the wire "in clear text"
 - E.g.: telnet, ftp, yp (yellow pages, for distributed login)
 - Result: Snooping programs widespread
- What do we need? Cannot rely on physical security!
 - Encryption: Privacy, restrict receivers
 - Authentication: Remote Authenticity, restrict senders

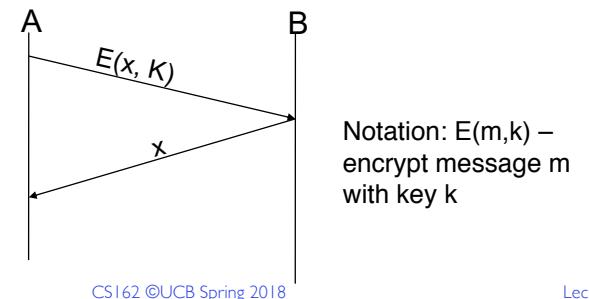
4/18/18

CS162 ©UCB Spring 2018

Lec 23.23

Authentication via Secret Key

- Main idea: entity proves identity by decrypting a secret encrypted with its own key
 - K – secret key shared only by A and B
- A can ask B to authenticate itself by decrypting a nonce, i.e., random value, x
 - Avoid replay attacks (attacker impersonating client or server)
- Vulnerable to man-in-the-middle attack



Administrivia

- Midterm 3 Review: **Fri 4/20 7:00-10:00PM (VLSB 2050)**
- Midterm 3 coming up on **Wed 4/25 6:30-8PM**
 - All topics up to and including Lecture 23
 - » Focus will be on Lectures 17 – 23 and associated readings, and Projects 3
 - » But expect 20-30% questions from materials from Lectures 1-16
 - LKS 245, Hearst Field Annex A1, VLSB 2060, Barrows 20, Wurster 102 (**see Piazza for your room assignment**)
 - Closed book
 - 2 pages hand-written notes both sides

4/18/18

CS162 ©UCB Spring 2018

Lec 23.25

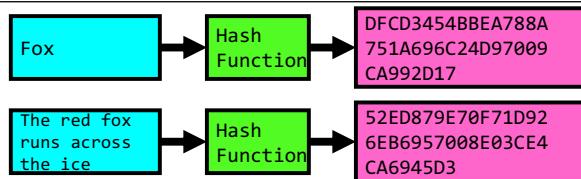
BREAK

4/18/18

CS162 ©UCB Spring 2018

Lec 23.26

Secure Hash Function



- Hash Function: Short summary of data (message)
 - For instance, $h_1 = H(M_1)$ is the hash of message M_1
 - » h_1 fixed length, despite size of message M_1
 - » Often, h_1 is called the “digest” of M_1
- Hash function H is considered secure if
 - It is infeasible to find M_2 with $h_1 = H(M_2)$; i.e., can't easily find other message with same digest as given message
 - It is infeasible to locate two messages, m_1 and m_2 , which “collide”, i.e. for which $H(m_1) = H(m_2)$
 - A small change in a message changes many bits of digest/can't tell anything about message given its hash

4/18/18

CS162 ©UCB Spring 2018

Lec 23.27

Integrity: Cryptographic Hashes

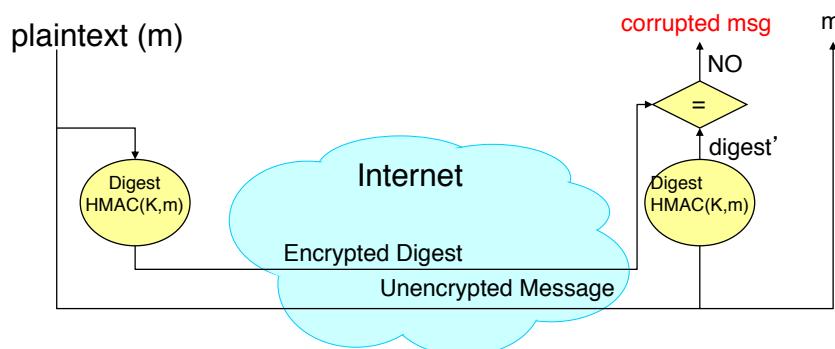
- Basic building block for integrity: cryptographic hashing
 - Associate hash with byte-stream, receiver verifies match
 - » Assures data hasn't been modified, either accidentally – or maliciously
- Approach:
 - Sender computes a secure digest of message m using $H(x)$
 - » $H(x)$ is a publicly known hash function
 - » Digest $d = \text{HMAC}(K, m) = H(K \mid H(K \mid m))$
 - » $\text{HMAC}(K, m)$ is a hash-based message authentication function
 - Send digest d and message m to receiver
 - Upon receiving m and d , receiver uses shared secret key, K , to recompute $\text{HMAC}(K, m)$ and see whether result agrees with d

4/18/18

CS162 ©UCB Spring 2018

Lec 23.28

Using Hashing for Integrity



Can encrypt m for confidentiality

4/18/18

CS162 ©UCB Spring 2018

Lec 23.29

Key Distribution

- How do you get shared secret to both places?
 - For instance: how do you send authenticated, secret mail to someone who you have never met?
 - Must negotiate key over private channel
 - » Exchange code book/key cards/memory stick/others
- Could use a third party

4/18/18

CS162 ©UCB Spring 2018

Lec 23.31

Standard Cryptographic Hash Functions

- MD5 (Message Digest version 5)
 - Developed in 1991 (Rivest), produces 128 bit hashes
 - Widely used (RFC 1321)
 - Broken (1996-2008): attacks that find collisions
- SHA-1 (Secure Hash Algorithm)
 - Developed in 1995 (NSA) as MD5 successor with 160 bit hashes
 - Widely used (SSL/TLS, SSH, PGP, IPSEC)
 - Broken in 2005, government use discontinued in 2010
- SHA-2 (2001)
 - Family of SHA-224, SHA-256, SHA-384, SHA-512 functions
- HMAC's are secure even with older "insecure" hash functions

4/18/18

CS162 ©UCB Spring 2018

Lec 23.30

Third Party: Authentication Server (Kerberos)

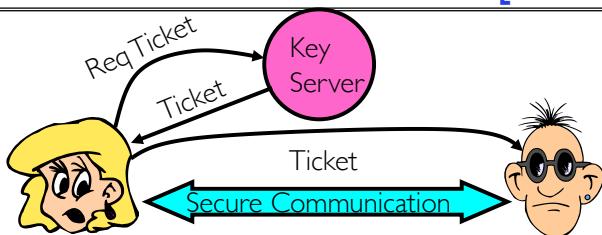
- Notation:
 - K_{xy} is key for talking between x and y
 - $(\dots)^K$ means encrypt message (...) with the key K
 - Clients: A and B, Authentication server S
- Usage:
 - A asks server for key:
 - » $A \rightarrow S$: [Hi! I'd like a key for talking between A and B]
 - » Not encrypted. Others can find out if A and B are talking
 - Server returns session key encrypted using B's key
 - » $S \rightarrow A$: Message [Use K_{ab} (This is A! Use $K_{ab}^{K_{sb}}$)] $^{K_{sa}}$
 - » This allows A to know, "S said use this key"
 - Whenever A wants to talk with B
 - » $A \rightarrow B$: Ticket [This is A! Use $K_{ab}^{K_{sb}}$] $^{K_{sa}}$
 - » Now, B knows that K_{ab} is sanctioned by S

4/18/18

CS162 ©UCB Spring 2018

Lec 23.32

Authentication Server Continued [Kerberos]



- Details
 - Both A and B use passwords (shared with key server) to decrypt return from key servers
 - Add in timestamps to limit how long tickets will be used to prevent attacker from replaying messages later
 - Also have to include encrypted checksums (hashed version of message) to prevent malicious user from inserting things into messages/changing messages
- *Kerberos is the system that holds all CalNet IDs and passphrases*

4/18/18

CS162 ©UCB Spring 2018

Lec 23.33

Asymmetric Encryption (Public Key)

- Idea: use two different keys, one to encrypt (e) and one to decrypt (d)
 - A **key pair**
- Crucial property: knowing e does not give away d
- Therefore e can be public: everyone knows it!
- If Alice wants to send to Bob, she fetches Bob's public key (say from Bob's home page) and encrypts with it
 - Alice can't decrypt what she's sending to Bob ...
 - ... but then, neither can anyone else (except Bob)

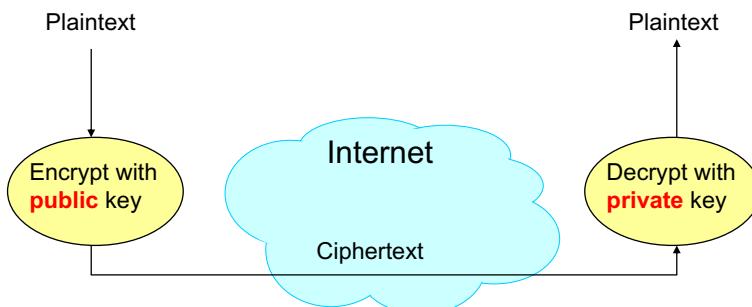
4/18/18

CS162 ©UCB Spring 2018

Lec 23.34

Public Key / Asymmetric Encryption

- Sender uses receiver's **public key**
 - Advertised to everyone
- Receiver uses complementary **private key**
 - Must be kept secret



4/18/18

CS162 ©UCB Spring 2018

Lec 23.35

Public Key Encryption Details

- Idea: K_{public} can be made public, keep K_{private} private
 - Insecure Channel
- Gives message privacy (restricted receiver):
 - Public keys (secure destination points) can be acquired by anyone/used by anyone
 - Only person with private key can decrypt message
- What about authentication?
 - Use combination of private and public key
 - $\text{Alice} \rightarrow \text{Bob}: [(I'm \text{ Alice})^{A_{\text{private}}} \text{ Rest of message}]^{B_{\text{public}}}$
 - Provides restricted sender and receiver
- But: how does Alice know that it was Bob who sent her B_{public} ? And vice versa...

4/18/18

CS162 ©UCB Spring 2018

Lec 23.36

Public Key Cryptography

- Invented in the 1970s
 - Revolutionized cryptography
 - (Was actually invented earlier by British intelligence)
- How can we construct an encryption/decryption algorithm using a key pair with the public/private properties?
 - Answer: Number Theory
- Most fully developed approach: RSA
 - Rivest / Shamir / Adleman, 1977; RFC 3447
 - Based on modular multiplication of very large integers
 - Very widely used (e.g., ssh, SSL/TLS for https)
- Also mature approach: Elliptic Curve Cryptography (ECC)
 - Based on curves in a Galois-field space
 - Shorter keys and signatures than RSA

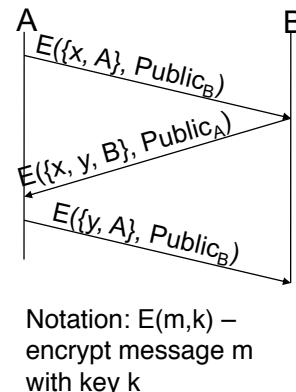
4/18/18

CS162 ©UCB Spring 2018

Lec 23.37

Simple Public Key Authentication

- Each side need only to know the other side's public key
 - No secret key need be shared
- A encrypts a nonce (random num.) x
 - Avoid **replay attacks**, e.g., attacker impersonating client or server
- B proves it can recover x , generates second nonce y
- A can authenticate itself to B in the same way
- **A and B have shared private secrets on which to build private key!**
 - We just did secure key distribution!
- Many more details to make this work securely in practice!



4/18/18

CS162 ©UCB Spring 2018

Lec 23.39

Properties of RSA

- Requires generating large, random prime numbers
 - Algorithms exist for quickly finding these (probabilistic!)
- Requires exponentiation of very large numbers
 - Again, fairly fast algorithms exist
- Overall, much slower than symmetric key crypto
 - **One general strategy: use public key crypto to exchange a (short) symmetric session key**
 - » Use that key then with AES or such
- How difficult is recovering d , the private key?
 - Equivalent to finding prime factors of a large number
 - » Many have tried - believed to be very hard
(= brute force only)
 - » (Though quantum computers could do so in polynomial time!)

4/18/18

CS162 ©UCB Spring 2018

Lec 23.38

Non-Repudiation: RSA Crypto & Signatures

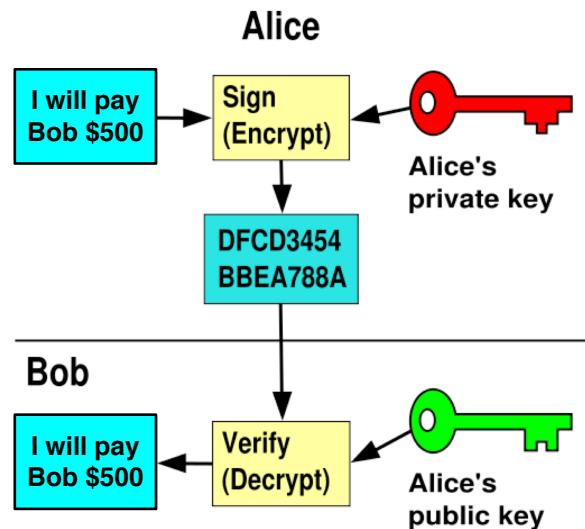
- Suppose Alice has published public key KE
- If she wishes to prove who she is, she can send a message x encrypted with her private key KD (i.e., she sends $E(x, KD)$)
 - Anyone knowing Alice's public key KE can recover x , verify that Alice must have sent the message
 - » It provides a **signature**
 - Alice can't deny it: **non-repudiation**
- Could simply encrypt a hash of the data to sign a document that you wanted to be in clear text
- Note that either of these signature techniques work perfectly well with any data (not just messages)
 - Could sign every datum in a database, for instance

4/18/18

CS162 ©UCB Spring 2018

Lec 23.40

RSA Crypto & Signatures (cont'd)



4/18/18

CS162 ©UCB Spring 2018

Lec 23.41

Digital Certificates

- How do you know K_E is Alice's public key?
 - Trusted authority (e.g., Verisign) signs binding between Alice and K_E with its private key KV_{private}
 - $C = E(\{Alice, K_E\}, KV_{\text{private}})$
 - C : digital certificate
- Alice: distribute her digital certificate, C
- Anyone: use trusted authority's KV_{public} to extract Alice's public key from C
 - $D(C, KV_{\text{public}}) = D(E(\{Alice, K_E\}, KV_{\text{private}}), KV_{\text{public}}) = \{Alice, K_E\}$

4/18/18

CS162 ©UCB Spring 2018

Lec 23.42

Summary of Our Crypto Toolkit

- If we can securely distribute a key, then
 - Symmetric ciphers (e.g., AES) offer fast, presumably strong confidentiality
- Public key cryptography does away with (potentially major) problem of secure key distribution
 - But: not as computationally efficient
 - » Often addressed by using public key crypto to exchange a **session key**
- Digital signature binds the public key to an entity

4/18/18

CS162 ©UCB Spring 2018

Lec 23.43

Putting It All Together - HTTPS

- What happens when you click on <https://www.amazon.com>?
 - https = “Use HTTP over SSL/TLS”
 - SSL = Secure Socket Layer
 - TLS = Transport Layer Security
 - » Successor to SSL
 - Provides security layer (authentication, encryption) on top of TCP
 - » Fairly transparent to applications

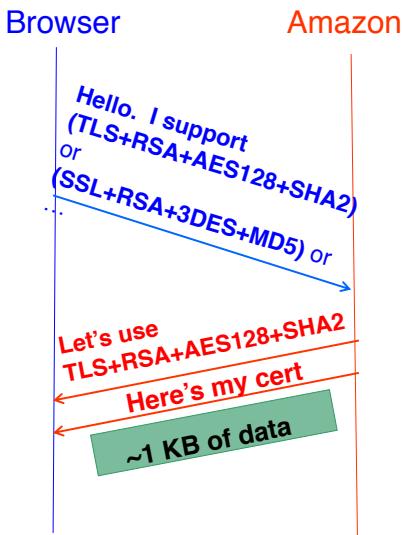
4/18/18

CS162 ©UCB Spring 2018

Lec 23.44

HTTPS Connection (SSL/TLS) (cont'd)

- Browser (client) connects via TCP to Amazon's HTTPS server
- Client sends over list of crypto protocols it supports
- Server picks protocols to use for this session
- Server sends over its certificate
- (all of this is in the clear)



4/18/18

CS162 ©UCB Spring 2018

Lec 23.45

Inside the Server's Certificate

- Name associated with cert (e.g., Amazon)
- Amazon's RSA public key
- A bunch of auxiliary info (physical address, type of cert, expiration time)
- Name of certificate's signatory (who signed it)
- A public-key signature of a hash (**SHA-256**) of all this
 - Constructed using the signatory's private RSA key, i.e.,
 - Cert = $E(H_{SHA256}(KA_{public}, \text{www.amazon.com}, \dots), KS_{private}))$
 - » KA_{public} : Amazon's public key
 - » $KS_{private}$: signatory (certificate authority) private key
- ...

4/18/18

CS162 ©UCB Spring 2018

Lec 23.46

Validating Amazon's Identity

- How does the browser authenticate certificate signatory?
 - Certificates of several certificate authorities (e.g., Verisign) are **hardwired into the browser (or OS)**
- If can't find cert, warn user that site has not been verified
 - And may ask whether to continue
 - Note, can still proceed, just **without authentication**
- Browser uses public key in signatory's cert to decrypt signature
 - Compares with its own **SHA-256** hash of Amazon's cert
- Assuming signature matches, now have high confidence it's indeed Amazon ... assuming signatory is trustworthy
 - DigiNotar CA breach (July-Sept 2011): Google, Yahoo!, Mozilla, Tor project, Wordpress, ... (**531 total certificates**)

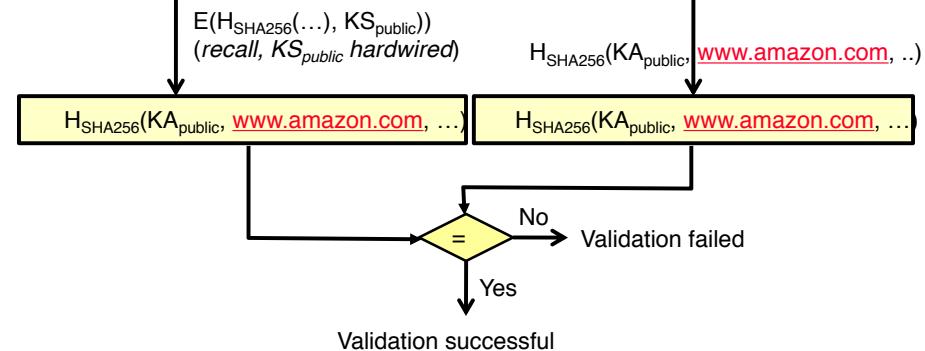
4/18/18

CS162 ©UCB Spring 2018

Lec 23.47

Certificate Validation

Certificate

 $E(H_{SHA256}(KA_{public}, \text{www.amazon.com}, \dots), KS_{private}),$
 $KA_{public}, \text{www.amazon.com}, \dots$ 

Can also validate using peer approach: <https://www.eff.org/observatory>

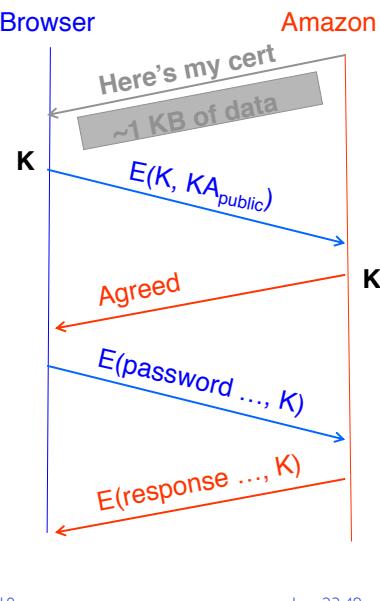
4/18/18

CS162 ©UCB Spring 2018

Lec 23.48

HTTPS Connection (SSL/TLS) cont'd

- Browser constructs a random session key K used for data communication
 - Private key for bulk crypto
- Browser encrypts K using Amazon's public key
- Browser sends $E(K, KA_{public})$ to server
- Browser displays 
- All subsequent comm. encrypted w/ symmetric cipher (e.g., AES128) using key K
 - Eg, client can authenticate using a password



4/18/18

CS162 ©UCB Spring 2018

Lec 23.49

Summary

- Key-Value Store:
 - Two operations
 - » put(key, value)
 - » value = get(key)
 - Challenges
 - » Fault Tolerance → replication
 - » Scalability → serve get()'s in parallel; replicate/cache hot tuples
 - » Consistency → quorum consensus to improve put() performance
- Security:
 - Many more challenges to building secure systems and applications
 - No fixed-point solutions
 - Adversaries constantly change and adapt
 - Defenses must also constantly change and adapt
 - Take CS 161 !

4/18/18

CS162 ©UCB Spring 2018

Lec 23.50