

CS162

Operating Systems and Systems Programming

Lecture 24

Research Talks

April 23rd, 2018

Profs. Anthony D. Joseph & Jonathan Ragan-Kelley
<http://cs162.eecs.Berkeley.edu>

Anthony's Research Projects

- SecML – Secure Machine Learning in the presence of adversaries
- Big Data Genomics – Algorithms and systems programming for population-scale genomics analyses

4/23/18

CS162 ©UCB Spring 2018

2

Adversarial Exploitation of ML

- Machine Learning is everywhere
 - Search rankings, spam detection, financial models, self driving cars, ...
- Adversaries everywhere! Traditional approach – *Evading Adversary*
 - Attacker determines decision boundary
 - Crafts (positive instance) content that is classified as negative
- Input data “drifts” over time, so must periodically retrain ML
 - Use previously classified input data
- Newer approach – *Influencing Adversary*
 - Patient attacker operates during periodic retraining stage by injecting “tricky” positive instances
 - Shifts decision boundary over time during retraining such that (positive instance) content is eventually classified as negative
- We need novel adaptive, robust ML techniques to defend against Influencing Adversaries
 - Also Active Learning with Humans-in-the-Loop

4/23/18

CS162 ©UCB Spring 2018

3

Our Current Domain and Dataset

- Windows malware detection
 - Ever changing, evading adversary with novel attacks
- Over 1 million unique executables from VirusTotal
 - Seen over 5 million times
 - Spanning January 2012 to June 2014
 - ~1% of VirusTotal's executables for the period
- Labeled by 32 AV providers' static engines
 - Threshold of four detections to label as malicious
- 85% malicious (3,000 to 406,000 families)
- **The largest academic malware detection dataset**

4/23/18

CS162 ©UCB Spring 2018

4

An ML Analysis Pipeline: From Labs to Production

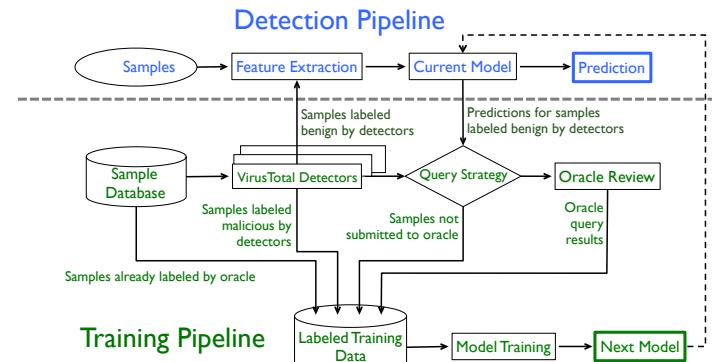
- Goal: research AND production environment
 - *Moonshot: Detect undetectable Advanced Persistent Threats*
- Significant challenges to deploying ML in practice:
 - Increased human insight – *how the system is making decisions*
 - Scalable, interactive analysis – *large-scale data, human-driven exploration*
 - Time series analysis – *repeatable scientific experiments to build confidence in ML*
 - Resource management – *variable cost analyses and experts*
 - Adversarial resilience – *detecting attempts to manipulate or evade detection*

4/23/18

CS162 ©UCB Spring 2018

5

Secure Active Learning Testbed (SALT)

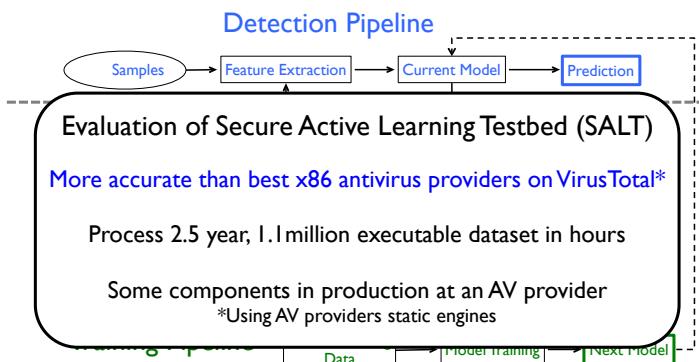


Built on Apache Spark for scalability

CS162 ©UCB Spring 2018

6

Secure Active Learning Testbed (SALT)

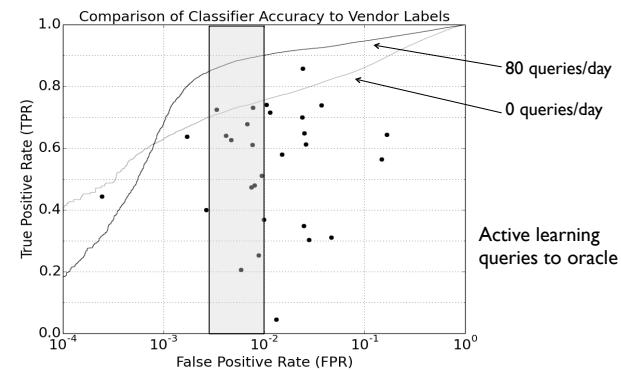


4/23/18

CS162 ©UCB Spring 2018

7

Active Learning Results

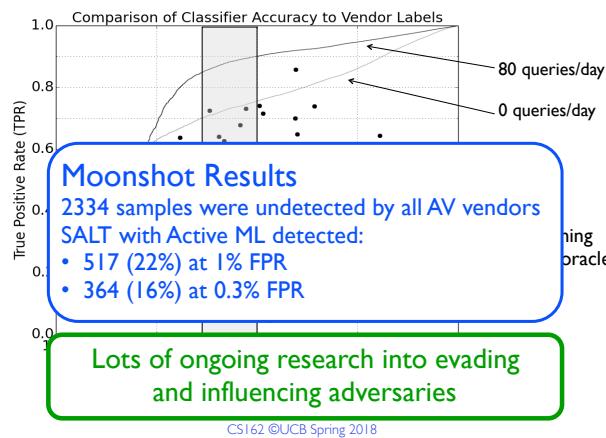


4/23/18

CS162 ©UCB Spring 2018

8

Active Learning Results



4/23/18

CS162 ©UCB Spring 2018

9

SecML Acknowledgements

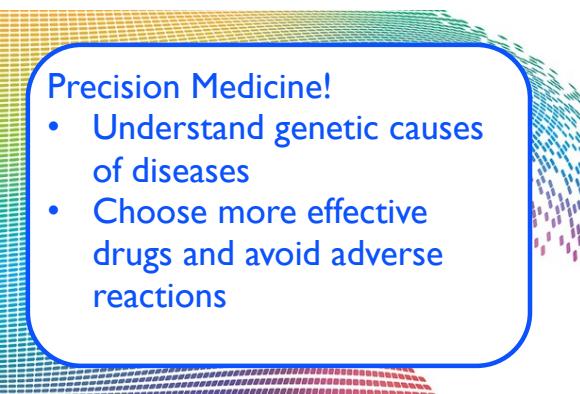
- **UC Berkeley:** David Fifield, Alex Kantchelian, Brad Miller, Qi Zhong, Vaishaal Shankar, Tony Wu, George Yu, Riyaz Faizullabhoj, Doug Tygar, Anthony D. Joseph
- **Drexel:** Eleanor Cawthon, Edwin Dauber, Rachel Greenstadt
- **ICSI:** Sadia Afroz, Michael Tshantz
- **Intel:** Ling Huang, Rekha Bachwani, Catherine Huang
- **Mitsubishi:** Takumi Yamamoto

4/23/18

CS162 ©UCB Spring 2018

10

Why Do We Need Big Genomic Data?



4/23/18

CS162 ©UCB Spring 2018

11

A Precision Medicine First

- Joshua Osborn (14 years old)
 - Visited ER 3 times in 4 months with encephalitis (brain swelling)
 - Hospitalized after 3rd visit and placed into a medically induced coma due to uncontrollable seizures
- Doctors tried traditional medicine approaches
 - Over 100 viral/fungal/bacterial pathogens cause encephalitis
 - Lots and lots of tests for different pathogens
 - 1 cm³ of brain tissue biopsied
- Last resort – sent spinal fluid to UCSF for DNA sequencing
 - Needle in a haystack – throw out human reads (99.984%)
 - 475 reads out of 3M from rare *Leptospira santarosae* bacteria → given antibiotics, out of coma in week, discharged 4 wks



4/23/18

CS162 ©UCB Spring 2018

12



Abstract
More than half of laboratory tests establishing a diagnosis can be difficult to interpret, especially for rare diseases, largely because bacterial causes contributed often to outcome.

CASE REPORT

A 14-year-old boy with severe combined immunodeficiency (SCID) caused by adenosine deaminase deficiency and partial immune reconstitution after he had undergone two hematopoietic bone marrow transplants developed the emergency department in April 2013 after having had headache and fevers, with temperatures up to 39.4°C, for 6 days (Figure 1A). He was admitted to the hospital and discharged 1 day later after resolution of his fever and headache.

The patient's immunodeficiencies included mostly infections of flaviviruses, immune globulin for hypogammaglobulinemia and immunophenotype, sulfamethoxazole or aztreonam for prophylaxis against *Pseudomonas aeruginosa*. He had no known colds but did have three pneumonia cases. He had gone on a missionary trip to Puerto Rico during the first 2 weeks of August 2012 (Figure 1A) where he became ill in a resort town in the mountains.

<https://amplab.cs.berkeley.edu/snap-helps-save-a-life/>

The New York Times

In First, Test of DNA Finds Root of Illness

BY CARL ZIMMER JUNE 4, 2014

Joshua Osborn, 14, lay in a coma at American Family Children's Hospital in Madison, Wisc. For weeks his brain had been swelling with fluid, and a battery of tests had failed to reveal why.

The doctors told his parents, Clark and Julie, that they wanted to run one more test with an experimental new technology. Scientists would search Joshua's cerebrospinal fluid for pieces of DNA. Some of them might belong to the pathogen causing his encephalitis.

The Osborns agreed, although they were skeptical that the test would succeed where so many others had failed.

It did. The test found a tiny amount of DNA from a bacterium called *Neuroleptospiro*.

That bacterium had never before been linked to encephalitis.

It was a breakthrough, the first time that a genetic test had solved a mystery diagnosis.

Joshua's case shows how genetic sequencing is revolutionizing medicine.

It also shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

Joshua's case shows how far we still have to go.

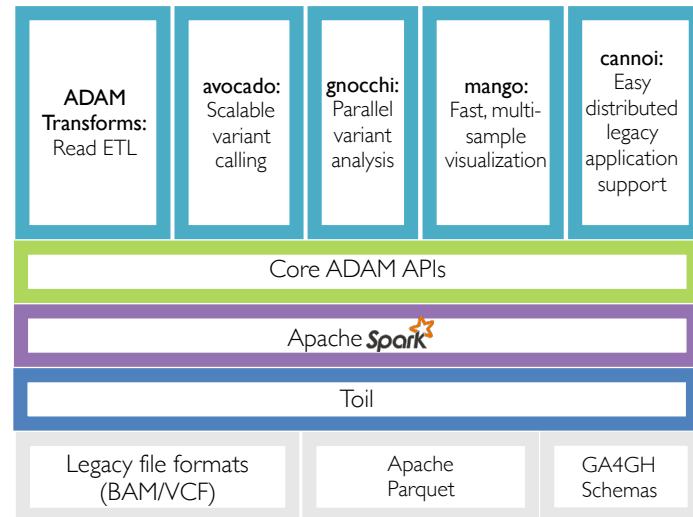
Cloud-scale Genomic Analyses

- Going from raw genomic reads (100+ GB) to aligned genomes (3+ GB) to the differences between us (1GB) to actionable insights
- Projects: [Toil](#), [ADAM](#), [avocado](#), [mango](#), [gnocchi](#), [cannoli](#)
- Many algorithmic challenges:
 - Complex preprocessing and filtering methods
 - ML algorithms at genome- and population-scale
- Many infrastructure challenges:
 - Volume of data: VA MVP, NIH AoU Project
 - How do we run geo-distributed workflows efficiently?
 - Can bioinformatics be made reproducible?

4/23/18

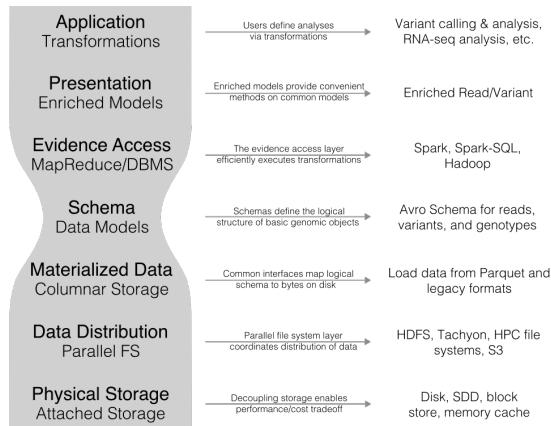
CS162 ©UCB Spring 2018

17



ADAM: A Stack-based Design

A Well Designed Stack Simplifies Application Design



4/23/18

CS162 ©UCB Spring 2018

19

Define a Schema, Query Genomic Data via Schema

```
record AlignmentRecord {
  union { int, long } contig = null;
  union { null, long } start = null;
  union { null, long } end = null;
  union { null, int } mapq = null;
  union { null, string } readname = null;
  union { null, string } referenceName = null;
  union { null, string } noteReference = null;
  union { null, long } noteAlignmentStart = null;
  union { null, string } cigar = null;
  union { null, string } strand = null;
  union { null, string } recordgroupName = null;
  union { int, null } basesTrimmedFromStart = 0;
  union { boolean, null } basesTrimmedFromEnd = 0;
  union { boolean, null } isProperPair = false;
  union { boolean, null } readMapped = false;
  union { boolean, null } mateMapped = false;
  union { boolean, null } firstPair = false;
  union { boolean, null } secondPair = false;
  union { boolean, null } isSecondary = false;
  union { boolean, null } failedVendorQualityChecks = false;
  union { boolean, null } duplicatedRead = false;
  union { boolean, null } isPaired = false;
  union { boolean, null } noteNegativeStrand = false;
  union { boolean, null } primaryAlignment = false;
  union { boolean, null } secondaryAlignment = false;
  union { boolean, null } isplementaryAlignment = false;
  union { null, string } mismatchingPositions = null;
  union { null, string } original = null;
  union { null, string } otherNames = null;
  union { null, string } recordGroupSequencingCenter = null;
  union { null, string } recordGroupDescription = null;
  union { null, long } recordGroupEndEpoch = null;
  union { null, string } recordGroupForward = null;
  union { null, string } recordGroupReverse = null;
  union { null, string } recordGroupLibrary = null;
  union { null, int } recordGroupPredictedMedianInsertSize = null;
  union { null, string } recordGroupPlatform = null;
  union { null, string } recordGroupPlatformId = null;
  union { null, string } recordGroupSample = null;
  union { null, Contig } mateContig = null;
}
```

Logical data representation

Schema Data Models

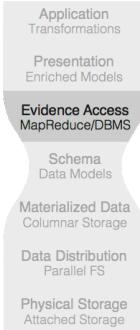
4/23/18

CS162 ©UCB Spring 2018

20

Parallelize Common Access Patterns

Efficient parallel access methods



- We provide parallel methods for processing genomic data with Spark
- Makes use of genomics specific reference genome structure to accelerate queries
- E.g., region join for high performance overlap queries

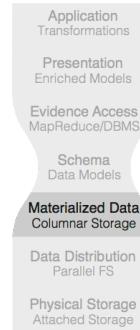
4/23/18

CS162 ©UCB Spring 2018

21

Pick Appropriate Storage

Efficient physical storage formats



- When accessing scientific datasets, we frequently slice and dice the dataset:
 - Algorithms may touch subsets of columns
 - We don't always touch the whole dataset
- This is a good match for columnar storage like Apache Parquet
- Also, can support legacy formats

4/23/18

CS162 ©UCB Spring 2018

22

Want to Learn More?

- Check out the code!
 - ADAM: <https://github.com/bigdatagenomics/adam>
- Check out a demo!
 - <https://databricks.com/blog/2016/05/24/genome-sequencing-in-a-nutshell.html>
- Run ADAM in Databricks Community Edition!
 - <http://goo.gl/xK8x7s>

4/23/18

CS162 ©UCB Spring 2018

23

Acknowledgements

- **UC Berkeley:** Matt Massie, Timothy Danford, André Schumacher, Jey Kottalam, Karen Feng, Eric Tu, Niranjan Kumar, Ananth Pallaseni, Michael Heuer, Alyssa Morrow, Frank A. Nothaft, Devin Petersohn, Gunjan Baid, Justin Paschall, Anthony D. Joseph, Dave Patterson
- **Mt. Sinai:** Arun Ahuja, Neal Sidhwani, Ryan Williams, Michael Linderman, Jeff Hammerbacher
- **GenomeBridge:** Carl Yeksigian
- **Cloudera:** Uri Laserson, Tom White
- **Microsoft Research:** Ravi Pandya, Bill Bolosky
- **UC Santa Cruz:** Benedict Paten, David Haussler, Hannes Schmidt, Beau Norgeot, Audrey Musselman-Brown, John Vivian, Jacob Pfeil
- And many other open source contributors, especially Neil Ferguson, Andy Petrella, Xavier Tordoir
- Total of >60 contributors to ADAM/BDG from >12 institutions

4/23/18

CS162 ©UCB Spring 2018

24

Administrivia

- Midterm 3 coming up on **Wed 4/25 6:30-8PM**
 - All topics up to and including Lecture 23
 - » Focus will be on Lectures 17 – 23 and associated readings, and Projects 3
 - » But expect 20-30% questions from materials from Lectures 1-16
 - LKS 245, Hearst Field Annex A1, VLSB 2060, Barrows 20, Wurster 102 ([see Piazza for your room assignment](#))
 - Closed book
 - 2 pages hand-written notes both sides
- Please fill out the online course evaluation

4/23/18

CS162 ©UCB Spring 2018

25

BREAK

4/23/18

CS162 ©UCB Spring 2018

26

Systems & Compilers for Graphics and Visual Computing

Jonathan Ragan-Kelley



Visual computing demands orders of magnitude more performance

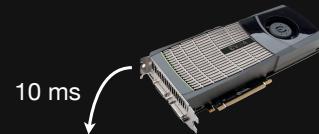
Rendering: insatiable demand for computation

Modern game

2 Mpixels

1 Mpolys

10 ms/frame



10 ms

5 hrs



Tintin, Avatar

8 Mpixels

5 Gpolys

5 hrs/frame



6 orders of magnitude
more computation

images by Valve, Weta

The biggest data is visual

YouTube: 400 hrs uploaded / min

[Brewer et al. 2016]

1.5 Terapixels/sec

250 M surveillance cameras,
2.5 B cell phone cameras, ...



Pervasive sensing: “the cloud” is not enough data transfer >> capture

Sensor + Read out

5 Mpixels

~1 mJ/frame



LTE radio

50 Mbit/sec

1 W

~1 J/frame

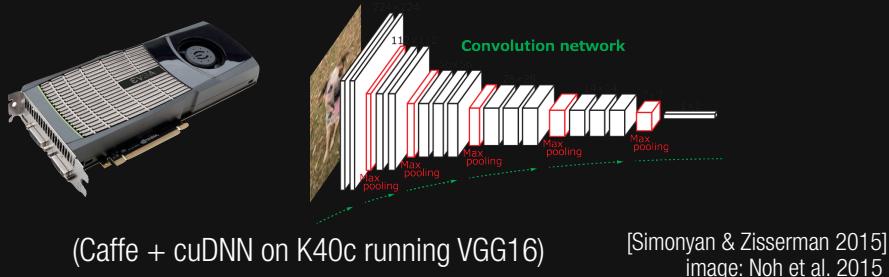


[Wu et al. 2012]

transmission power costs 1,000x capture

Visual data analysis is expensive

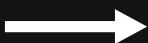
One object recognition neural net:
250 Watt GPU → 0.05 megapixels at video rate



Your data-intensive problem here...

Programmer productivity has exploded

1990s



C/C++



2010s

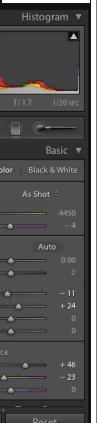


Writing high-performance code is hard

Reference:

300 lines C++

Adobe: 1500 lines
3 months of work
10x faster



My group's research:

Compilers, systems, architectures, and algorithms
for high-performance graphics & visual computing.



- applications
- algorithms
- data structures
- languages
- compilers
- hardware

Reorganize computations & data.

- Simpler programs
- Order of magnitude faster
- Scalable on future architectures

Communication dominates computation in both energy and time

Operation (32-bit operands)	Energy/Op (28 nm)	Cost (vs. ALU)
ALU op	1 pJ	-
Load from SRAM	5 pJ	5x
Move 10mm on-chip	32 pJ	32x
Send off-chip	500 pJ	500x
Send to DRAM	1 nJ	1,000x
Send over LTE	> 50 µJ	50,000,000x

data from John Brunhaver, Bill Dally, Mark Horowitz

How can we increase performance and efficiency?

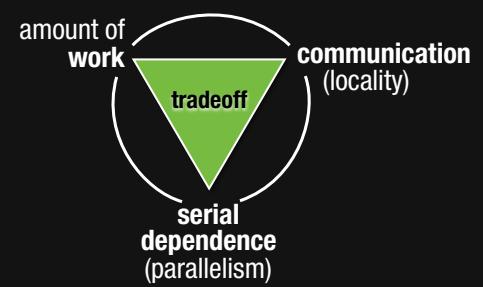
Parallelism

“Moore’s law” scaling requires
exponentially more parallelism.

Locality

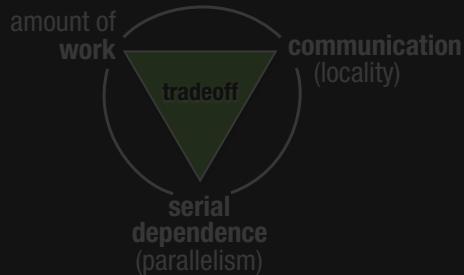
Data should move as little as possible.

Message #1: Performance requires complex tradeoffs



Where does performance come from?

Program
Hardware

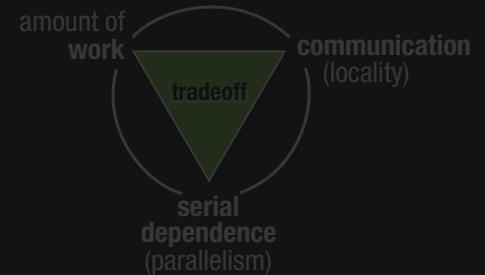


Message #2: organization of computation is a first-class issue

Program:

Algorithm
Organization of computation

Hardware

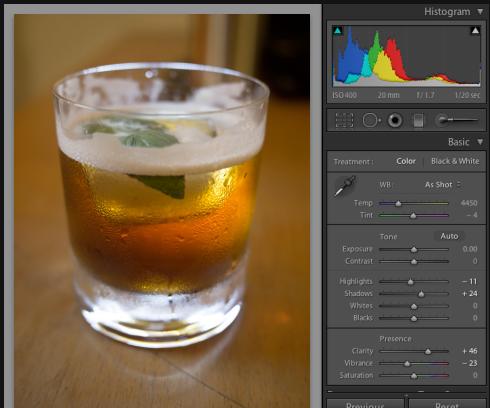


Reorganizing computation is painful



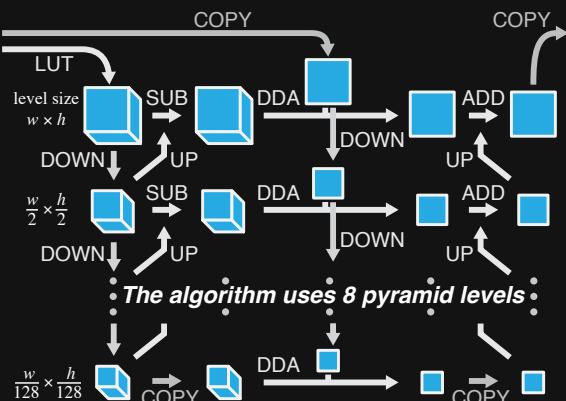
Reference:
300 lines C++

Adobe: 1500 lines
3 months of work
10x faster (vs. reference)



Same algorithm,
Different organization

Global reorganization breaks modularity



Halide

a language and compiler
for image processing & vision

[SIGGRAPH 2012,
PLDI 2013,
SIGGRAPH 2016,
SIGGRAPH 2018, ...]

Algorithm vs. Organization: 3x3 blur

```
for (int x = 0; x < input.width(); x++)
    for (int y = 0; y < input.height(); y++)
        blurH(x, y) = (input(x-1, y) + input(x, y) + input(x+1, y))/3;

for (int x = 0; x < input.width(); x++)
    for (int y = 0; y < input.height(); y++)
        blurV(x, y) = (blurH(x, y-1) + blurH(x, y) + blurH(x, y+1))/3;
```

Algorithm vs. Organization: 3x3 blur

```
for (int y = 0; y < input.height(); y++)
    for (int x = 0; x < input.width(); x++)
        blurH(x, y) = (input(x-1, y) + input(x, y) + input(x+1, y))/3;

for (int y = 0; y < input.height(); y++)
    for (int x = 0; x < input.width(); x++)
        blurV(x, y) = (blurH(x, y-1) + blurH(x, y) + blurH(x, y+1))/3;
```

Same algorithm, different organization
One of them is 15x faster

Algorithm vs. Organization: 3x3 blur

```
for (int y = 0; y < input.height(); y++)
    for (int x = 0; x < input.width(); x++)
        blurH(x, y) = (input(x-1, y) + input(x, y) + input(x+1, y))/3;

for (int y = 0; y < input.height(); y++)
    for (int x = 0; x < input.width(); x++)
        blurV(x, y) = (blurH(x, y-1) + blurH(x, y) + blurH(x, y+1))/3;
```

Hand-optimized C++

9.9 → 0.9 ms/megapixel

```
void box_filter_3x3(const Image &in, Image &blurV) {
    _m128i one_third = _mm_set_epi16(1184);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        _m128i a, b, c, sum, avg;
        _m128i blurH[(256/8)*(32/8)]; // allocate tile blurH array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            _m128i *blurxPtr = blurH;
            for (int y = 0; y < 32; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));
                    c = _mm_load_si28((__m128i*)(inPtr));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulu_ep16(sum, one_third);
                    _mm_store_si128(blurxPtr++, avg);
                    inPtr += 8;
                }
                blurxPtr = blurx;
                for (int y = 0; y < 32; y++) {
                    _m128i *outPtr = (_m128i*)(&(blurV[yTile+y][xTile]));
                    for (int x = 0; x < 256; x += 8) {
                        a = _mm_load_si128(blurxPtr+(x*32)/8);
                        b = _mm_load_si128(blurxPtr+(x*32)/8);
                        c = _mm_load_si128(blurxPtr+(x*32)/8);
                        sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                        avg = _mm_mulu_ep16(sum, one_third);
                        _mm_store_si128(outPtr++, avg);
                    }
                }
            }
        }
    }
}
```

11x faster
(quad core x86)

Tiled, fused
Vectorized
Multithreaded
Redundant computation
Near roof-line optimum

Optimized 3x3 blur in C++

```
void box_filter_3x3(const Image &in, Image &blurV) {
    _m128i one_third = _mm_set_epi16(1184);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        _m128i blurH[(256/8)*(32/8)]; // allocate tile blurH array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            _m128i *blurHPtr = blurH;
            for (int y = 0; y < 32; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));
                    c = _mm_load_si28((__m128i*)(inPtr));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulu_ep16(sum, one_third);
                    _mm_store_si128(blurHPtr++, avg);
                    inPtr += 8;
                }
                blurHPtr = blurH;
                for (int y = 0; y < 32; y++) {
                    _m128i *outPtr = (_m128i*)(&(blurV[yTile+y][xTile]));
                    for (int x = 0; x < 256; x += 8) {
                        a = _mm_load_si128(blurHPtr+(x*32)/8);
                        b = _mm_load_si128(blurHPtr+(x*32)/8);
                        c = _mm_load_si128(blurHPtr+(x*32)/8);
                        sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                        avg = _mm_mulu_ep16(sum, one_third);
                        _mm_store_si128(outPtr++, avg);
                    }
                }
            }
        }
    }
}
```

parallelism
distribute across threads
SIMD parallel vectors

Traditional languages conflate algorithm & organization

```
void box_filter_3x3(const Image &in, Image &blurV) {
    _m128i one_third = _mm_set_epi16(1184);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        _m128i blurH[(256/8)*(32/8)]; // allocate tile blurH array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            _m128i *blurHPtr = blurH;
            for (int y = 0; y < 32; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));
                    c = _mm_load_si28((__m128i*)(inPtr));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulu_ep16(sum, one_third);
                    _mm_store_si128(blurHPtr++, avg);
                    inPtr += 8;
                }
            }
        }
    }
}
```

```
void box_filter_3x3(const Image &in, Image &blurV) {
    Image blurH(in.width(), in.height()); // allocate blurH array
    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blurH(x, y) = (in(x-, y) + in(x, y) + in(x+, y))/3;
    for (int x = 0; x < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blurV(x, y) = (blurH(x, y-) + blurH(x, y) + blurH(x, y+))/3;
}
```

not readable
architecture-specific
hard to change organization
or algorithm

Optimized 3x3 blur in C++

```
void box_filter_3x3(const Image &in, Image &blurV) {
    _m128i one_third = _mm_set_epi16(1184);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        _m128i a, b, c, sum, avg;
        _m128i blurH[(256/8)*(32/8)]; // allocate tile blurH array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            _m128i *blurHPtr = blurH;
            for (int y = 0; y < 32; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));
                    c = _mm_load_si28((__m128i*)(inPtr));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulu_ep16(sum, one_third);
                    _mm_store_si128(blurHPtr++, avg);
                    inPtr += 8;
                }
            }
        }
    }
}
```

parallelism
distribute across threads
SIMD parallel vectors

locality
reorganize computation:
fuse two blurs,
compute in tiles

(Re)organizing computation is hard

Optimizing parallelism, locality requires transforming program & data structure.

What transformations are *legal*?

What transformations are *beneficial*?

libraries don't solve this:

BLAS, MKL, OpenCV, PyTorch, TensorFlow

optimized kernels compose into inefficient pipelines (no fusion)

Halide's answer: decouple algorithm from schedule

Algorithm: *what* is computed

Schedule: *where* and *when* it's computed

The algorithm defines pipelines as pure functions

Pipeline stages are functions from coordinates to values

Execution order and storage are unspecified

no explicit loops or arrays

3x3 blur as a Halide algorithm:

```
blurH(x, y) = (input(x-1, y) + input(x, y) + input(x+1, y))/3;  
blurV(x, y) = (blurH(x, y-1) + blurH(x, y) + blurH(x, y+1))/3;
```

Domain scope
of the programming model

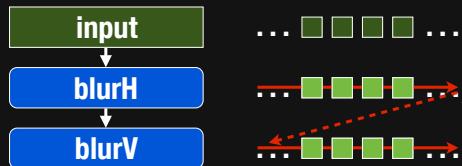
- All computation is over **regular grids**.
- not Turing complete**
 - Only **feed-forward pipelines**
 - Iterative computations are a (partial) escape hatch.
 - Iteration** must have **bounded depth**.
- Dependence must be inferable.**
 - User-defined clamping can impose tight bounds, when needed.
- Long, heterogeneous pipelines.**
 - Complex graphs, deeper than traditional stencil computations.

How can we organize this computation?

Organizing a data-parallel pipeline



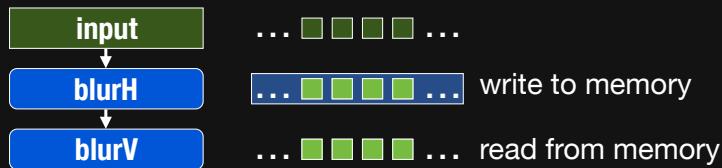
Simple loops execute **breadth-first** across stages



Simple loops execute **breadth-first** across stages

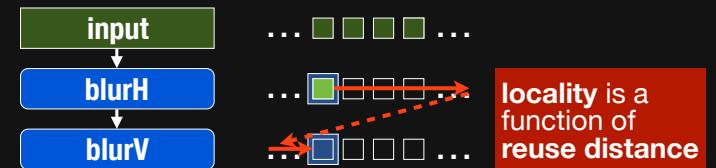


Breadth-first execution **sacrifices locality**



locality
parallelism

Breadth-first execution **sacrifices locality**



locality
parallelism

Interleaved execution (fusion) **improves locality**

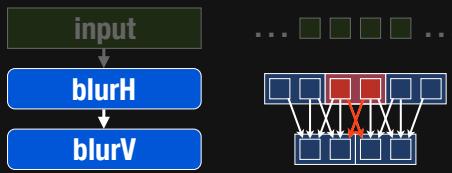


locality
parallelism

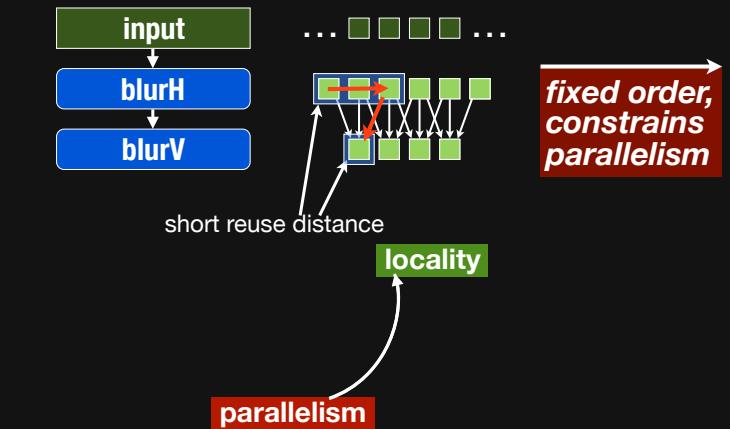
Understanding dependencies



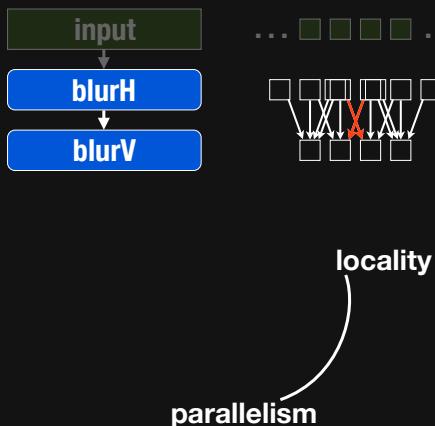
Stencils have overlapping dependencies



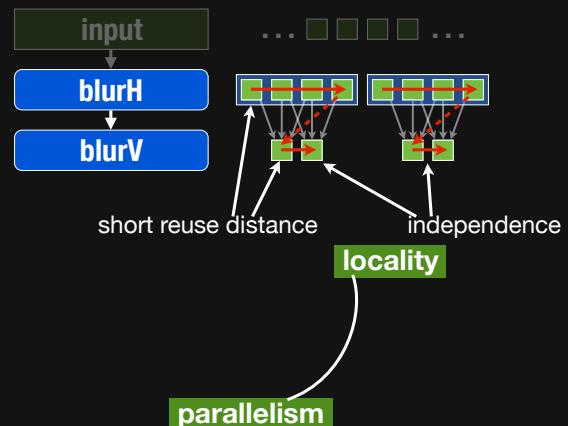
Sliding window execution **sacrifices parallelism**



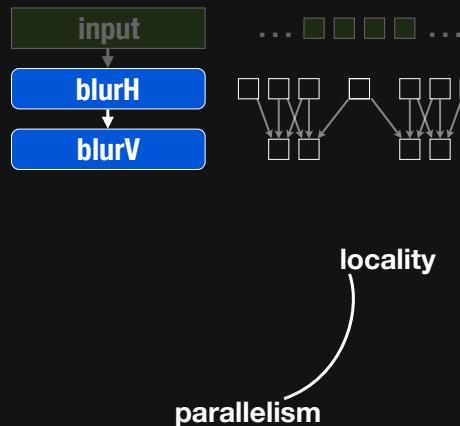
Breaking dependencies with tiling



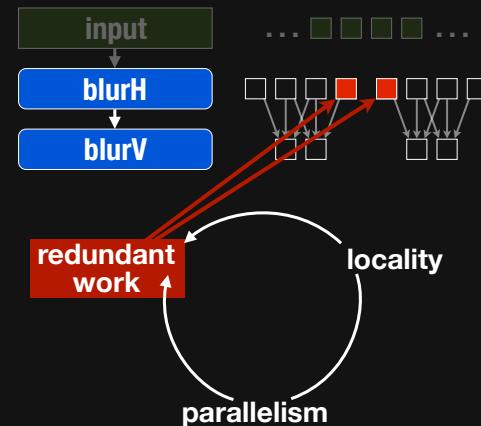
Decoupled tiles optimize **parallelism & locality**



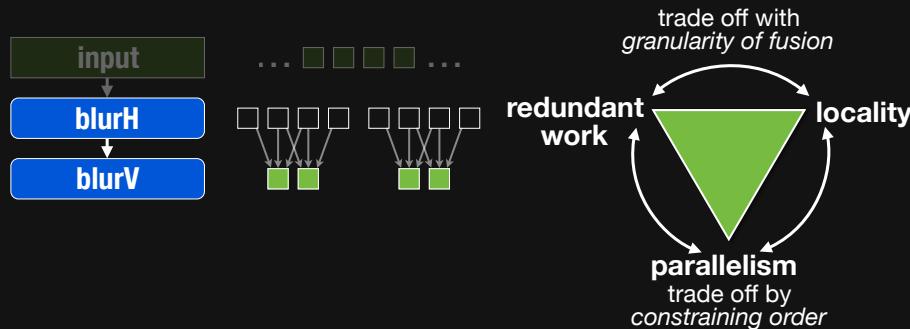
Breaking dependencies introduces redundant work



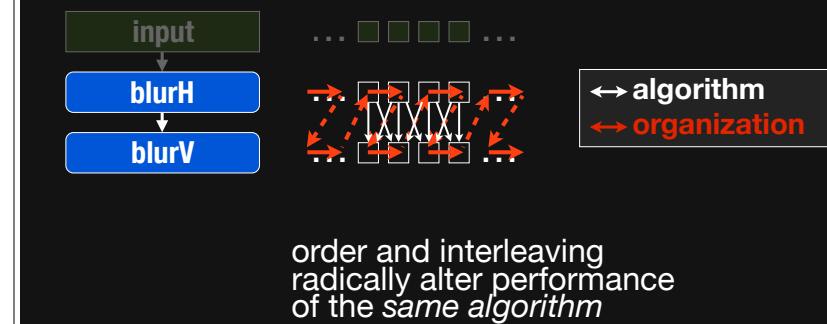
Breaking dependencies introduces redundant work



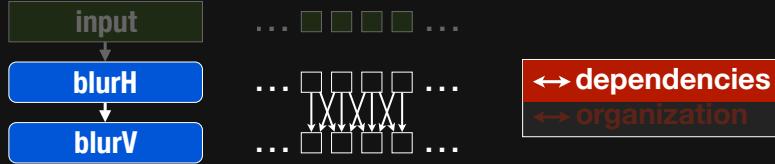
Message #1: performance requires tradeoffs



Message #2: algorithm vs. organization



Dependencies limit choices of organization



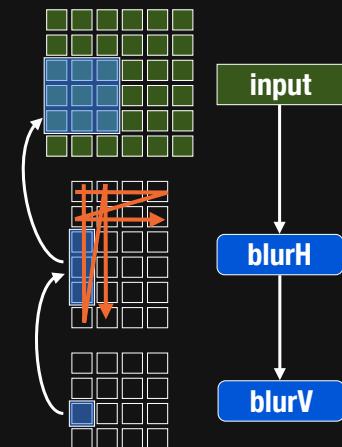
A language of schedules

The schedule defines intra-stage order, inter-stage interleaving

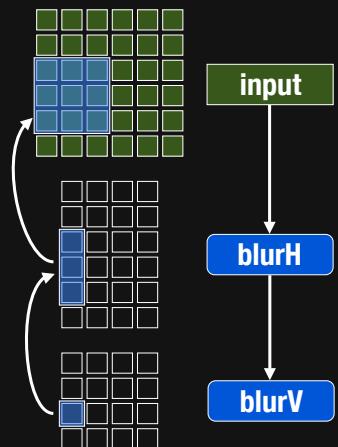
For each stage:

- 1) In what order should we compute its values?
- 2) When should we compute its inputs?

This is a **language** for scheduling choices.



The schedule defines intra-stage order, inter-stage interleaving



Schedule primitives **compose** to create many organizations



Schedule primitives **compose** to create many organizations

```
void box_filter_3x3(const Image &in, Image &blurV) {
    _m128i one_third = _mm_set1_epi16(21846);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        _m128i a, b, c, sum, avg;
        _m128i blurH((256/8)*(3+));
        // allocate tile blurH array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            _m128i *blurHPtr = blurH;
            for (int y = -1; y < 3+; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 4) {
                    a = _mm_load_si128((__m128i*)(inPtr-));
                    b = _mm_load_si128((__m128i*)(inPtr+));
                    c = _mm_load_si128((__m128i*)(inPtr+));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulhi_epi16(sum, one_third);
                    _mm_store_si128(blurHPtr++, avg);
                    inPtr += 4;
                }
            }
            blurHPtr = blurH;
        }
        for (int y = 0; y < 32; y++) {
            _m128i *outPtr = (_m128i*)(&(blurV[yTile+y][xTile]));
            for (int x = 0; x < 256; x += 4) {
                a = _mm_load_si128(blurHPtr+(2*256)/4);
                b = _mm_load_si128(blurHPtr+(2*256)/4);
                c = _mm_load_si128(blurHPtr++);
                sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                avg = _mm_mulhi_epi16(sum, one_third);
                _mm_store_si128(outPtr++, avg);
            }
        }
    }
}
```

Halide

0.9 ms/megapixel

```
Func box_filter_3x3(Func in) {
    Func blurH, blurV;
    Var x, y, xi, yi;

    // The algorithm - no storage, order
    blurH(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;
    blurV(x, y) = (blurH(x, y-1) + blurH(x, y) + blurH(x, y+1))/3;

    // The schedule - defines order, locality; implies storage
    blurV.tile(x, y, xi, yi, 256, 32)
        .vectorize(xi, 8).parallel(y);
    blurH.compute_at(blurV, x).store_at(blurV, x).vectorize(x, 8);

    return blurV;
}
```

Halide

0.9 ms/megapixel

```
Func box_filter_3x3(Func in) {
    Func blurH, blurV;
    Var x, y, xi, yi;

    // The algorithm - no storage, order
    blurH(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;
    blurV(x, y) = (blurH(x, y-1) + blurH(x, y) + blurH(x, y+1))/3;

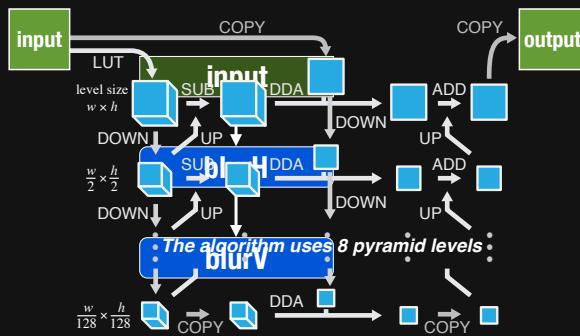
    // The schedule - defines order, locality; implies storage
    blurV.tile(x, y, xi, yi, 256, 32)
        .vectorize(xi, 8).parallel(y);
    blurH.compute_at(blurV, x).store_at(blurV, x).vectorize(x, 8);
    return blurV;
}
```

C++

0.9 ms/megapixel

```
void box_filter_3x3(const Image &in, Image &blurV) {
    _m128i one_third = _mm_set1_epi16(21846);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        _m128i a, b, c, sum, avg;
        _m128i blurH((256/8)*(3+));
        // allocate tile blurH array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            _m128i *blurHPtr = blurH;
            for (int y = -1; y < 3+; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 4) {
                    a = _mm_load_si128((__m128i*)(inPtr-));
                    b = _mm_load_si128((__m128i*)(inPtr+));
                    c = _mm_load_si128((__m128i*)(inPtr+));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulhi_epi16(sum, one_third);
                    _mm_store_si128(blurHPtr++, avg);
                    inPtr += 4;
                }
            }
            blurHPtr = blurH;
        }
        for (int y = 0; y < 32; y++) {
            _m128i *outPtr = (_m128i*)(&(blurV[yTile+y][xTile]));
            for (int x = 0; x < 256; x += 4) {
                a = _mm_load_si128(blurHPtr+(2*256)/4);
                b = _mm_load_si128(blurHPtr+(2*256)/4);
                c = _mm_load_si128(blurHPtr++);
                sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                avg = _mm_mulhi_epi16(sum, one_third);
                _mm_store_si128(outPtr++, avg);
            }
        }
    }
}
```

Organization requires global tradeoffs



local Laplacian filters
[Parker et al. 2011]

LUT: look-up table
 $O(x,y) \leftarrow \text{lut}(I(x,y) - k\sigma)$

UP: upsample
 $T_1(2x,2y) \leftarrow I(x,y)$
 $T_2 \leftarrow T_1 \otimes_{\downarrow} [1 \ 3 \ 3 \ 1]$
 $O \leftarrow T_2 \otimes_{\downarrow} [1 \ 3 \ 3 \ 1]$

ADD: addition
 $O(x,y) \leftarrow I_1(x,y) + I_2(x,y)$

DOWNSAMPLE
 $T_1 \leftarrow I \otimes_{\downarrow} [1 \ 3 \ 3 \ 1]$
 $T_2 \leftarrow T_1 \otimes_{\downarrow} [1 \ 3 \ 3 \ 1]$
 $O(x,y) \leftarrow T_2(2x,2y)$

SUB: subtraction
 $O(x,y) \leftarrow I_1(x,y) - I_2(x,y)$

DDA: data-dependent access
 $k \leftarrow \text{floor}(I_1(x,y) / \sigma)$
 $\alpha \leftarrow (I_1(x,y) / \sigma) - k$
 $O(x,y) \leftarrow (1-\alpha) I_1(x,y,k) + \alpha I_1(x,y,k+1)$

Adobe: 1500 lines

expert-tuned C++

multi-threaded, SSE

3 months of work

10x faster than original C++

Halide: 60 lines

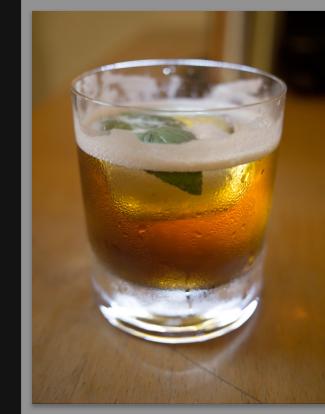
1 day

Halide vs. Adobe:

2x faster on same CPU

10x faster on GPU

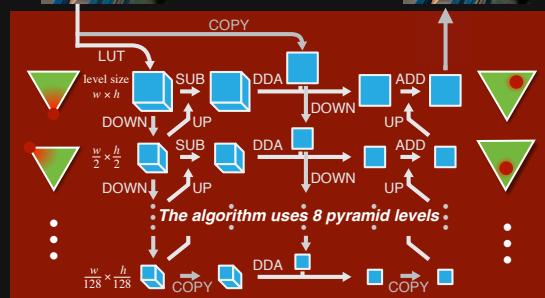
Local Laplacian Filters in Adobe Photoshop



Adobe: 1500 lines
expert-tuned C++
multi-threaded, SSE
3 months of work
10x faster than original C++

Halide: 60 lines
1 day

Halide vs. Adobe:
2x faster on same CPU
10x faster on GPU



Halide is all around you!

open source at <http://halide-lang.org>



Google

> 1000 pipelines
10s of kLOC in production



facebook

PYTORCH

Tensor Comprehensions

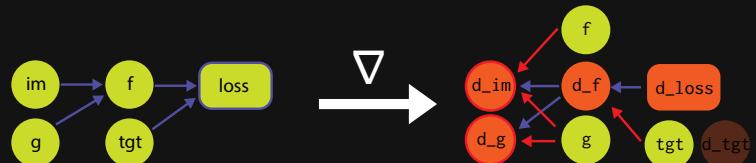
You Tube



What's next?

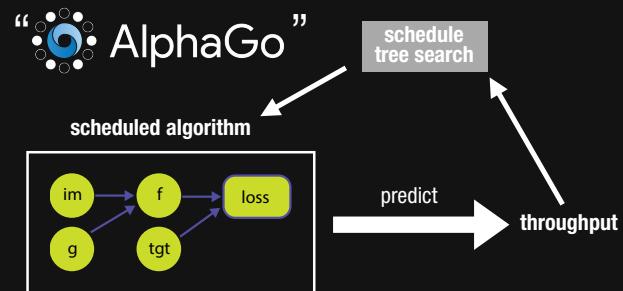
What's next?

Differentiable programming languages:
bringing gradient-based optimization
to arbitrary programs.



What's next?

Training compilers to **automatically optimize organization** using **reinforcement learning**.



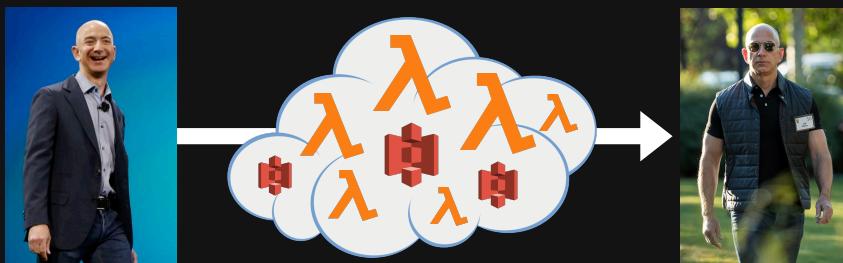
What's next?

Building the “**programmable GPU**”
of image processing & vision.



What's next?

Turning **AWS Lambda** into a **supercomputer** accessible from your laptop



Where you can learn more:

CS 164 - Compilers & Programming Languages

CS 294-jrk - Domain-Specific Languages

Talk to me about research!

Thank you!

